

# Methods and Algorithms

Spring 2026



## From Rules to Trees

- Decision trees encode if-then-else rules
- Each node splits data based on a feature threshold
- Leaves contain predictions (class or value)

## Key Questions

- How to choose the best split?
- When to stop splitting?
- How to make predictions?

*Decision trees: the building blocks of Random Forests*

**Gini Impurity** measures class mixture at a node:

$$G = 1 - \sum_{k=1}^K p_k^2$$

where  $p_k$  is the proportion of class  $k$  samples.

**Properties:**

- $G = 0$ : pure node (all samples same class)
- $G = 0.5$ : maximum impurity for binary classification
- Lower Gini = better split

*Gini impurity: probability of misclassifying a random sample*

**Entropy** measures disorder:

$$H = - \sum_{k=1}^K p_k \log_2(p_k)$$

**Information Gain:**

$$IG = H(\text{parent}) - \sum_j \frac{n_j}{n} H(\text{child}_j)$$

**Comparison:**

- Gini: faster to compute, tends to isolate most frequent class
- Entropy: more balanced trees, slightly slower
- In practice: similar performance

*Both criteria aim to create pure child nodes*

**For regression**, use Mean Squared Error:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2$$

**Split quality**:

$$\text{Reduction} = \text{MSE}(\text{parent}) - \sum_j \frac{n_j}{n} \text{MSE}(\text{child}_j)$$

**Leaf prediction**: mean of samples in leaf

*Trees can handle both classification and regression tasks*

## Recursive Partitioning:

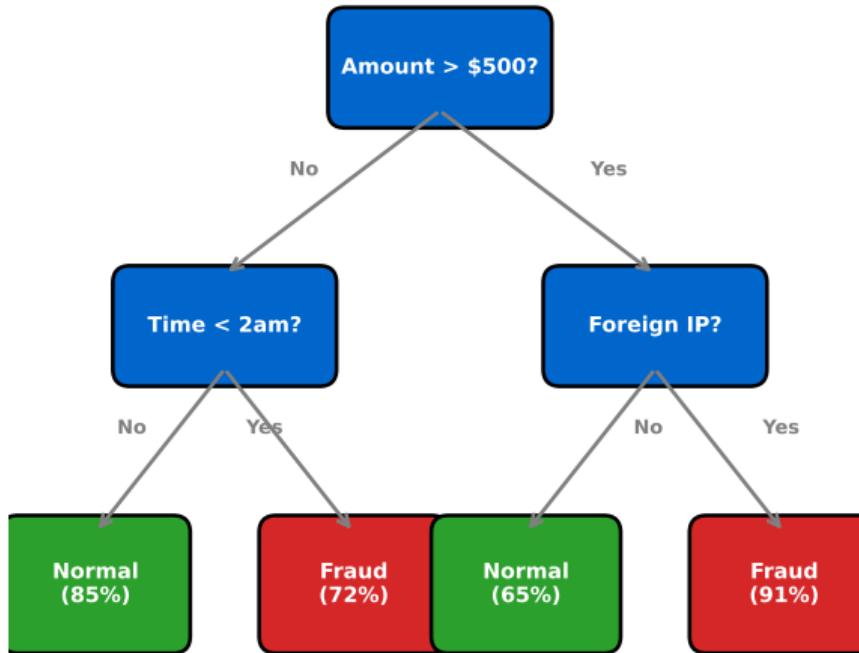
- ① Start with all samples at root
- ② For each feature and threshold:
  - Calculate impurity reduction
  - Select split with maximum reduction
- ③ Create child nodes with split samples
- ④ Recurse until stopping criterion met

## Stopping Criteria:

- Maximum depth reached
- Minimum samples per leaf
- No improvement in impurity

*Greedy algorithm: locally optimal splits at each step*

## Decision Tree for Fraud Detection



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/01\\_decision\\_tree](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/01_decision_tree)

Each path through the tree represents a fraud detection rule

## **Problem with Single Trees:**

- High variance: small data changes → very different trees
- Prone to overfitting
- Unstable predictions

## **Ensemble Solution:**

- Train multiple diverse models
- Combine predictions
- Reduce variance while maintaining low bias

*“Wisdom of crowds”: aggregate many weak learners into strong learner*

**Bootstrap:** sample with replacement from original data

**Properties:**

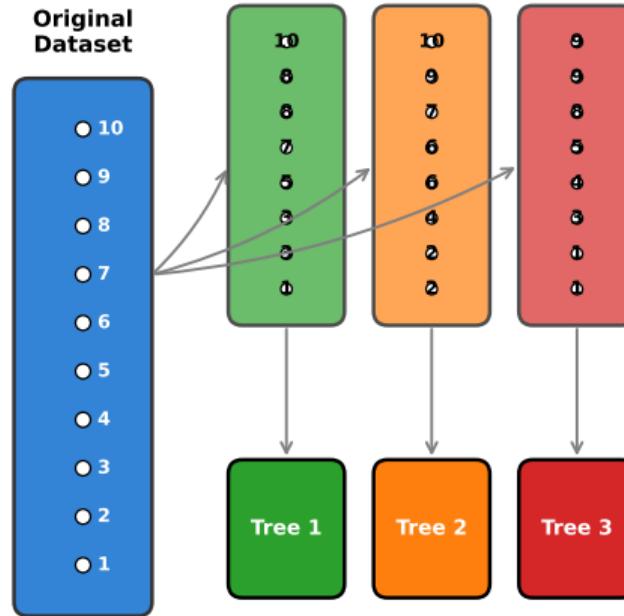
- Each sample: same size as original ( $n$  observations)
- Expected unique samples:  $\approx 63.2\%$  (probability  $1 - (1 - 1/n)^n$ )
- Remaining  $\sim 37\%$ : out-of-bag (OOB) samples

**Effect:**

- Each tree sees different data subset
- Creates diversity among trees
- OOB samples provide validation

*Bootstrap: key ingredient for reducing variance through aggregation*

## Bootstrap Aggregating (Bagging)



Each tree trained on ~63% unique samples (with replacement)

[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/03\\_bootstrap](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/03_bootstrap)

Bootstrap Aggregating: train on random samples, aggregate predictions

## Variance Reduction by Averaging

For  $B$  independent predictions with variance  $\sigma^2$ :

$$\text{Var} \left( \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x) \right) = \frac{\sigma^2}{B}$$

With correlation  $\rho$ :

$$\text{Var} = \rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$$

**Key insight:** Reduce correlation between trees to maximize variance reduction

*Lower correlation between trees = greater ensemble benefit*

## Two Sources of Randomness:

- ① **Bootstrap sampling**: each tree trained on random sample
- ② **Feature randomization**: each split considers random subset

## Feature Subset Size (at each split):

- Classification:  $\sqrt{p}$  features (default)
- Regression:  $p/3$  features (default)
- Decorrelates trees more than bagging alone

*Feature randomization: Breiman's key innovation over bagging*

## Training:

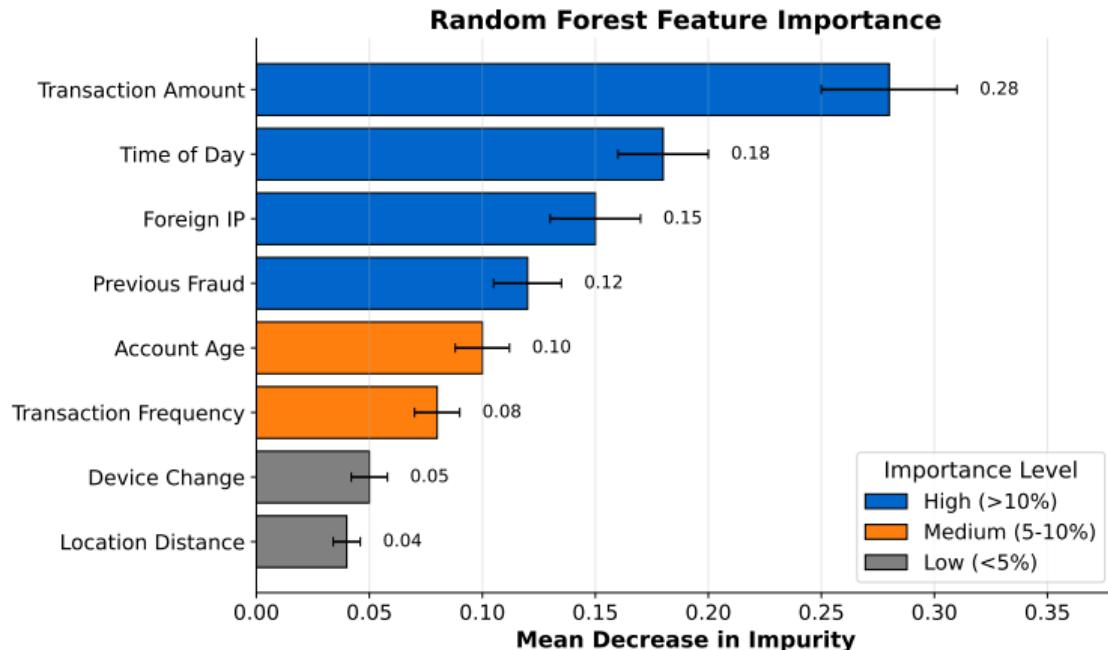
- ① For  $b = 1$  to  $B$  trees:

- Draw bootstrap sample of size  $n$
- Grow tree:
  - At each node, select  $m$  features randomly
  - Find best split among  $m$  features
  - Split until stopping criterion

## Prediction:

- Classification: majority vote across trees
- Regression: average predictions

*Typical: 100-500 trees, but more trees never hurts (just slower)*



*Mean Decrease in Impurity: sum of impurity reductions from splits on feature*

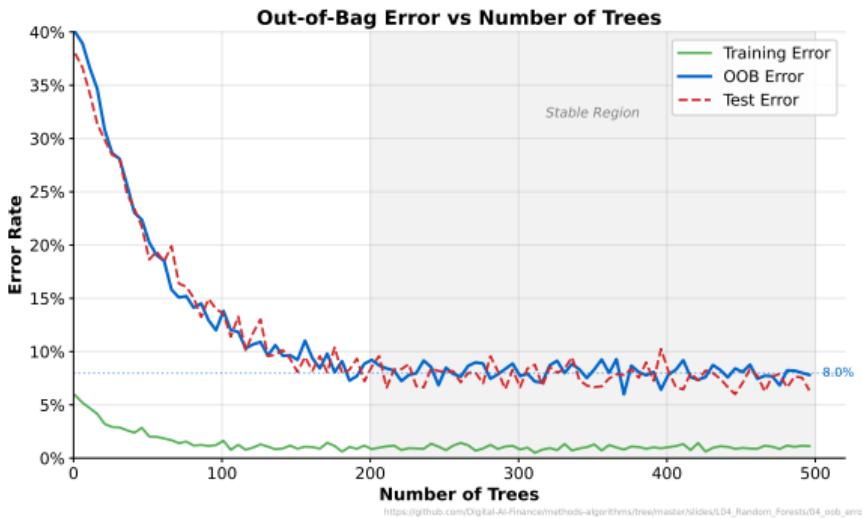
## 1. Mean Decrease in Impurity (MDI):

- Sum of Gini/entropy reductions from splits on feature
- Fast to compute (comes free from training)
- Bias toward high-cardinality features

## 2. Permutation Importance:

- Permute feature values, measure accuracy drop
- More reliable, less biased
- Slower (requires re-evaluation)

*Permutation importance preferred for final feature selection*



*OOB error: free cross-validation using samples not in bootstrap*

**For each observation  $i$ :**

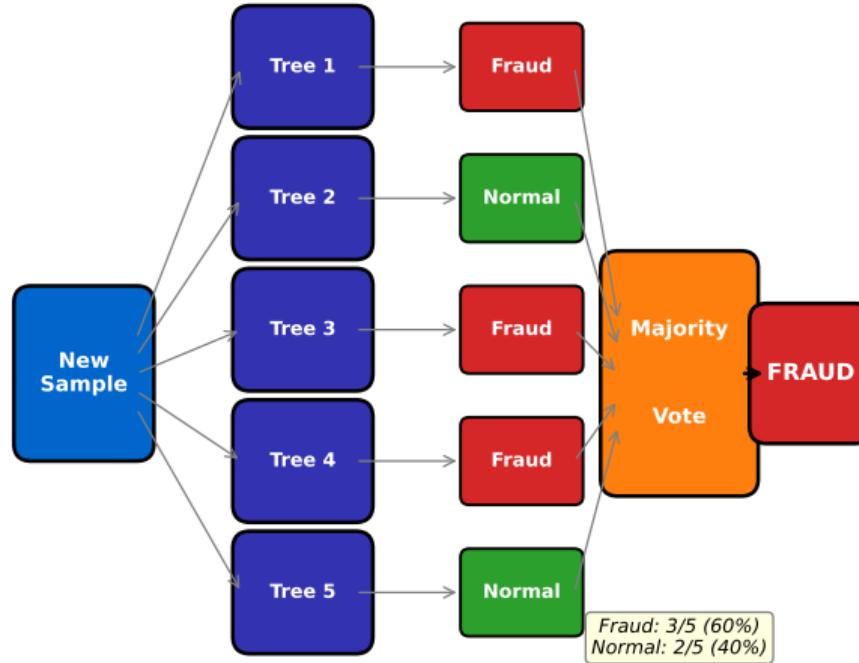
- ① Identify trees where  $i$  was OOB (not in bootstrap sample)
- ② Aggregate predictions from only those trees
- ③ Compare to true label

**Benefits:**

- No separate validation set needed
- Uses  $\sim 37\%$  of trees per sample
- Unbiased estimate of generalization error

*OOB error converges to leave-one-out cross-validation error*

## Ensemble Voting (Classification)



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/05\\_ensemble\\_voting](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/05_ensemble_voting)

Classification: majority vote. Regression: average prediction

### Expected Prediction Error:

$$\mathbb{E}[(y - \hat{f}(x))^2] = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$

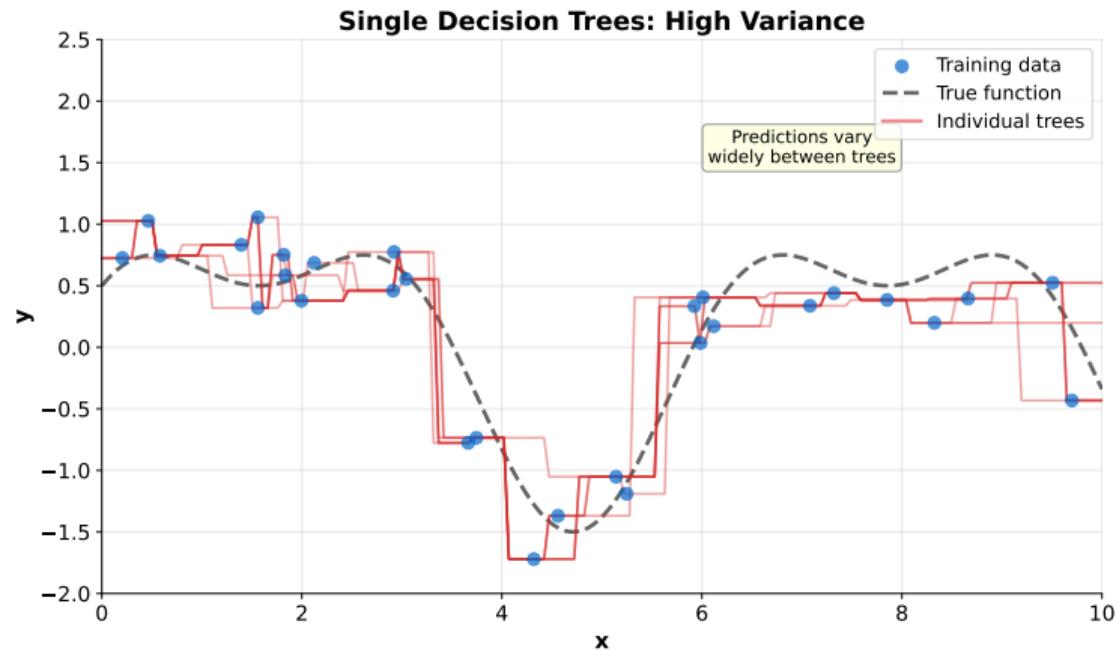
### Single Tree:

- Low bias (can fit complex patterns)
- High variance (sensitive to training data)

### Random Forest:

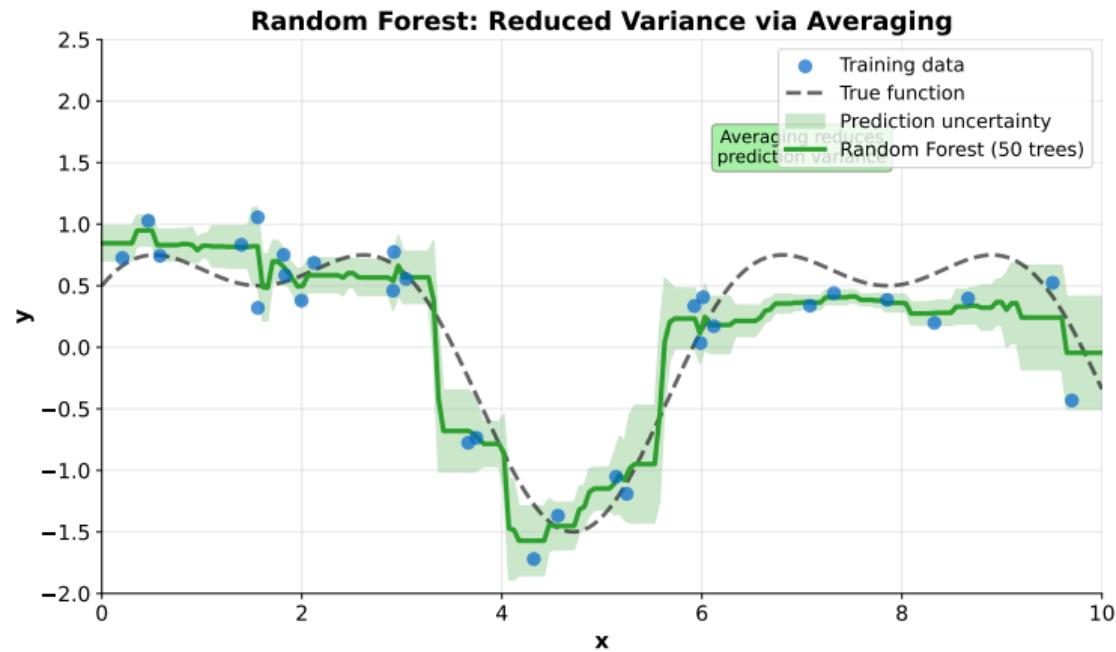
- Bias: similar to single tree
- Variance: reduced by  $\approx$  factor of  $1/B$  (with decorrelation)

*Ensembles reduce variance without increasing bias*



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/06a\\_single\\_tree\\_variance](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/06a_single_tree_variance)

*Individual trees overfit to their bootstrap samples, producing erratic predictions*



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/06b\\_random\\_forest\\_variance](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/06b_random_forest_variance)

Averaging decorrelated trees dramatically reduces prediction variance

## **n\_estimators** (number of trees):

- More trees = lower variance, never overfits
- Diminishing returns after 100-500 trees
- Cost: linear increase in training/prediction time

## **Guidelines:**

- Start with 100, increase if OOB error still decreasing
- For production: balance accuracy vs. latency
- More trees always better (if time permits)

*Unlike most hyperparameters, more trees cannot hurt accuracy*

**max\_depth:** Maximum tree depth

- Deeper = more complex patterns, higher variance
- Default: unlimited (grow full trees)

**min\_samples\_split:** Minimum samples to split

- Higher = simpler trees, more regularization
- Default: 2 (full trees)

**min\_samples\_leaf:** Minimum samples in leaf

- Higher = smoother predictions
- Default: 1 (full trees)

*Full trees (default) often work well due to bagging's variance reduction*

**max\_features**: Features considered at each split

- Lower = more decorrelated trees, higher bias
- Higher = less decorrelated, lower bias

**Defaults:**

- Classification:  $\sqrt{p}$  (e.g., 10 features  $\rightarrow 3$ )
- Regression:  $p/3$  (e.g., 30 features  $\rightarrow 10$ )

**Tuning:**

- Try:  $\sqrt{p}$ ,  $\log_2(p)$ ,  $p/3$
- Cross-validate to find optimal

*Feature randomization is key differentiator from bagged trees*

## **Advantages:**

- Excellent accuracy out-of-the-box
- Handles mixed feature types
- Built-in feature importance
- Robust to outliers and missing values
- Parallelizable (trees independent)

## **Limitations:**

- Less interpretable than single tree
- Memory intensive (stores all trees)
- Slower prediction than linear models
- Cannot extrapolate beyond training range

*Random Forests: excellent default choice for tabular data*

Aspect	RF	Single Tree	Logistic	KNN
Accuracy	High	Medium	Medium	Medium
Interpretability	Medium	High	High	Low
Training Speed	Medium	Fast	Fast	Fast
Feature Importance	Yes	Yes	Yes	No
Non-linear	Yes	Yes	No	Yes

*RF trades some interpretability for significant accuracy gains*

### **Use When:**

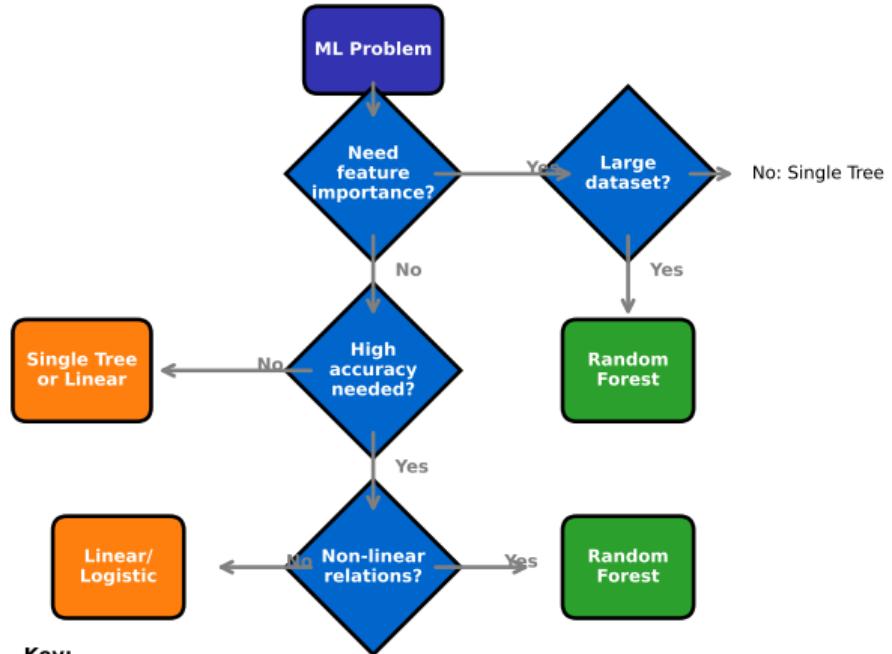
- Tabular data with mixed feature types
- Non-linear relationships expected
- Feature importance needed
- Out-of-the-box performance matters

### **Consider Alternatives When:**

- Need fully interpretable model (use single tree)
- Very high-dimensional sparse data (use linear models)
- Extrapolation required (use parametric models)
- Need fastest prediction (use linear/shallow tree)

*Random Forest: often the first model to try on tabular data*

## When to Use Random Forests



**Key:**

Random Forest: Best for accuracy + feature importance

Alternative: When interpretability is paramount

[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/07\\_decision\\_flowchart](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/07_decision_flowchart)

*Start with Random Forest; switch if specific constraints require it*

## Open the Colab Notebook

- Exercise 1: Train a decision tree on credit data
- Exercise 2: Build a random forest and analyze feature importance
- Exercise 3: Tune hyperparameters with cross-validation

Link: <https://colab.research.google.com/> [TBD]

## Classification:

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=100, max_features='sqrt',
                            oob_score=True, random_state=42)
rf.fit(X_train, y_train)
print(f"OOB Score: {rf.oob_score_:.3f}")
```

## Key Parameters:

- n\_estimators: number of trees
- max\_features: features per split
- oob\_score: compute OOB error
- n\_jobs: parallel trees (-1 for all cores)

*Set random\_state for reproducibility*

## Core Concepts:

- Ensemble of decision trees with bootstrap + feature randomization
- Reduces variance while maintaining low bias
- OOB error provides free cross-validation

## Practical Takeaways:

- Excellent default for tabular data
- Feature importance aids interpretation
- More trees never hurts (just slower)
- Hyperparameter tuning usually optional

Next: *PCA and t-SNE for dimensionality reduction*

### Textbooks:

- James et al. (2021). *ISLR*, Chapter 8: Tree-Based Methods
- Hastie et al. (2009). *ESL*, Chapter 15: Random Forests

### Original Papers:

- Breiman, L. (2001). Random Forests. *Machine Learning*, 45(1), 5-32.
- Breiman, L. (1996). Bagging Predictors. *Machine Learning*, 24(2), 123-140.

### Documentation:

- scikit-learn: `sklearn.ensemble.RandomForestClassifier`

*Breiman's 2001 paper: one of the most cited in ML*