

# Linear Algebra for Machine Learning

## Mini-Lecture: The Mathematical Foundation

### Methods and Algorithms

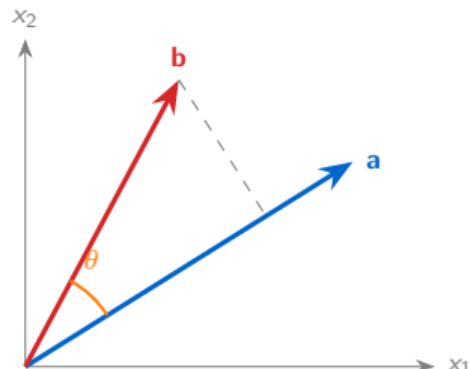
MSc Data Science



XKCD #1838 by Randall Munroe (CC BY-NC 2.5)

# Vectors: Direction and Magnitude

- A **vector**  $x \in \mathbb{R}^n$  represents one data point — e.g. a loan applicant: [income, age, debt\_ratio, credit\_score]
- **Dot product**  $a^\top b = \sum_i a_i b_i$  measures similarity between two feature vectors
- **Norm**  $\|x\| = \sqrt{x^\top x}$  measures magnitude — large norm = extreme observation



---

Every ML algorithm starts with data encoded as vectors — master this representation.

- The **design matrix**  $\mathbf{X} \in \mathbb{R}^{n \times p}$ :  $n$  observations (rows),  $p$  features (columns)
- Each row is one customer; each column is one attribute (income, age, ...)
- Matrix–vector product computes **all predictions at once**:

$$\hat{\mathbf{y}} = \mathbf{X}\boldsymbol{\beta} = \begin{pmatrix} x_{11} & \cdots & x_{1p} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{np} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \vdots \\ \beta_p \end{pmatrix}$$

- Finance example: 500 stocks  $\times$  10 factors  $\Rightarrow$  **factor exposure matrix**

---

Think of  $\mathbf{X}\boldsymbol{\beta}$  as “apply the model to every observation simultaneously.”

- **Transpose  $\mathbf{A}^\top$ :** flips rows and columns — needed for  $\mathbf{X}^\top \mathbf{X}$
- **Inverse  $\mathbf{A}^{-1}$ :** “undo” a transformation —  $\mathbf{A}\mathbf{A}^{-1} = \mathbf{I}$
- The product  $\mathbf{X}^\top \mathbf{X}$  is covariance-like and powers the OLS solution (derived in L01)
- $\det(\mathbf{A}) = 0$  means **singular** — no unique inverse, features are linearly dependent

## OLS Normal Equation (Preview)

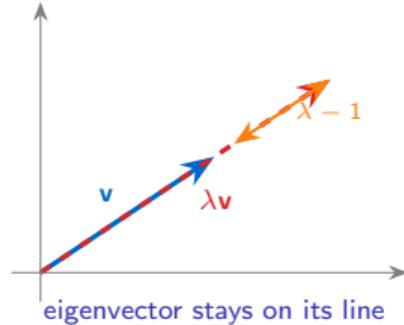
$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

$\mathbf{X}^\top \mathbf{X}$  is the single most important matrix product in regression — you will see it in L01.

- **Eigen-equation:**  $\mathbf{Av} = \lambda v$
- $\lambda$  (eigenvalue) = *how much* the vector is stretched
- $v$  (eigenvector) = *which direction* stays unchanged
- PCA uses eigenvectors of  $\Sigma$  to find principal components (details in L05)

## Finance Insight

Eigenvalues of a correlation matrix reveal independent **risk factors**  
— large  $\lambda$  = large variance explained.



---

Eigenvectors point where data varies most — PCA ranks them by eigenvalue magnitude.

- **Matrix multiplication** = linear transformation of the input space
- $\hat{y} = \mathbf{X}\beta$  transforms the feature space into a prediction space
- **Rank** of  $\mathbf{X}$  = number of linearly independent columns
- If  $\text{rank}(\mathbf{X}) < p$ : **multicollinearity** — OLS becomes unstable

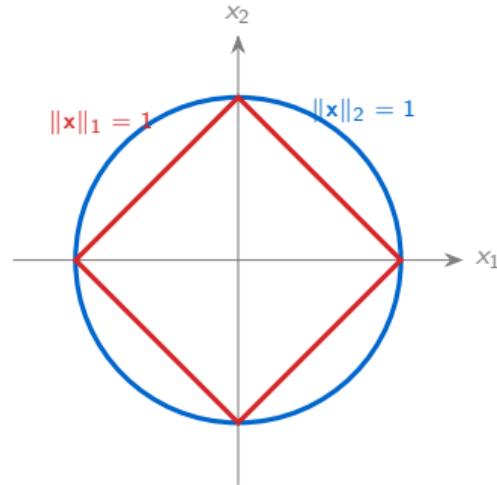
## Practical Check

$$\text{rank}(\mathbf{X}) = p \iff \mathbf{X}^\top \mathbf{X} \text{ is invertible} \iff \text{unique OLS solution exists}$$

---

Multicollinearity is the #1 numerical pitfall in linear regression — regularization (L01) fixes it.

- **L2 norm** (Euclidean):  $\|\mathbf{x}\|_2 = \sqrt{\sum x_i^2}$  — used in Ridge regression
- **L1 norm** (Manhattan):  $\|\mathbf{x}\|_1 = \sum |x_i|$  — used in Lasso (sparse solutions)
- Frobenius norm for matrices:  $\|\mathbf{A}\|_F = \sqrt{\sum_{ij} a_{ij}^2}$
- Finance: tracking error =  $\|\mathbf{r}_{\text{port}} - \mathbf{r}_{\text{bench}}\|_2$



L1 produces sparse models (feature selection); L2 shrinks all coefficients evenly.

- **Covariance matrix:**  $\Sigma_{ij} = \text{Cov}(R_i, R_j)$  captures pairwise return co-movements
- **Portfolio variance:**  $\sigma_p^2 = \mathbf{w}^\top \boldsymbol{\Sigma} \mathbf{w}$ , where  $\mathbf{w}$  = asset weights
- Eigendecomposition of  $\boldsymbol{\Sigma}$  reveals independent **risk factors**

$$\boldsymbol{\Sigma} = \mathbf{V} \boldsymbol{\Lambda} \mathbf{V}^\top \quad \text{where } \boldsymbol{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_p)$$

- **Cholesky decomposition**  $\boldsymbol{\Sigma} = \mathbf{L} \mathbf{L}^\top$  generates correlated Monte Carlo samples

---

Nearly every risk model in quantitative finance starts with  $\boldsymbol{\Sigma}$  and its eigenstructure.

1. **Vectors** are data points; **matrices** are datasets —  $\mathbf{X} \in \mathbb{R}^{n \times p}$
2.  $\mathbf{X}^\top \mathbf{X}$  powers the OLS solution and appears throughout regression (L01)
3. **Eigendecomposition** reveals variance directions — foundation of PCA (L05)
4. **Norms** (L1, L2) drive regularization and distance-based methods (L01, L03)

## Coming Up

P02: Supervised & Unsupervised Learning — the two paradigms that organize this entire course.

---

These four ideas — vectors,  $\mathbf{X}^\top \mathbf{X}$ , eigenvalues, norms — recur in every lecture.