

Methods and Algorithms

Spring 2026

1 K-Nearest Neighbors

2 Practice

3 K-Means Clustering

4 Comparison and Applications

5 Summary

Key Insight

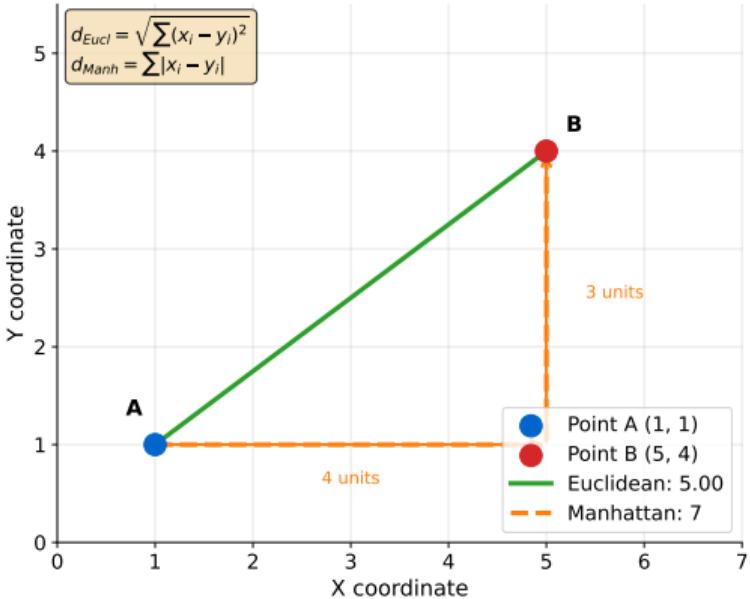
- No explicit model training (store all data)
- Classification by majority vote of K nearest neighbors
- “Lazy” because work is done at prediction time

The Algorithm

- ① Store all training examples
- ② For new query x : find K nearest training points
- ③ Return majority class among neighbors

Instance-based learning: the training data IS the model

Distance Metrics: Euclidean vs Manhattan



https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L03_KNN_KMeans/02_distance_metrics

Common: Euclidean $\sqrt{\sum(x_i - y_i)^2}$, Manhattan $\sum |x_i - y_i|$

Minkowski Distance

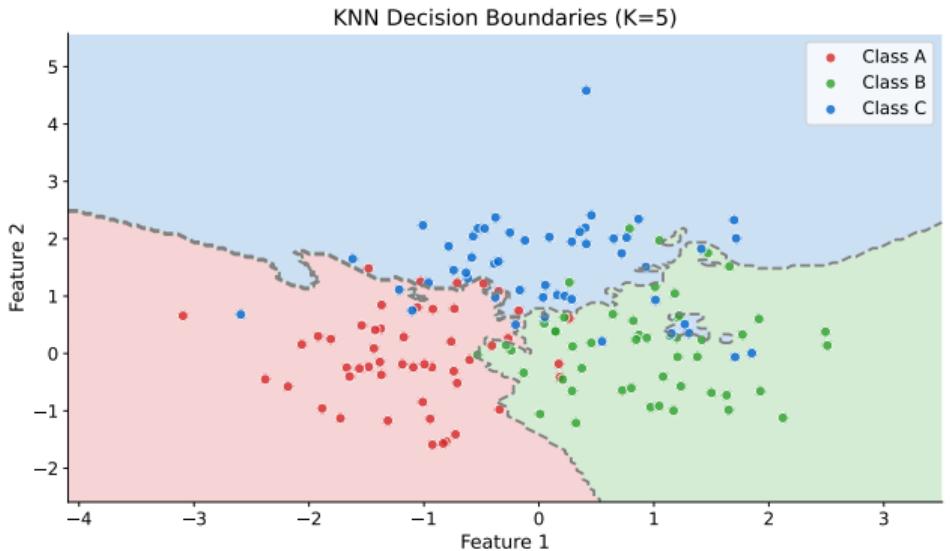
$$d_p(x, y) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

- $p = 1$: Manhattan (L1)
- $p = 2$: Euclidean (L2)
- $p = \infty$: Chebyshev (max absolute difference)

Choosing p

- $p = 2$: Default, works well in most cases
- $p = 1$: More robust to outliers
- Higher p : Sensitive to large single differences

In high dimensions, all distances become similar (curse of dimensionality)



Boundary Properties

- Non-linear, locally adaptive
- Small K: complex boundary, may overfit
- Large K: smoother boundary, may underfit

The Bias-Variance Trade-off

- $K = 1$: High variance, low bias (very flexible)
- $K = n$: High bias, low variance (always predicts majority class)

Practical Guidelines

- Start with $K = \sqrt{n}$ where n is training size
- Use odd K for binary classification (avoid ties)
- Cross-validation to find optimal K

Common Choices: $K \in \{3, 5, 7, 11\}$

Small K for complex patterns, larger K for noisy data

Problem with Equal Voting

- All K neighbors have equal influence
- A distant neighbor counts as much as closest neighbor

Solution: Distance Weighting

$$w_i = \frac{1}{d(x, x_i)^2}$$

- Closer neighbors get higher weight
- Reduces sensitivity to K choice

In scikit-learn

- `weights='uniform'`: equal weights (default)
- `weights='distance'`: inverse distance weighting

Distance weighting often improves performance

Why Scaling Matters

Without scaling:

- Income: ranges 20,000–200,000
- Age: ranges 20–80
- Distance dominated by income (larger scale)!

Scaling Methods

- **Standardization:** $z = \frac{x-\mu}{\sigma}$ (mean=0, std=1)
- **Min-Max:** $x' = \frac{x-x_{min}}{x_{max}-x_{min}}$ (range [0,1])

Rule: Always scale features for distance-based methods!

StandardScaler for Gaussian-like features, MinMaxScaler for bounded

The Problem

In high dimensions:

- All points become approximately equidistant
- “Nearest neighbor” becomes meaningless
- Exponentially more data needed

Solutions

- **Dimensionality reduction:** PCA before KNN
- **Feature selection:** keep only relevant features
- **Use domain knowledge:** select meaningful features

KNN works best with moderate number of features (<20)

Time Complexity

- Training: $O(1)$ – just store data!
- Prediction: $O(nd)$ for brute force (n samples, d features)

Acceleration Techniques

- **KD-Tree**: $O(d \log n)$ average for low d
- **Ball Tree**: Works better in higher dimensions
- **Approximate NN**: Trade accuracy for speed

In scikit-learn

- `algorithm='auto'`: automatically chooses best
- `algorithm='brute'`: force brute force

For large datasets: consider approximate methods or trees

Classification

- `from sklearn.neighbors import KNeighborsClassifier`
- `knn = KNeighborsClassifier(n_neighbors=5)`
- `knn.fit(X_train, y_train)`
- `y_pred = knn.predict(X_test)`

Key Parameters

- `n_neighbors`: K value
- `weights`: 'uniform' or 'distance'
- `metric`: 'euclidean', 'manhattan', etc.
- `algorithm`: 'auto', 'ball_tree', 'kd_tree', 'brute'

Also available: `KNeighborsRegressor` for regression tasks

Open the Colab Notebook

- Exercise 1: Implement KNN classifier from scratch
- Exercise 2: Apply K-Means to customer segmentation data
- Exercise 3: Compare distance metrics and k values

Link: <https://colab.research.google.com/> [TBD]

Goal: Partition n points into K clusters

Objective: Minimize within-cluster sum of squares (WCSS)

$$\sum_{k=1}^K \sum_{x \in C_k} \|x - \mu_k\|^2$$

where μ_k is the centroid of cluster C_k

Key Insight

- Each point assigned to nearest centroid
- Centroids are cluster means
- Iterative refinement until convergence

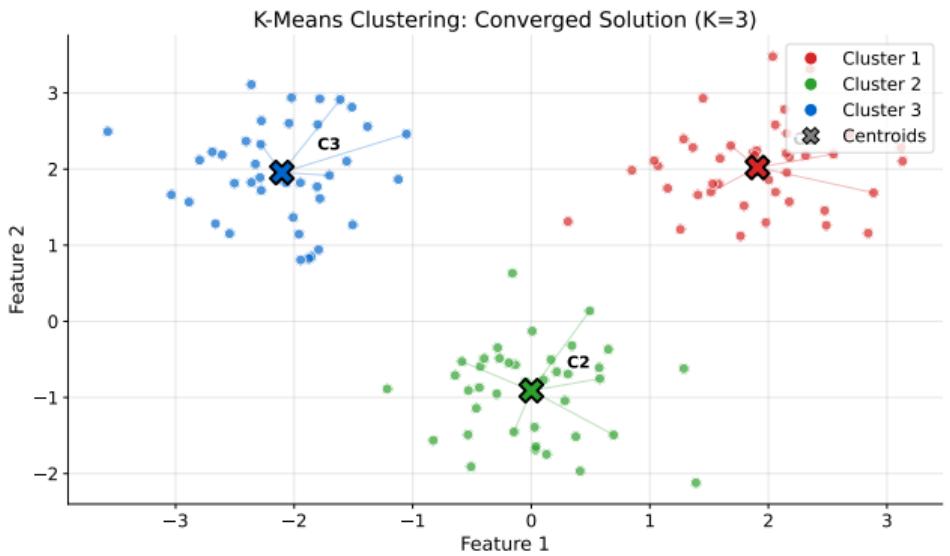
K-Means finds locally optimal solution (not guaranteed global)

- 1: **Input:** Data X , number of clusters K
- 2: Initialize K centroids randomly
- 3: **repeat**
- 4: **Assignment:** assign each point to nearest centroid
- 5: **Update:** recompute centroids as cluster means
- 6: **until** centroids don't change (or max iterations)
- 7: **return** cluster assignments, centroids

Convergence

- Guaranteed to converge (WCSS decreases each iteration)
- May converge to local optimum

Each iteration: $O(nKd)$ where $n = \text{samples}$, $K = \text{clusters}$, $d = \text{features}$



https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L03_KNN_KMeans/03_kmeans_iteration

Final state after convergence: points colored by cluster, X marks centroids

Random Initialization

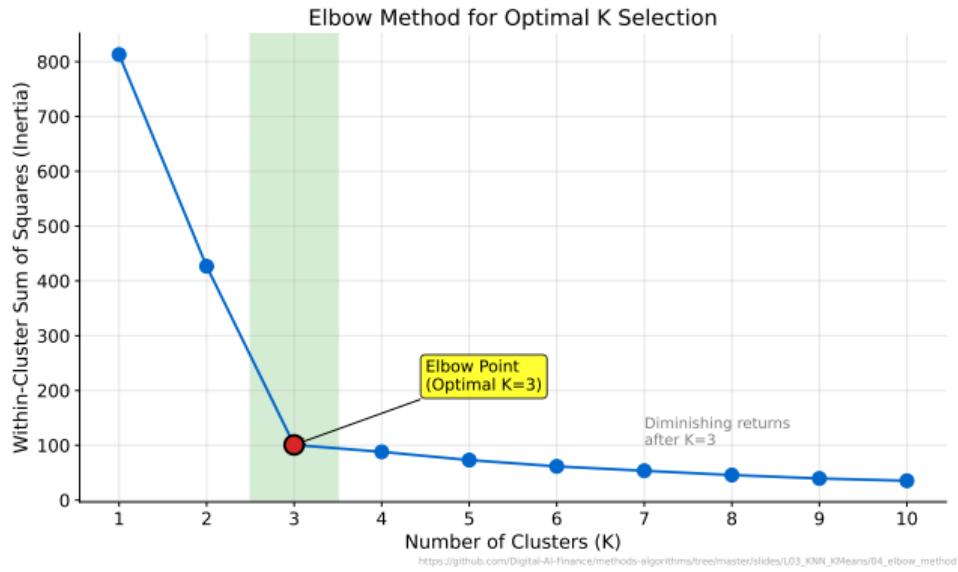
- Pick K random points as initial centroids
- Sensitive to choice, may get poor solution
- Run multiple times, keep best result

K-Means++ (Default in scikit-learn)

- Smart initialization: spread out initial centroids
- First centroid: random
- Next centroids: probability proportional to squared distance

Result: Much better starting point, fewer iterations

K-Means++ gives provably better initialization with theoretical guarantees



Interpretation: Look for the “elbow” where WCSS stops dropping sharply

For each point i

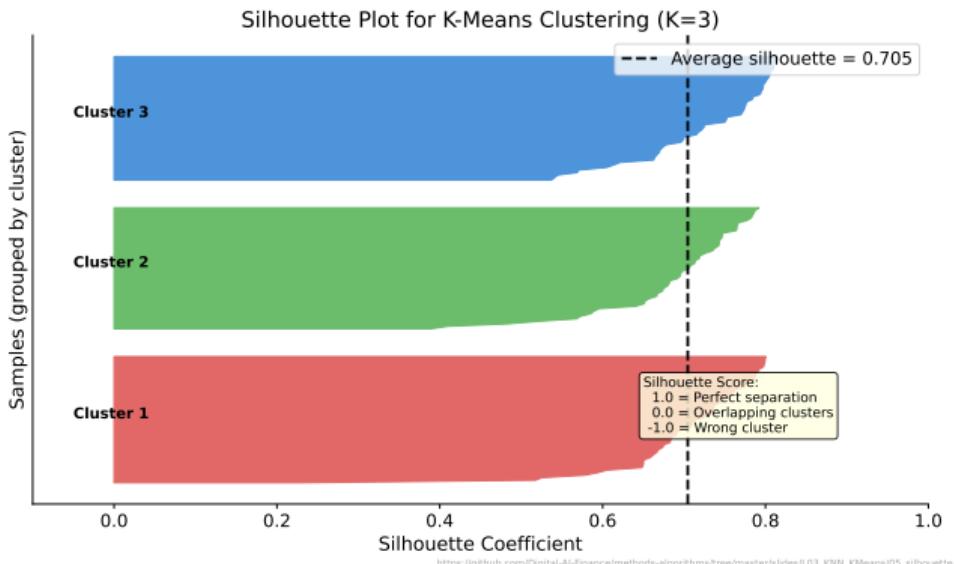
- $a(i)$: average distance to points in same cluster
- $b(i)$: average distance to points in nearest other cluster

Silhouette score

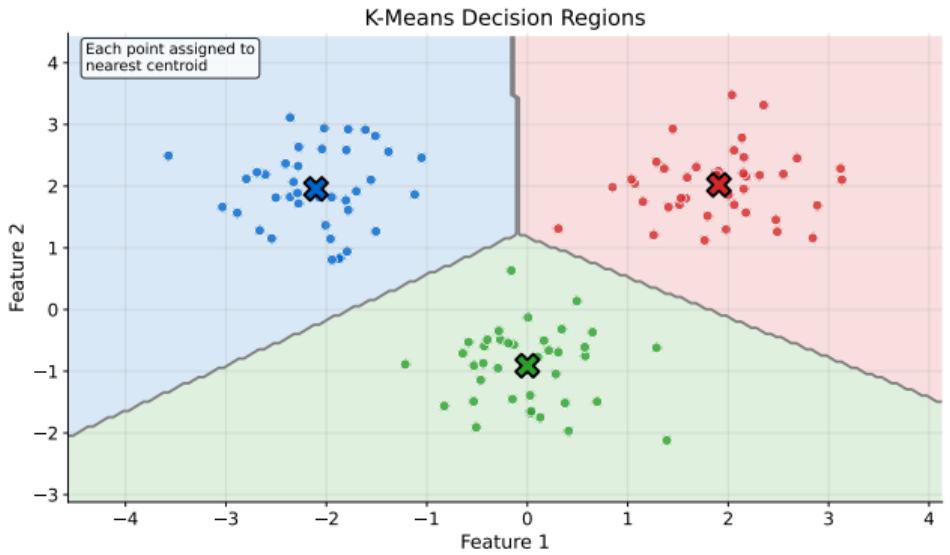
$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

- Range: $[-1, 1]$
- $s \approx 1$: point is well-matched to cluster
- $s \approx 0$: point is on boundary
- $s < 0$: point may be in wrong cluster

Average silhouette score summarizes overall clustering quality



All clusters should have similar width and scores above average line



K-Means creates Voronoi tessellation around centroids

What K-Means Assumes

- Clusters are spherical (isotropic)
- Clusters have similar sizes
- Clusters have similar densities

When K-Means Fails

- Non-convex shapes (e.g., crescents)
- Very different cluster sizes
- Different densities
- Outliers (pull centroids away)

Consider DBSCAN or Gaussian Mixture Models for these cases

Mini-Batch K-Means

- Uses random subsets for updates
- Much faster for large datasets
- Slightly worse results

K-Medoids

- Centroids must be actual data points
- More robust to outliers
- Slower than K-Means

K-Means for Mixed Data

- K-Modes: for categorical data
- K-Prototypes: mixed continuous and categorical

Mini-Batch K-Means: good for >10k samples

Basic Usage

- from sklearn.cluster import KMeans
- kmeans = KMeans(n_clusters=3, random_state=42)
- labels = kmeans.fit_predict(X)
- centroids = kmeans.cluster_centers_

Key Parameters

- n_clusters: K (required)
- init: 'k-means++' (default) or 'random'
- n_init: number of runs (default 10)
- max_iter: max iterations per run

inertia_ attribute gives WCSS after fitting

Aspect	KNN	K-Means
Task	Classification/Regression	Clustering
Learning	Supervised (needs labels)	Unsupervised
K meaning	Number of neighbors	Number of clusters
Training	None (lazy)	Iterative optimization
Prediction	Compute distances	Assign to centroid
Output	Class label	Cluster ID

The “K” in KNN and K-Means mean completely different things!

Problem: Group customers for targeted marketing

Features

- Transaction frequency
- Average transaction amount
- Account balance
- Product holdings

K-Means Solution

- Cluster into segments (e.g., K=4)
- Profile each segment
- Tailor offerings to segment needs

Example segments: High-value frequent, Dormant, New/Growing, Price-sensitive

Problem: Flag suspicious transactions

KNN Approach

- Features: amount, time, location, merchant category
- Find K similar past transactions
- If most neighbors are fraud, flag as suspicious

K-Means Approach

- Cluster “normal” transactions
- New transaction far from all centroids = anomaly
- Combined with distance threshold

KNN needs labeled fraud examples, K-Means detects deviation from normal

Use KNN When

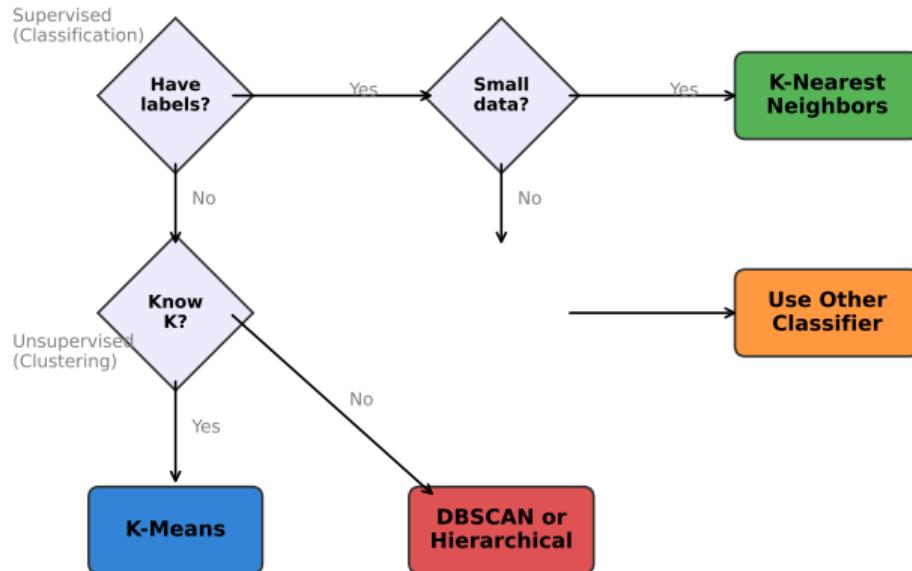
- You have labeled training data
- Local patterns matter (non-linear boundaries)
- Interpretability: "similar to these examples"
- Moderate dataset size

Use K-Means When

- No labels available
- Looking for natural groupings
- Clusters are roughly spherical
- Need fast clustering of large data

K-Means often used as preprocessing before supervised learning

KNN vs K-Means Decision Guide



https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L03_KNN_KMeans/07_decision_flowchart

Start with KNN for classification, K-Means for clustering

Instead of KNN

- Large data: use ball tree or approximate NN
- Need probability: logistic regression
- Many features: random forest

Instead of K-Means

- Unknown K: DBSCAN (density-based)
- Hierarchical structure: agglomerative clustering
- Soft assignments: Gaussian Mixture Models
- Non-spherical: spectral clustering

DBSCAN automatically determines number of clusters

K-Nearest Neighbors

- Instance-based, lazy learner
- Scale features, choose K via cross-validation
- Works best with moderate features, moderate data size

K-Means

- Iterative: assign points, update centroids
- K-Means++ for initialization, elbow/silhouette for K
- Assumes spherical, similar-size clusters

Common Considerations

- Feature scaling is critical for both
- "K" means different things in each algorithm

Both are foundational algorithms: simple, interpretable, widely used

Textbooks

- James et al. (2021). *ISLR*, Chapters 2, 12
- Hastie et al. (2009). *ESL*, Chapters 13, 14

Key Papers

- Arthur & Vassilvitskii (2007). K-Means++
- Cover & Hart (1967). Nearest Neighbor Pattern Classification

Next Lecture

- L04: Random Forests
- Ensemble methods and feature importance