

## L03: KNN & K-Means

Deep Dive: Mathematical Foundations and Implementation

### Methods and Algorithms

MSc Data Science

Spring 2026

# Outline

# The Math Behind “Similar Things Behave Similarly”

## Today's Deep Dive

In the overview, we saw that KNN classifies by neighbors and K-Means discovers clusters. Now we go deeper:

- How do we **formalize** distance and prove KNN's theoretical guarantees?
- How do we **prove** K-Means convergence and understand its connection to EM?
- How do we **validate** clusters statistically before trusting them?



XKCD #1838 by Randall Munroe (CC BY-NC 2.5) – The math behind the “pile of linear algebra”

**By the end of this lecture, you will be able to:**

1. **Analyze** the bias-variance tradeoff in KNN and derive its asymptotic error bounds
2. **Evaluate** clustering validity using statistical tests (Hopkins, Gap statistic, silhouette)
3. **Prove** K-Means convergence and analyze computational complexity
4. **Compare** distance metrics and assess their suitability for high-dimensional finance data

**Finance Applications:** Customer segmentation (RFM), fraud detection

---

Bloom's Level 4–5: Analyze, Evaluate, Prove, Compare

# Why Is KNN Called a Lazy Learner?

## Key Insight

- No explicit model training – store all training data
- Classification by majority vote of  $K$  nearest neighbors
- “Lazy” because all computation happens at prediction time

## The Algorithm in Words

1. Store all training examples  $(\mathbf{x}_i, y_i)$
2. For new query  $\mathbf{x}$ : find  $K$  nearest training points
3. Return majority class among those  $K$  neighbors

---

Instance-based learning: the training data IS the model – no parameters to estimate

# How Does the KNN Algorithm Work?

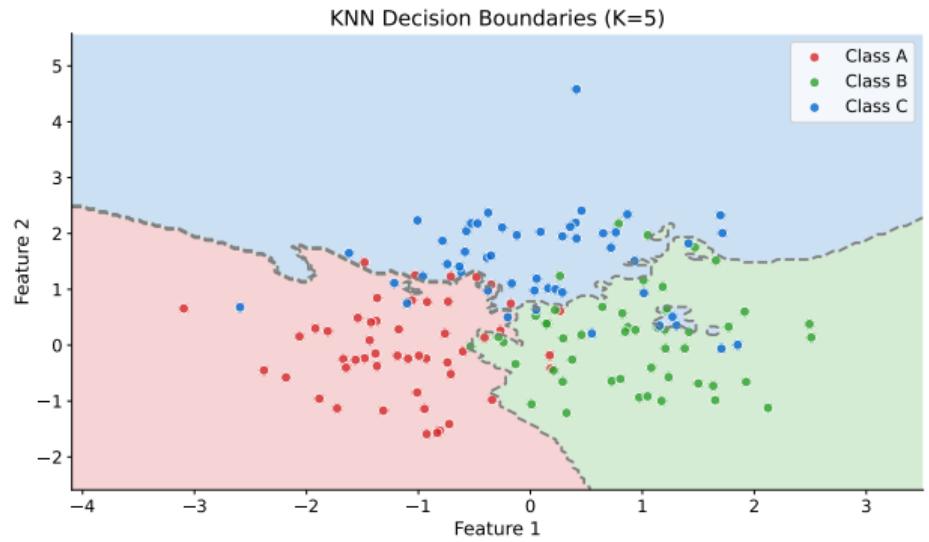
- 1: **Input:** Training set  $\mathcal{D}$ , query point  $\mathbf{x}$ , number of neighbors  $K$
- 2: Compute  $d(\mathbf{x}, \mathbf{x}_i)$  for all  $\mathbf{x}_i \in \mathcal{D}$
- 3: Select  $K$  points with smallest distances:  $N_K(\mathbf{x})$
- 4: Count votes:  $v_c = |\{i \in N_K(\mathbf{x}) : y_i = c\}|$  for each class  $c$
- 5: **if** unique maximum class exists **then**
- 6:      $\hat{y} = \arg \max_c v_c$
- 7: **else**
- 8:     **Tie-breaking:**  $\hat{y} = \arg \max_c \sum_{i:y_i=c} \frac{1}{d(\mathbf{x}, \mathbf{x}_i)}$
- 9: **end if**
- 10: **return**  $\hat{y}$

**Complexity:**  $O(nd)$  per query (brute force over  $n$  samples,  $d$  features)

---

Tie-breaking via distance-weighted vote prevents non-determinism on decision boundaries

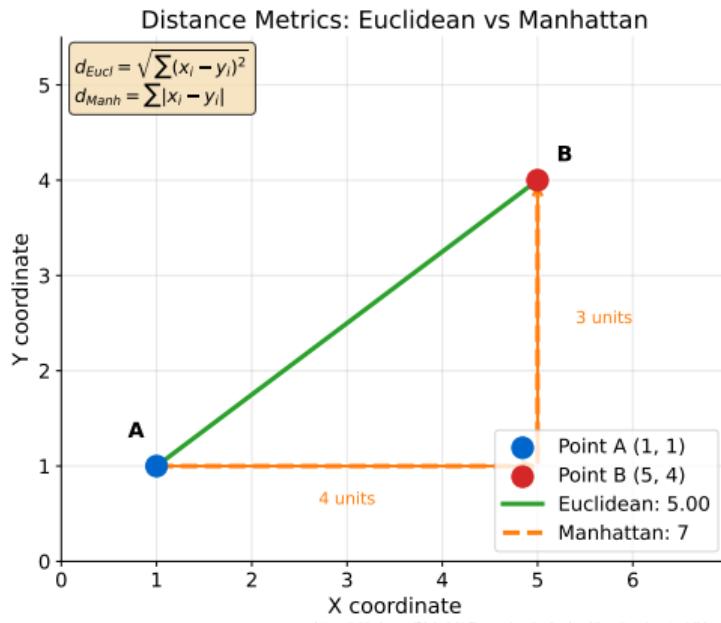
# How Do KNN Boundaries Change with K?



- Small K creates jagged, complex boundaries that follow noise
- Large K creates smooth boundaries that may miss local patterns
- The optimal K balances flexibility with stability

Decision boundaries are non-parametric – they adapt to the local data structure

# What Do Different Distance Metrics Look Like?



**Key Insight:** Different metrics define different “neighborhoods” – the choice of metric shapes the decision boundary and determines which points count as nearest neighbors.

Euclidean (circular), Manhattan (diamond), Chebyshev (square) unit balls in 2D

# How Do We Generalize Distance?

## Minkowski Distance

$$d_p(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^d |x_i - y_i|^p \right)^{1/p}$$

- $p = 1$ : Manhattan (L1) – robust to outliers, sparse feature differences
- $p = 2$ : Euclidean (L2) – default, geometrically intuitive
- $p = \infty$ : Chebyshev – maximum absolute difference along any axis

**Triangle Inequality:** For  $p \geq 1$ ,  $d_p(\mathbf{x}, \mathbf{z}) \leq d_p(\mathbf{x}, \mathbf{y}) + d_p(\mathbf{y}, \mathbf{z})$

- Enables pruning in KD-Trees and Ball Trees
- For  $p < 1$ : violates triangle inequality (not a true metric)

---

Higher  $p$  amplifies large single-feature differences;  $p = 2$  is the default for most applications

# When Do We Need Non-Euclidean Distance?

## Cosine Similarity (for text and embeddings)

$$\cos(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}, \quad \text{range: } [-1, 1]$$

- Measures angle, not magnitude – ideal for sparse, high-dimensional data
- Use for: document similarity, word embeddings, TF-IDF vectors

## Mahalanobis Distance (accounts for correlation)

$$d_M(\mathbf{x}, \mathbf{y}) = \sqrt{(\mathbf{x} - \mathbf{y})^T \Sigma^{-1} (\mathbf{x} - \mathbf{y})}$$

- Uses feature covariance  $\Sigma$  – unit-less, scale-invariant
- Detects outliers accounting for correlation structure

---

Choose metric based on data type: Euclidean for dense, Cosine for sparse/text, Mahalanobis for correlated features

# How Does K Control the Bias-Variance Tradeoff?

## The Fundamental Tradeoff

- $K = 1$ : High variance, low bias – very flexible, memorizes noise
- $K = n$ : High bias, low variance – always predicts majority class

## Practical Guidelines

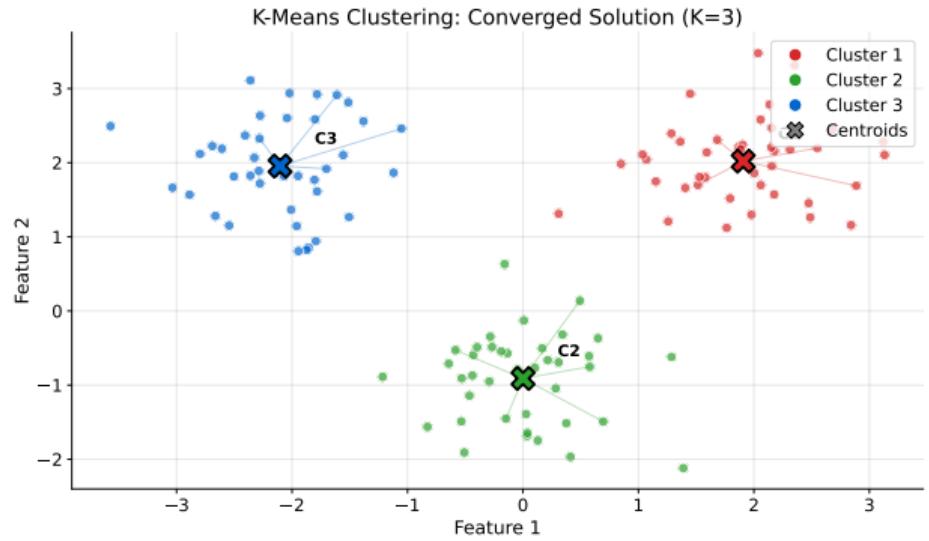
- Start with  $K = \sqrt{n}$  where  $n$  is training size
- Use odd  $K$  for binary classification (avoids ties)
- Always validate with cross-validation

**Common Choices:**  $K \in \{3, 5, 7, 11\}$

---

Small K for complex patterns with clean data; larger K for noisy data requiring smoothing

# How Do Centroids Move During Iteration?



- Each iteration: assign points to nearest centroid, then recompute centroids
- Centroids migrate toward cluster centers until stable
- Convergence typically occurs within 10–50 iterations

Visualizing iterations helps verify that K-Means is finding meaningful clusters

# Why Weight Neighbors by Distance?

**Problem:** All  $K$  neighbors have equal influence – a distant neighbor counts as much as the closest one.

## Inverse-Distance Weighting

$$\hat{y} = \arg \max_c \sum_{i \in N_K(\mathbf{x})} w_i \cdot \mathbf{1}[y_i = c] \quad \text{where } w_i = \frac{1}{d(\mathbf{x}, \mathbf{x}_i)}$$

- Closer neighbors receive exponentially higher weight
- Reduces sensitivity to the choice of  $K$
- For  $d = 0$  (exact match): return that point's class directly

In scikit-learn: `weights='uniform'` (default) vs `weights='distance'`

---

Distance weighting often improves performance, especially when  $K$  is not perfectly tuned

# How Do We Select K Objectively?

## GridSearchCV for Optimal K

- `param_grid = {'n_neighbors': range(1, 21, 2)}`
- `grid = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5)`
- `grid.fit(X_train, y_train)`
- `best_k = grid.best_params_['n_neighbors']`

## Validation Curve Interpretation

- Plot accuracy vs K for both training and validation sets
- Choose K where validation accuracy peaks (before overfitting gap widens)
- Large gap between train/val curves signals overfitting

---

Cross-validation provides objective, data-driven K selection – never choose K by "eyeballing"

# Why Must We Scale Features for KNN?

## Why Scaling Matters – Without scaling:

- Income: ranges 20,000–200,000
- Age: ranges 20–80
- Distance dominated by income ( $180,000 \times$  larger scale!)

## Scaling Methods

- **Standardization:**  $z = \frac{x - \mu}{\sigma}$  (mean=0, std=1)
- **Min-Max:**  $x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$  (range [0,1])

## Edge Cases

- If  $\sigma = 0$  or  $x_{\max} = x_{\min}$  (constant feature): drop feature or set to 0

---

Rule: **ALWAYS** scale features for distance-based methods. StandardScaler for Gaussian-like, MinMaxScaler for bounded.

# Why Does KNN Struggle in High Dimensions?

## The Problem (Beyer et al., 1999)

In high dimensions, for any query point  $\mathbf{x}$ :

$$\lim_{d \rightarrow \infty} \frac{d_{\max} - d_{\min}}{d_{\min}} \rightarrow 0$$

- All points become approximately equidistant
- “Nearest neighbor” becomes meaningless
- Volume of unit hypersphere  $\rightarrow 0$  as  $d \rightarrow \infty$

## Solutions

- **Dimensionality reduction:** PCA before KNN (see L05)
- **Feature selection:** keep only relevant features
- **Domain knowledge:** select meaningful features from subject matter

---

KNN works best with moderate dimensionality ( $d < 15-20$ , problem-dependent)

# How Good Can KNN Theoretically Be?

**Theorem:** As  $n \rightarrow \infty$ , the 1-NN error rate satisfies:

$$R^* \leq R_{\text{1-NN}} \leq 2R^*(1 - R^*)$$

where  $R^*$  is the Bayes optimal error rate.

## Key Results

- 1-NN error is at most  $2 \times$  Bayes error – remarkably strong for such a simple method
- KNN is **universally consistent**:  $R_{K\text{-NN}} \rightarrow R^*$  exactly
- Requires:  $K \rightarrow \infty$  and  $K/n \rightarrow 0$  simultaneously

**Practical Implication:** KNN is a strong baseline – if it performs badly, the features may be poor rather than the algorithm.

---

Universal consistency holds under mild conditions; see appendix for full proof

## How Do We Prove the 1-NN Error Bound?

**Proof idea** (1-NN case, binary classification):

1. As  $n \rightarrow \infty$ , the nearest neighbor  $\mathbf{x}_{(1)} \rightarrow \mathbf{x}$  (converges to query)
2. Let  $\eta(\mathbf{x}) = P(Y = 1|\mathbf{x})$ . The 1-NN error at  $\mathbf{x}$ :

$$R_{\text{1-NN}}(\mathbf{x}) = \eta(\mathbf{x})(1 - \eta(\mathbf{x})) + (1 - \eta(\mathbf{x}))\eta(\mathbf{x}) = 2\eta(\mathbf{x})(1 - \eta(\mathbf{x}))$$

Since  $R^*(\mathbf{x}) = \min(\eta, 1 - \eta)$  and  $2\eta(1 - \eta) \leq 2 \min(\eta, 1 - \eta)$ :

$$R^* \leq R_{\text{1-NN}} \leq 2R^*(1 - R^*)$$

For  $K$ -NN with  $K \rightarrow \infty$ ,  $K/n \rightarrow 0$ : majority vote over  $K$  neighbors converges to  $\mathbf{1}[\eta(\mathbf{x}) > 0.5]$ , so  $R_{K\text{-NN}} \rightarrow R^*$  exactly.

---

The  $2 \times$  bound is tight: equality when  $\eta(\mathbf{x}) = 0.5$  everywhere (maximum ambiguity)

## How Does K Affect Bias and Variance?

For KNN regression,  $\hat{f}_K(\mathbf{x}) = \frac{1}{K} \sum_{i \in N_K(\mathbf{x})} y_i$ , the expected error decomposes as:

$$E[(\hat{f}_K(\mathbf{x}) - Y)^2] = \underbrace{\text{Bias}^2(\hat{f}_K(\mathbf{x}))}_{\text{increases with } K} + \underbrace{\text{Var}(\hat{f}_K(\mathbf{x}))}_{\text{decreases with } K} + \sigma^2$$

- **Variance:**  $\text{Var}(\hat{f}_K) = \frac{\sigma^2}{K}$  (averaging  $K$  neighbors reduces variance)
- **Bias:**  $\text{Bias}(\hat{f}_K) \approx f(\mathbf{x}) - \frac{1}{K} \sum_{i \in N_K} f(\mathbf{x}_i)$  (larger  $K$  = farther neighbors = more bias)
- **Optimal  $K$ :** balances this tradeoff; found via cross-validation

---

$K = 1$ : zero bias, variance =  $\sigma^2$ .     $K = n$ : high bias (global mean), variance =  $\sigma^2/n$ .

## What Is K-Means Trying to Optimize?

**Goal:** Partition  $n$  points into  $K$  clusters by minimizing within-cluster sum of squares (WCSS):

$$J = \sum_{k=1}^K \sum_{\mathbf{x}_i \in C_k} \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2$$

where  $\boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{\mathbf{x}_i \in C_k} \mathbf{x}_i$  is the centroid of cluster  $C_k$ .

### Key Properties

- Each point assigned to its nearest centroid
- Iterative refinement: assign → update → repeat
- Converges to local (not global) optimum

---

Finding the globally optimal K-Means solution is NP-hard (Aloise et al., 2009)

# How Does Lloyd's Algorithm Work?

- 1: **Input:** Data  $\mathbf{X}$ , number of clusters  $K$
- 2: Initialize  $K$  centroids (K-Means++ or random)
- 3: **repeat**
- 4:   **Assignment:**  $C_k = \{\mathbf{x}_i : \|\mathbf{x}_i - \boldsymbol{\mu}_k\| \leq \|\mathbf{x}_i - \boldsymbol{\mu}_j\| \forall j\}$
- 5:   **Update:**  $\boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{\mathbf{x}_i \in C_k} \mathbf{x}_i$
- 6:   **Handle empty:** if  $|C_k| = 0$ , reinitialize  $\boldsymbol{\mu}_k$  from farthest point
- 7: **until** centroid change  $< \epsilon$  or max iterations reached
- 8: **return** cluster assignments, centroids

**Convergence:** Guaranteed – WCSS decreases monotonically each iteration. May converge to local optimum; run multiple restarts.

---

Each iteration costs  $O(nKd)$  where  $n$  = samples,  $K$  = clusters,  $d$  = features

# Why Does Smart Initialization Matter?

**Problem:** Random initialization is sensitive to starting positions and often converges to poor local optima.

**K-Means++ Algorithm** (Arthur & Vassilvitskii, 2007)

1. Choose first centroid uniformly at random from data
2. For each subsequent centroid: select point  $x$  with probability  $\propto d(x)^2$  (squared distance to nearest existing centroid)
3. Repeat until  $K$  centroids chosen

**Guarantee:** Expected cost  $\leq 8(\ln K + 2) \times$  optimal cost – an  $O(\log K)$ -competitive approximation.

---

**K-Means++ is the default in scikit-learn (`init='k-means++'`); typically converges in fewer iterations**

# How Do We Measure Cluster Quality?

For each point  $i$ :

- $a(i)$ : average distance to all other points in **same** cluster (cohesion)
- $b(i)$ : average distance to points in **nearest other** cluster (separation)

Silhouette coefficient:

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}, \quad s(i) \in [-1, 1]$$

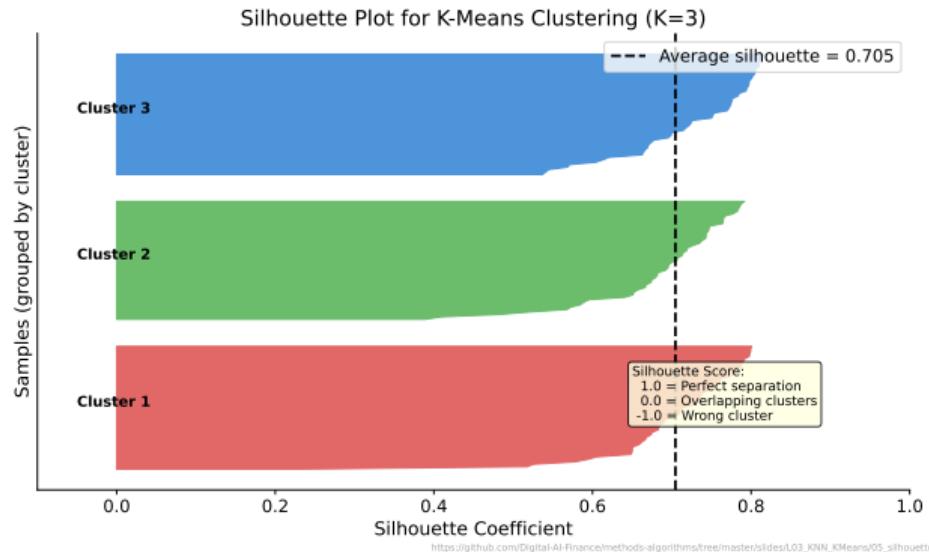
Interpretation:

- $s \approx 1$ : point is well-clustered (far from neighboring clusters)
- $s \approx 0$ : point lies on the boundary between clusters
- $s < 0$ : point is likely assigned to the wrong cluster

---

Average silhouette score across all points summarizes overall clustering quality. Singleton:  $s(i) = 0$  by convention.

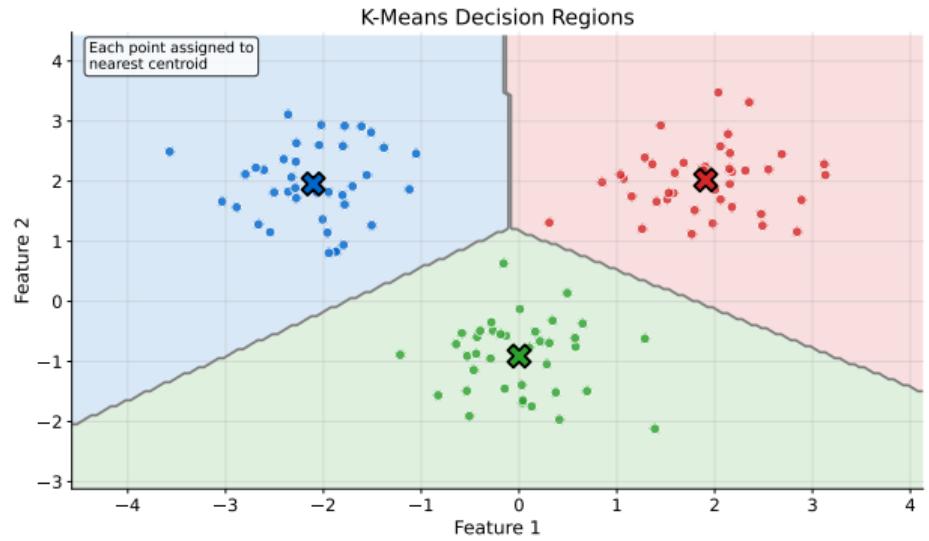
# What Does a Silhouette Plot Reveal?



**Reading the Plot:** Each cluster should have similar width (balanced sizes) and scores above the average line. Thin or negative clusters indicate poor assignments.

Compare silhouette plots across different K values to find the optimal number of clusters

# What Shape Are K-Means Clusters?



**Voronoi Tessellation:** K-Means partitions the feature space into convex polygonal regions. Every point within a region is closest to that region's centroid.

---

Voronoi cells are always convex – this is why K-Means cannot discover non-convex cluster shapes

# When Does K-Means Fail?

## What K-Means Assumes

- Clusters are **spherical** (isotropic variance in all directions)
- Clusters have **similar sizes** (roughly equal number of points)
- Clusters have **similar densities**

## When K-Means Fails

- Non-convex shapes (e.g., crescents, rings) – use DBSCAN
- Very different cluster sizes or densities – use GMM
- Strong outliers – use K-Medoids (centroids must be data points)

---

Always visualize clusters in 2D (via PCA/t-SNE) to check if assumptions hold

## Why Must K-Means Converge?

**Theorem:** Lloyd's algorithm converges in a finite number of iterations.

**Proof** (coordinate descent argument):

The WCSS objective  $J = \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$  satisfies:

1. **Assignment step** (fix  $\mu$ , optimize  $C$ ): Reassigning each point to its nearest centroid can only decrease  $J \Rightarrow J$  non-increasing
2. **Update step** (fix  $C$ , optimize  $\mu$ ): The mean minimizes within-cluster SSE  $\Rightarrow J$  non-increasing

Since  $J \geq 0$  and non-increasing, and there are finitely many partitions of  $n$  points into  $K$  clusters: the algorithm must terminate. □

**NP-hardness:** Global optimum is NP-hard; K-Means++ provides  $O(\log K)$  approximation.

---

Block coordinate descent: alternating optimization of two blocks ( $C$  and  $\mu$ ) on a bounded objective

# How Does K-Means Relate to EM?

## Connection to Expectation-Maximization

- K-Means = “Hard EM” for Gaussian Mixture Models (GMM)
- Assumes: spherical Gaussians with equal variance  $\sigma^2 I$
- **E-step:** assign points to nearest centroid (hard assignment,  $\gamma_{ik} \in \{0, 1\}$ )
- **M-step:** update centroids as cluster means

## Why This Matters

- Explains WHY convergence is guaranteed (EM always converges)
- Explains WHY K-Means assumes spherical clusters
- Opens door to **soft clustering** via full GMM (probabilistic assignments)

---

Soft EM:  $\gamma_{ik} \in [0, 1]$  gives probability of membership – more flexible but computationally heavier

# What Alternatives Exist to Standard K-Means?

## Mini-Batch K-Means

- Uses random subsets (mini-batches) for centroid updates
- 10–100× faster for large datasets; slightly worse results

## K-Medoids (PAM)

- Centroids must be actual data points (medoids)
- More robust to outliers; works with any distance metric

## K-Modes / K-Prototypes

- K-Modes: for categorical data (uses mode instead of mean)
- K-Prototypes: mixed continuous and categorical features

---

Mini-Batch K-Means is recommended for  $n > 10,000$  samples; K-Medoids for non-Euclidean distances

## How Fast Are These Algorithms?

Method	Training	Prediction
KNN (brute force)	$O(1)$	$O(nd)$
KNN (KD-Tree)	$O(nd \log n)$	$O(d \log n)$
KNN (Ball Tree)	$O(nd \log n)$	$O(d \log n)$
K-Means (Lloyd's)	$O(nKdT)$	$O(Kd)$
Mini-Batch K-Means	$O(bKdT)$	$O(Kd)$

$n$  = samples,  $d$  = features,  $K$  = neighbors/clusters,  $T$  = iterations,  $b$  = batch size.

- KD-Tree degrades to  $O(nd)$  when  $d > 15\text{--}20$
- For very large  $n$ : approximate nearest neighbors (LSH, FAISS)

---

scikit-learn `algorithm='auto'` selects the best method based on  $n$  and  $d$

# Should We Cluster This Data at All?

**Before Clustering:** Test whether data has cluster tendency at all.

$$H = \frac{\sum_{i=1}^m u_i^d}{\sum_{i=1}^m u_i^d + \sum_{i=1}^m w_i^d}$$

- $u_i$  = distance from random point (uniform in data range) to nearest data point
- $w_i$  = distance from randomly sampled data point to its nearest neighbor

**Interpretation:**

- $H \approx 0.5$ : uniform distribution (no clusters) – do not cluster
- $H > 0.75$ : significant clustering tendency – proceed with clustering

---

Always run Hopkins test before K-Means – clustering uniform data produces meaningless results

## How Does the Gap Statistic Choose K?

**Gap Statistic** (Tibshirani et al., 2001)

$$\text{Gap}_n(k) = E_n^*[\log W_k] - \log W_k$$

where  $W_k = \sum_{r=1}^k \frac{1}{2|C_r|} \sum_{i,j \in C_r} \|x_i - x_j\|^2$  and  $E_n^*$  is the expectation under a reference (uniform) distribution.

**Selection Rule:** Choose smallest  $K$  such that:

$$\text{Gap}(k) \geq \text{Gap}(k+1) - s_{k+1}$$

where  $s_{k+1}$  is the standard error from  $B$  bootstrap reference samples.

**Advantage over Elbow:** Provides a statistical criterion rather than subjective visual inspection.

---

Tibshirani et al. recommend  $B = 20\text{--}50$  bootstrap samples for stable estimates

# Which K Selection Method Should We Use?

Method	Objectivity	Speed	Best For
Elbow (WCSS)	Low (subjective)	Fast	Quick exploration
Silhouette	High	Moderate	Cluster quality
Gap Statistic	High (statistical)	Slow	Rigorous analysis

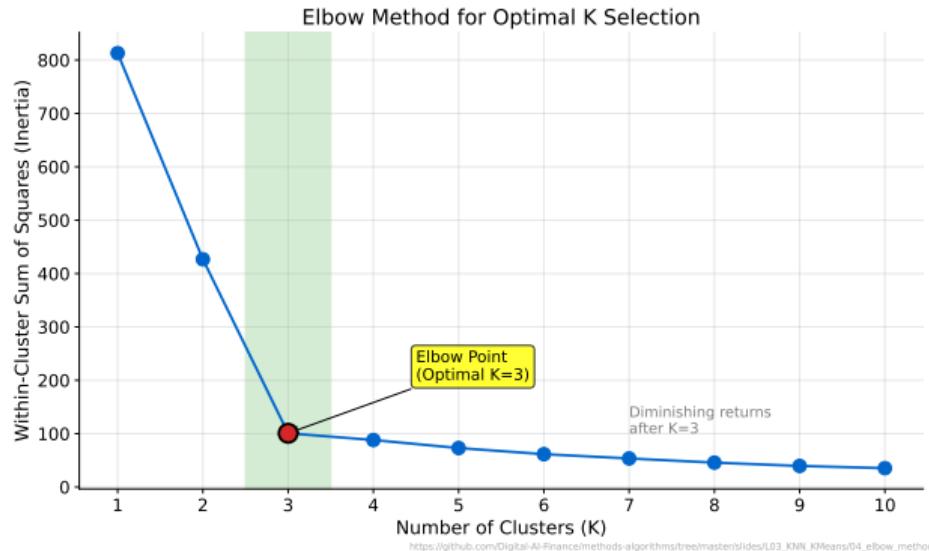
## Practical Recommendation:

- Start with elbow plot for quick overview
- Validate with silhouette score and silhouette plots
- Use Gap statistic when formal justification is needed (e.g., publications)

---

No single method is perfect – use multiple methods and check if they agree

# Where Is the Elbow in the WCSS Curve?



- WCSS always decreases with more clusters – the question is *how fast*
- The “elbow” marks the point of diminishing returns
- Combine with silhouette analysis for a more objective decision

The elbow is often ambiguous – that is why we also use silhouette and Gap statistics

## How Do KNN and K-Means Differ Fundamentally?

Aspect	KNN	K-Means
Task	Classification / Regression	Clustering
Learning Type	Supervised (needs labels)	Unsupervised (no labels)
K meaning	Number of neighbors	Number of clusters
Training	None (lazy learner)	Iterative optimization
Prediction	Compute distances to all	Assign to nearest centroid
Output	Class label / value	Cluster assignment
Scalability	$O(nd)$ per query	$O(nKdT)$ total

---

The “K” in KNN and K-Means mean completely different things – a common source of confusion!

# How Do Banks Segment Customers with RFM?

## RFM Analysis (Industry Standard)

- Recency: Days since last transaction
- Frequency: Number of transactions in period
- Monetary: Total / average transaction value

**Pipeline:** Standardize RFM → K-Means ( $K=4-6$ ) → Profile segments **Typical Segments:**

- “Champions” (high R, F, M) → loyalty rewards
- “At Risk” (low R, declining F) → retention campaigns
- “New Customers” (high R, low F) → onboarding programs

**Business Value:** Customer Lifetime Value (CLV) prediction per segment

---

RFM segmentation is foundational for CRM, marketing analytics, and credit risk profiling

# How Does KNN Detect Financial Fraud?

## CRITICAL: Class Imbalance Problem

- Fraud is typically <1% of transactions (100:1 ratio)
- Naive KNN majority vote: ALWAYS predicts non-fraud (99% accuracy, zero fraud detected!)

## Solutions for Imbalanced Data

- **SMOTE**: Synthetic Minority Oversampling TEchnique
- **Weighted KNN**: Higher weight for minority class neighbors
- **Anomaly score**: Use distance to  $K$ -th neighbor instead of majority vote

**Evaluation:** Use **Precision-Recall AUC**, NOT accuracy!

---

Always check class distribution before applying majority vote – accuracy is misleading for imbalanced data

# What If Clusters Aren't Spherical?

## DBSCAN (Density-Based Spatial Clustering):

- **Core point:**  $\geq \text{minPts}$  neighbors within radius  $\varepsilon$
- **Border point:** within  $\varepsilon$  of a core point but not itself core
- **Noise:** neither core nor border  $\Rightarrow$  automatic outlier detection
- Does not require  $K$ ; discovers non-spherical clusters

## Hierarchical (Agglomerative) Clustering:

- Bottom-up: each point starts as its own cluster, iteratively merge closest pairs
- Linkage: single (min), complete (max), Ward (minimize variance)
- **Dendrogram:** cut at desired height to get  $K$  clusters

---

DBSCAN for irregular shapes with outliers; hierarchical when exploring multiple K values via dendrogram

# Which Algorithm Should You Choose?

## Use KNN When

- You have labeled training data (supervised)
- Local patterns matter (non-linear boundaries)
- Interpretability: “classified because similar to these examples”

## Use K-Means When

- No labels available (unsupervised)
- Looking for natural groupings in data
- Clusters are roughly spherical and similar in size

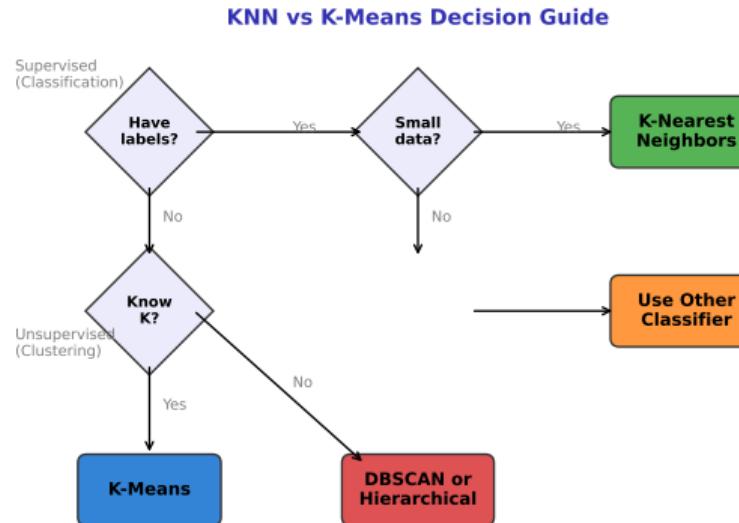
## Use Alternatives When

- Non-spherical clusters → DBSCAN or spectral clustering
- Soft assignments needed → Gaussian Mixture Models
- Very large  $n$  → Mini-Batch K-Means or approximate NN

---

K-Means is often used as preprocessing (cluster features) before supervised learning

# Decision Flowchart: KNN vs K-Means vs Alternatives



- Labels available → supervised methods (KNN, logistic regression, random forests)
- No labels, spherical clusters → K-Means with K-Means++ initialization
- Non-spherical or unknown K → DBSCAN or hierarchical clustering

Start simple (K-Means) and add complexity only when validation metrics demand it

# How Do We Implement KNN in Python?

## Classification

- `from sklearn.neighbors import KNeighborsClassifier`
- `knn = KNeighborsClassifier(n_neighbors=5, weights='distance')`
- `knn.fit(X_train, y_train)`
- `y_pred = knn.predict(X_test)`

## Key Parameters

- `n_neighbors`: K value (default 5)
- `weights`: 'uniform' or 'distance'
- `metric`: 'euclidean', 'manhattan', 'minkowski'
- `algorithm`: 'auto', 'ball\_tree', 'kd\_tree', 'brute'

---

Also available: `KNeighborsRegressor` for regression; `RadiusNeighborsClassifier` for radius-based

# How Do We Implement K-Means in Python?

## Basic Usage

- `from sklearn.cluster import KMeans`
- `km = KMeans(n_clusters=3, init='k-means++', random_state=42)`
- `labels = km.fit_predict(X)`
- `centroids = km.cluster_centers_`

## Key Parameters

- `n_clusters`:  $K$  (required, no default)
- `init`: '`k-means++`' (default) or '`random`'
- `n_init`: number of restarts (default 10; keeps best)
- `max_iter`: max iterations per run (default 300)

---

`km.inertia_` gives WCSS after fitting; `km.n_iter_` gives iterations until convergence

# Why Must Scaling Be Inside the Pipeline?

**Why Pipelines?** Prevent data leakage – scaling must be fit on training data only.

- from sklearn.pipeline import Pipeline
- from sklearn.preprocessing import StandardScaler
- pipe = Pipeline([('scaler', StandardScaler()),  
                  ('knn', KNeighborsClassifier())])
- pipe.fit(X\_train, y\_train)

## GridSearchCV with Pipeline

- params = {'knn\_\_n\_neighbors': [3,5,7,11],  
              'knn\_\_weights': ['uniform','distance']}
- grid = GridSearchCV(pipe, params, cv=5, scoring='f1')

---

Pipelines ensure scaling is part of cross-validation – fitting scaler on full data before CV causes leakage

## Open the Colab Notebook

- **Exercise 1:** Implement weighted KNN with distance weighting; compare uniform vs distance on fraud data
- **Exercise 2:** Compare Gap statistic vs Elbow vs Silhouette for K selection on customer data
- **Exercise 3:** Apply SMOTE + KNN pipeline for imbalanced fraud detection; evaluate with PR AUC

**Link:** See course materials for Colab notebook

---

All exercises use real-world-inspired financial datasets with class imbalance and mixed scales

# What Should You Remember?

## K-Nearest Neighbors

- Instance-based lazy learner; Cover & Hart bound:  $R_{1\text{-NN}} \leq 2R^*$
- Scale features, choose K via cross-validation, consider weighted voting

## K-Means

- Iterative assign-update with guaranteed convergence (local optimum)
- K-Means++ for initialization; silhouette/Gap for K selection

## Common Considerations

- Feature scaling is critical for both methods
- “K” means completely different things in each algorithm

**Finance:** RFM segmentation (K-Means), fraud detection with SMOTE (KNN)

---

Both are foundational algorithms: simple, interpretable, widely used across industries

*"Even K-Means would struggle to cluster the ways students misuse K-Means."*

With KNN and K-Means, you can now classify the known and discover the unknown.

**Next Session:** L04 – Random Forests (from distance-based to tree-based methods)

---

XKCD #2731 callback – clustering is easy, knowing when to cluster is the hard part

## Appendix: Advanced Topics and Proofs

## Distance Metric Properties: Triangle Inequality

**Statement:** For any metric  $d$  and points  $x, y, z$ :

$$d(x, z) \leq d(x, y) + d(y, z)$$

**Proof for Euclidean (via Cauchy-Schwarz):**

1.  $\|x - z\| = \|(x - y) + (y - z)\|$
2.  $\leq \|x - y\| + \|y - z\|$  (by Minkowski inequality, itself a consequence of Cauchy-Schwarz)

**Why It Enables KD-Tree Pruning:**

- If current best distance is  $d^*$  and node center is  $c$ :
- Prune subtree if  $d(x, c) - r > d^*$  (where  $r$  = node radius)
- Triangle inequality guarantees no closer point can exist in that subtree

---

Without triangle inequality ( $p < 1$ ), spatial indexing structures cannot prune and brute force is required

## K-Means Convergence: Formal Proof

**Objective:**  $J(C, \mu) = \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$

**Block Coordinate Descent:**

1. **Fix  $\mu$ , optimize  $C$ :** For each  $x_i$ ,  $\frac{\partial J}{\partial C}$  is minimized by  $C_k^* = \{x_i : k = \arg \min_j \|x_i - \mu_j\|\} \Rightarrow J^{(t+1)} \leq J^{(t)}$
2. **Fix  $C$ , optimize  $\mu$ :**  $\frac{\partial J}{\partial \mu_k} = -2 \sum_{x_i \in C_k} (x_i - \mu_k) = 0$  gives  $\mu_k^* = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i \Rightarrow J^{(t+1)} \leq J^{(t)}$

**Termination:**  $J \geq 0$  and strictly non-increasing. The number of distinct partitions of  $n$  points into  $K$  groups is  $\leq K^n$  (finite). Each partition visited at most once  $\Rightarrow$  convergence in  $\leq K^n$  steps (in practice, much fewer).

---

Worst-case  $K^n$  iterations is exponential but never observed in practice; typical convergence in 10–50 iterations

# Empty Cluster Handling Strategies

**Problem:** During K-Means iteration, a cluster may lose all its points.

## Strategy 1: Farthest Point Reinitialization

- Find the point farthest from its assigned centroid
- Use it as the new centroid for the empty cluster
- Rationale: poorly fit points are good candidates for new clusters

## Strategy 2: Split Largest Cluster

- Find cluster with highest WCSS, split it into two
- Maintains total number of clusters at  $K$

## Strategy 3: Random Reinitialization

- Randomly select a new data point as centroid
- Simplest but least principled approach

---

scikit-learn uses a reassignment strategy; K-Means++ initialization makes empty clusters rare

**Setup:** Binary classification,  $\eta(\mathbf{x}) = P(Y = 1|\mathbf{x})$ , Bayes rule:  $g^*(\mathbf{x}) = \mathbf{1}[\eta(\mathbf{x}) > 0.5]$ .

## 1-NN Error Rate Derivation:

- As  $n \rightarrow \infty$ , nearest neighbor  $\mathbf{x}_{(1)} \rightarrow \mathbf{x}$ , so  $\eta(\mathbf{x}_{(1)}) \rightarrow \eta(\mathbf{x})$
- $P(\text{error}|\mathbf{x}) = P(Y \neq Y_{(1)}|\mathbf{x})$
- $= \eta(\mathbf{x})(1 - \eta(\mathbf{x})) + (1 - \eta(\mathbf{x}))\eta(\mathbf{x}) = 2\eta(\mathbf{x})(1 - \eta(\mathbf{x}))$

## Bounding:

$$R^*(\mathbf{x}) = \min(\eta, 1 - \eta) \leq 2\eta(1 - \eta) \leq 2 \min(\eta, 1 - \eta) = 2R^*(\mathbf{x})$$

Integrating over  $\mathbf{x}$ :  $R^* \leq R_{\text{1-NN}} \leq 2R^*(1 - R^*)$

**K-NN Extension:** With  $K \rightarrow \infty$ ,  $K/n \rightarrow 0$ , majority vote converges to  $\mathbf{1}[\eta > 0.5]$  by law of large numbers  
 $\Rightarrow R_{K\text{-NN}} \rightarrow R^*$ .

---

The tighter bound  $2R^*(1 - R^*) \leq 2R^*$  shows 1-NN is at most 2× suboptimal

Algorithm	Build	Query (avg)	Query (worst)
Brute Force	$O(1)$	$O(nd)$	$O(nd)$
KD-Tree	$O(nd \log n)$	$O(d \log n)$	$O(nd)$
Ball Tree	$O(nd \log n)$	$O(d \log n)$	$O(nd)$
LSH (approx.)	$O(nd)$	$O(d)$	$O(nd)$

## When Each Breaks Down:

- KD-Tree: degrades to brute force for  $d > 15\text{--}20$  (curse of dimensionality)
- Ball Tree: better than KD-Tree in moderate-high  $d$  but still degrades
- LSH: approximate; tunable accuracy-speed tradeoff via hash functions

---

Locality-Sensitive Hashing (LSH) is used by FAISS, Annoy, and other large-scale NN libraries

## Gap Statistic: Mathematical Details

**Reference Distribution:** Generate  $B$  datasets from uniform distribution over bounding box of data (or from PCA-aligned box for better null model).

**Bootstrap Procedure:**

1. For each  $k = 1, \dots, K_{\max}$ : compute  $\log W_k$  on real data
2. For  $b = 1, \dots, B$ : generate reference data, compute  $\log W_k^{*(b)}$
3.  $\text{Gap}(k) = \frac{1}{B} \sum_{b=1}^B \log W_k^{*(b)} - \log W_k$
4.  $s_k = \text{sd}(\log W_k^{*(b)}) \cdot \sqrt{1 + 1/B}$

**Formal Selection Rule:** Choose smallest  $k$  such that:

$$\text{Gap}(k) \geq \text{Gap}(k+1) - s_{k+1}$$

**Intuition:** Gap measures how much better the clustering is compared to clustering uniform (structureless) data.

---

Tibshirani et al. (2001) recommend  $B = 20-50$ ; the  $\sqrt{1 + 1/B}$  correction accounts for simulation error

## Textbooks

- James et al. (2021). *ISLR*, Ch. 2 (KNN), Ch. 12 (Clustering)
- Hastie et al. (2009). *ESL*, Ch. 13 (Prototypes/NN), Ch. 14 (Unsupervised)

## Key Papers

- Arthur & Vassilvitskii (2007). K-Means++: The Advantages of Careful Seeding
- Cover & Hart (1967). Nearest Neighbor Pattern Classification
- Tibshirani et al. (2001). Estimating the Number of Clusters via the Gap Statistic
- Beyer et al. (1999). When Is “Nearest Neighbor” Meaningful?

**Next Lecture:** L04 – Random Forests (from distance-based to tree-based ensemble methods)

---

All papers are available via university library access or open-access preprints