

## L03: KNN & K-Means

Mathematical Foundations and Implementation

Methods and Algorithms

Spring 2026

# Outline

1 K-Nearest Neighbors

2 Practice

3 K-Means Clustering

4 Comparison and Applications

5 Summary

## Key Insight

- No explicit model training (store all data)
- Classification by majority vote of K nearest neighbors
- “Lazy” because work is done at prediction time


## The Algorithm

1. Store all training examples
2. For new query  $x$ : find K nearest training points
3. Return majority class among neighbors

---

Instance-based learning: the training data IS the model

# Distance Metrics



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L03\\_KNN\\_KMeans/02\\_distance\\_metrics](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L03_KNN_KMeans/02_distance_metrics)

**Common:** Euclidean  $\sqrt{\sum(x_i - y_i)^2}$ , Manhattan  $\sum|x_i - y_i|$

# Distance Metrics: Minkowski Family

## Minkowski Distance

$$d_p(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

- $p = 1$ : Manhattan (L1)
- $p = 2$ : Euclidean (L2)
- $p = \infty$ : Chebyshev (max absolute difference)


## Choosing p

- $p = 2$ : Default, works well in most cases
- $p = 1$ : More robust to outliers
- Higher  $p$ : Sensitive to large single differences

---

In high dimensions, all distances become similar (curse of dimensionality)

# KNN Decision Boundaries



## Boundary Properties

- Non-linear, locally adaptive
- Small K: complex boundary, may overfit
- Large K: smoother boundary, may underfit

# Choosing K

## The Bias-Variance Trade-off

- $K = 1$ : High variance, low bias (very flexible)
- $K = n$ : High bias, low variance (always predicts majority class)

## Practical Guidelines

- Start with  $K = \sqrt{n}$  where  $n$  is training size
- Use odd K for binary classification (avoid ties)
- Cross-validation to find optimal K

**Common Choices:**  $K \in \{3, 5, 7, 11\}$

---

Small K for complex patterns, larger K for noisy data

## Problem with Equal Voting

- All K neighbors have equal influence
- A distant neighbor counts as much as closest neighbor

## Solution: Distance Weighting

$$w_i = \frac{1}{d(\mathbf{x}, \mathbf{x}_i)^2}$$

- Closer neighbors get higher weight
- Reduces sensitivity to K choice

## In scikit-learn

- `weights='uniform'`: equal weights (default)
- `weights='distance'`: inverse distance weighting

---

Distance weighting often improves performance

## Why Scaling Matters

Without scaling:

- Income: ranges 20,000–200,000
- Age: ranges 20–80
- Distance dominated by income (larger scale)!

## Scaling Methods

- **Standardization:**  $z = \frac{x - \mu}{\sigma}$  (mean=0, std=1)
- **Min-Max:**  $x' = \frac{x - x_{min}}{x_{max} - x_{min}}$  (range [0,1])

**Rule:** Always scale features for distance-based methods!

---

StandardScaler for Gaussian-like features, MinMaxScaler for bounded

## The Problem

In high dimensions:

- All points become approximately equidistant
- “Nearest neighbor” becomes meaningless
- Exponentially more data needed

## Solutions

- **Dimensionality reduction:** PCA before KNN
- **Feature selection:** keep only relevant features
- **Use domain knowledge:** select meaningful features

---

KNN works best with moderate number of features ( $\leq 20$ )

## Time Complexity

- Training:  $O(1)$  – just store data!
- Prediction:  $O(nd)$  for brute force ( $n$  samples,  $d$  features)

## Acceleration Techniques

- **KD-Tree**:  $O(d \log n)$  average for low  $d$
- **Ball Tree**: Works better in higher dimensions
- **Approximate NN**: Trade accuracy for speed

## In scikit-learn

- `algorithm='auto'`: automatically chooses best
- `algorithm='brute'`: force brute force

---

For large datasets: consider approximate methods or trees

## Classification

- `from sklearn.neighbors import KNeighborsClassifier`
- `knn = KNeighborsClassifier(n_neighbors=5)`
- `knn.fit(X_train, y_train)`
- `y_pred = knn.predict(X_test)`

## Key Parameters

- `n_neighbors`: K value
- `weights`: 'uniform' or 'distance'
- `metric`: 'euclidean', 'manhattan', etc.
- `algorithm`: 'auto', 'ball\_tree', 'kd\_tree', 'brute'

---

Also available: `KNeighborsRegressor` for regression tasks

## Open the Colab Notebook

- Exercise 1: Implement KNN classifier from scratch
- Exercise 2: Apply K-Means to customer segmentation data
- Exercise 3: Compare distance metrics and k values

Link: <https://colab.research.google.com/> [TBD]

# K-Means: The Idea

**Goal:** Partition  $n$  points into  $K$  clusters

**Objective:** Minimize within-cluster sum of squares (WCSS)

$$\sum_{k=1}^K \sum_{x \in C_k} \|x - \mu_k\|^2$$

where  $\mu_k$  is the centroid of cluster  $C_k$

## Key Insight

- Each point assigned to nearest centroid
- Centroids are cluster means
- Iterative refinement until convergence

---

K-Means finds locally optimal solution (not guaranteed global)

# K-Means Algorithm

```
1: Input: Data  $\mathbf{X}$ , number of clusters  $K$ 
2: Initialize  $K$  centroids randomly
3: repeat
4:   Assignment: assign each point to nearest centroid
5:   Update: recompute centroids as cluster means
6: until centroids don't change (or max iterations)
7: return cluster assignments, centroids
```


## Convergence

- Guaranteed to converge (WCSS decreases each iteration)
- May converge to local optimum

---

Each iteration:  $O(nKd)$  where  $n$  = samples,  $K$  = clusters,  $d$  = features

# K-Means: Visualization



Final state after convergence: points colored by cluster, X marks centroids

# Initialization Strategies

## Random Initialization

- Pick K random points as initial centroids
- Sensitive to choice, may get poor solution
- Run multiple times, keep best result

## K-Means++ (Default in scikit-learn)


- Smart initialization: spread out initial centroids
- First centroid: random
- Next centroids: probability proportional to squared distance

**Result:** Much better starting point, fewer iterations

---

K-Means++ gives provably better initialization with theoretical guarantees

## Choosing K: Elbow Method



**Interpretation:** Look for the “elbow” where WCSS stops dropping sharply

## Choosing K: Silhouette Score

For each point  $i$

- $a(i)$ : average distance to points in same cluster
- $b(i)$ : average distance to points in nearest other cluster

Silhouette score


$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

- Range:  $[-1, 1]$
- $s \approx 1$ : point is well-matched to cluster
- $s \approx 0$ : point is on boundary
- $s < 0$ : point may be in wrong cluster

---


Average silhouette score summarizes overall clustering quality

# Silhouette Plot



All clusters should have similar width and scores above average line

# K-Means: Decision Regions



**K-Means creates Voronoi tessellation around centroids**

# K-Means Assumptions

## What K-Means Assumes

- Clusters are spherical (isotropic)
- Clusters have similar sizes
- Clusters have similar densities

## When K-Means Fails

- Non-convex shapes (e.g., crescents)
- Very different cluster sizes
- Different densities
- Outliers (pull centroids away)

---

Consider DBSCAN or Gaussian Mixture Models for these cases

## Mini-Batch K-Means

- Uses random subsets for updates
- Much faster for large datasets
- Slightly worse results

## K-Medoids

- Centroids must be actual data points
- More robust to outliers
- Slower than K-Means

## K-Means for Mixed Data

- K-Modes: for categorical data
- K-Prototypes: mixed continuous and categorical

---

Mini-Batch K-Means: good for  $\geq 10k$  samples

## Basic Usage

- `from sklearn.cluster import KMeans`
- `kmeans = KMeans(n_clusters=3, random_state=42)`
- `labels = kmeans.fit_predict(X)`
- `centroids = kmeans.cluster_centers_`

## Key Parameters

- `n_clusters`: K (required)
- `init`: 'k-means++' (default) or 'random'
- `n_init`: number of runs (default 10)
- `max_iter`: max iterations per run

---

`inertia_` attribute gives WCSS after fitting

## KNN vs K-Means: Key Differences

Aspect	KNN	K-Means
Task	Classification/Regression	Clustering
Learning	Supervised (needs labels)	Unsupervised
K meaning	Number of neighbors	Number of clusters
Training	None (lazy)	Iterative optimization
Prediction	Compute distances	Assign to centroid
Output	Class label	Cluster ID

---

The “K” in KNN and K-Means mean completely different things!

**Problem:** Group customers for targeted marketing

## Features

- Transaction frequency
- Average transaction amount
- Account balance
- Product holdings

## K-Means Solution

- Cluster into segments (e.g., K=4)
- Profile each segment
- Tailor offerings to segment needs

---

Example segments: High-value frequent, Dormant, New/Growing, Price-sensitive

**Problem:** Flag suspicious transactions

## KNN Approach

- Features: amount, time, location, merchant category
- Find K similar past transactions
- If most neighbors are fraud, flag as suspicious

## K-Means Approach

- Cluster “normal” transactions
- New transaction far from all centroids = anomaly
- Combined with distance threshold

---

KNN needs labeled fraud examples, K-Means detects deviation from normal

## Use KNN When

- You have labeled training data
- Local patterns matter (non-linear boundaries)
- Interpretability: “similar to these examples”
- Moderate dataset size


## Use K-Means When

- No labels available
- Looking for natural groupings
- Clusters are roughly spherical
- Need fast clustering of large data

---

K-Means often used as preprocessing before supervised learning

## KNN vs K-Means Decision Guide



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/l03\\_KNN\\_KMeans/07\\_decision\\_flowchart](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/l03_KNN_KMeans/07_decision_flowchart)

Start with KNN for classification, K-Means for clustering

## Instead of KNN

- Large data: use ball tree or approximate NN
- Need probability: logistic regression
- Many features: random forest

## Instead of K-Means

- Unknown K: DBSCAN (density-based)
- Hierarchical structure: agglomerative clustering
- Soft assignments: Gaussian Mixture Models
- Non-spherical: spectral clustering

---

DBSCAN automatically determines number of clusters

## K-Nearest Neighbors

- Instance-based, lazy learner
- Scale features, choose K via cross-validation
- Works best with moderate features, moderate data size

## K-Means

- Iterative: assign points, update centroids
- K-Means++ for initialization, elbow/silhouette for K
- Assumes spherical, similar-size clusters

## Common Considerations

- Feature scaling is critical for both
- “K” means different things in each algorithm

---

Both are foundational algorithms: simple, interpretable, widely used

# References

## Textbooks

- James et al. (2021). *ISLR*, Chapters 2, 12
- Hastie et al. (2009). *ESL*, Chapters 13, 14

## Key Papers

- Arthur & Vassilvitskii (2007). K-Means++
- Cover & Hart (1967). Nearest Neighbor Pattern Classification

## Next Lecture

- L04: Random Forests
- Ensemble methods and feature importance