

## L03: KNN & K-Means

Mathematical Foundations and Implementation

Methods and Algorithms

Spring 2026

# Outline

1 K-Nearest Neighbors

2 Practice

3 K-Means Clustering

4 Comparison and Applications

5 Summary

**By the end of this lecture, you will be able to:**

1. Analyze the bias-variance tradeoff in KNN and derive its asymptotic error bounds
2. Evaluate clustering validity using statistical tests (Hopkins, Gap statistic, silhouette)
3. Prove K-Means convergence and analyze computational complexity
4. Compare distance metrics and assess their suitability for high-dimensional finance data

**Finance Applications:** Customer segmentation, fraud detection

---

Bloom's Level 4–5: Analyze, Evaluate, Prove, Compare

## Key Insight

- No explicit model training (store all data)
- Classification by majority vote of K nearest neighbors
- “Lazy” because work is done at prediction time


## The Algorithm

1. Store all training examples
2. For new query  $x$ : find K nearest training points
3. Return majority class among neighbors

---

Instance-based learning: the training data IS the model

# Distance Metrics



**Common:** Euclidean  $\sqrt{\sum(x_i - y_i)^2}$ , Manhattan  $\sum|x_i - y_i|$

## Minkowski Distance

$$d_p(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

- $p = 1$ : Manhattan (L1)
- $p = 2$ : Euclidean (L2)
- $p = \infty$ : Chebyshev (max absolute difference)
- Note:  $p < 1$  violates triangle inequality (not a true metric)

## Choosing p

- $p = 2$ : Default, works well in most cases
- $p = 1$ : More robust to outliers
- Higher  $p$ : Sensitive to large single differences

---

In high dimensions, all distances become similar (curse of dimensionality)

## Cosine Similarity (critical for text/embeddings)

$$\cos(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$$

- Measures angle, not magnitude
- Range:  $[-1, 1]$  where 1 = identical direction
- Use for: document similarity, word embeddings, high-dimensional sparse data

## Mahalanobis Distance (accounts for correlation)


$$d_M(\mathbf{x}, \mathbf{y}) = \sqrt{(\mathbf{x} - \mathbf{y})^T \Sigma^{-1} (\mathbf{x} - \mathbf{y})}$$

- Considers feature covariance  $\Sigma$
- Unit-less, scale-invariant
- Detects outliers accounting for correlation structure

---

Choose metric based on data type: Euclidean for dense, Cosine for sparse/text

# KNN Decision Boundaries



## Boundary Properties

- Non-linear, locally adaptive
- Small K: complex boundary, may overfit
- Large K: smoother boundary, may underfit

# Choosing K

## The Bias-Variance Trade-off

- $K = 1$ : High variance, low bias (very flexible)
- $K = n$ : High bias, low variance (always predicts majority class)

## Practical Guidelines

- Start with  $K = \sqrt{n}$  where  $n$  is training size
- Use odd K for binary classification (avoid ties)
- Cross-validation to find optimal K

**Common Choices:**  $K \in \{3, 5, 7, 11\}$

---

Small K for complex patterns, larger K for noisy data

## GridSearchCV for Optimal K

- from sklearn.model\_selection import GridSearchCV
- param\_grid = {'n\_neighbors': range(1, 21, 2)}
- grid = GridSearchCV(KNeighborsClassifier(), param\_grid, cv=5)
- grid.fit(X\_train, y\_train)
- best\_k = grid.best\_params\_['n\_neighbors']

## Validation Curve

- Plot accuracy vs K for training and validation sets
- Choose K where validation accuracy peaks (before overfitting)

---

Cross-validation provides objective K selection

## Problem with Equal Voting

- All K neighbors have equal influence
- A distant neighbor counts as much as closest neighbor

## Solution: Inverse-Distance Weighting

$$\hat{y} = \frac{\sum_{i \in N_k(\mathbf{x})} w_i y_i}{\sum_{i \in N_k(\mathbf{x})} w_i} \quad \text{where } w_i = \frac{1}{d(\mathbf{x}, \mathbf{x}_i)}$$

- Closer neighbors get higher weight
- Reduces sensitivity to K choice
- For  $d = 0$  (exact match): return that point's class directly

## In scikit-learn

- `weights='uniform'`: equal weights (default)
- `weights='distance'`: inverse distance weighting

---

Distance weighting often improves performance

## Why Scaling Matters

Without scaling:

- Income: ranges 20,000–200,000
- Age: ranges 20–80
- Distance dominated by income (larger scale)!

## Scaling Methods

- **Standardization:**  $z = \frac{x - \mu}{\sigma}$  (mean=0, std=1)
- **Min-Max:**  $x' = \frac{x - x_{min}}{x_{max} - x_{min}}$  (range [0,1])

## Edge Cases

- If  $\sigma = 0$  (constant feature): drop feature or set  $z = 0$
- If  $x_{max} = x_{min}$  (constant): drop feature or set  $x' = 0$

Rule: Always scale features for distance-based methods!

---

StandardScaler for Gaussian-like features, MinMaxScaler for bounded

## The Problem

In high dimensions:

- All points become approximately equidistant (Beyer et al., 1999)
- “Nearest neighbor” becomes meaningless
- Exponentially more data needed

## Solutions

- **Dimensionality reduction:** PCA before KNN
- **Feature selection:** keep only relevant features
- **Use domain knowledge:** select meaningful features

---

KNN works best with moderate number of features (typically  $d < 15 - 20$ , problem-dependent)

## Time Complexity

- Training:  $O(1)$  – just store data!
- Prediction:  $O(nd)$  for brute force ( $n$  samples,  $d$  features)

## Acceleration Techniques

- **KD-Tree**:  $O(d \log n)$  average for  $d < 15$  (degrades to  $O(n)$  in high dimensions)
- **Ball Tree**: Works better in higher dimensions
- **Approximate NN**: Trade accuracy for speed

## In scikit-learn

- `algorithm='auto'`: automatically chooses best
- `algorithm='brute'`: force brute force

---

For large datasets: consider approximate methods or trees

## Classification

- `from sklearn.neighbors import KNeighborsClassifier`
- `knn = KNeighborsClassifier(n_neighbors=5)`
- `knn.fit(X_train, y_train)`
- `y_pred = knn.predict(X_test)`

## Key Parameters

- `n_neighbors`: K value
- `weights`: 'uniform' or 'distance'
- `metric`: 'euclidean', 'manhattan', etc.
- `algorithm`: 'auto', 'ball\_tree', 'kd\_tree', 'brute'

---

Also available: `KNeighborsRegressor` for regression tasks

## Cover & Hart (1967) Consistency Theorem

- As  $n \rightarrow \infty$ : 1-NN error rate  $\leq 2 \times$  Bayes error
- KNN is **universally consistent**: converges to optimal
- Requires:  $K \rightarrow \infty$  but  $K/n \rightarrow 0$

## VC Dimension

- 1-NN has infinite VC dimension in general
- Shattering depends on geometry, not sample size
- Implications: prone to overfitting without regularization (large K)

**Practical Implication:** KNN's asymptotic optimality requires **LARGE** datasets

---

The  $2 \times$  Bayes error bound makes KNN a strong baseline

**Theorem:** As  $n \rightarrow \infty$ ,  $K \rightarrow \infty$ ,  $K/n \rightarrow 0$ , the KNN error rate  $R_{\text{KNN}}^* \rightarrow R^*$  (Bayes rate).

**Proof idea** (1-NN case, binary classification):

1. As  $n \rightarrow \infty$ , the nearest neighbor  $\mathbf{x}_{(1)} \rightarrow \mathbf{x}$  (converges to query point)
2. The 1-NN error at  $\mathbf{x}$  is:

$$R_{\text{1-NN}}(\mathbf{x}) = \eta(\mathbf{x})(1 - \eta(\mathbf{x})) + (1 - \eta(\mathbf{x}))\eta(\mathbf{x}) = 2\eta(\mathbf{x})(1 - \eta(\mathbf{x})) \quad (1)$$

where  $\eta(\mathbf{x}) = P(Y = 1|\mathbf{x})$ .

Since  $R^*(\mathbf{x}) = \min(\eta, 1 - \eta)$  and  $2\eta(1 - \eta) \leq 2R^*(1 - R^*) \leq 2R^*$ :

$$R^* \leq R_{\text{1-NN}} \leq 2R^*(1 - R^*) \quad (2)$$

---

The 1-NN error is at most twice the Bayes rate. For K-NN with  $K/n \rightarrow 0$ :  $R_{\text{K-NN}} \rightarrow R^*$  exactly.

## Bias-Variance Decomposition for KNN

For KNN regression, the expected prediction error at  $\mathbf{x}$  decomposes as:

$$E[(\hat{f}_K(\mathbf{x}) - Y)^2] = \underbrace{\text{Bias}^2(\hat{f}_K(\mathbf{x}))}_{\text{increases with } K} + \underbrace{\text{Var}(\hat{f}_K(\mathbf{x}))}_{\text{decreases with } K} + \sigma^2 \quad (3)$$

For the KNN estimator  $\hat{f}_K(\mathbf{x}) = \frac{1}{K} \sum_{i \in N_K(\mathbf{x})} y_i$ :

- **Variance:**  $\text{Var}(\hat{f}_K) = \frac{\sigma^2}{K}$  (averaging  $K$  neighbors reduces variance)
- **Bias:**  $\text{Bias}(\hat{f}_K) \approx f(\mathbf{x}) - \frac{1}{K} \sum_{i \in N_K} f(\mathbf{x}_i)$  (larger  $K \Rightarrow$  neighbors farther away  $\Rightarrow$  more bias)
- **Optimal  $K$ :** balances this tradeoff; found via cross-validation

---

At  $K = 1$ : zero bias, variance =  $\sigma^2$ . At  $K = n$ : high bias (predicts global mean), variance =  $\sigma^2/n$ .

## Open the Colab Notebook

- Exercise 1: Implement weighted KNN with distance weighting
- Exercise 2: Compare Gap statistic vs Elbow for K selection
- Exercise 3: Apply SMOTE for imbalanced fraud detection

**Link:** See course materials for Colab notebook

# K-Means: The Idea

**Goal:** Partition  $n$  points into  $K$  clusters

**Objective:** Minimize within-cluster sum of squares (WCSS)

$$\sum_{k=1}^K \sum_{x \in C_k} \|x - \mu_k\|^2$$

where  $\mu_k$  is the centroid of cluster  $C_k$

## Key Insight

- Each point assigned to nearest centroid
- Centroids are cluster means
- Iterative refinement until convergence

---

K-Means finds locally optimal solution (not guaranteed global)

# K-Means Algorithm

```
1: Input: Data  $\mathbf{X}$ , number of clusters  $K$ 
2: Initialize  $K$  centroids randomly
3: repeat
4:   Assignment: assign each point to nearest centroid
5:   Update: recompute centroids as cluster means
6:   Handle empty: if cluster empty, reinitialize from farthest point
7: until centroid change  $< \epsilon$  (e.g.,  $10^{-4}$ ) or max iterations
8: return cluster assignments, centroids
```

## Convergence

- Guaranteed to converge (WCSS decreases each iteration)
- May converge to local optimum

---

Each iteration:  $O(nKd)$  where  $n$  = samples,  $K$  = clusters,  $d$  = features

# K-Means Convergence Proof

**Theorem:** Lloyd's algorithm converges in a finite number of iterations.

**Proof:** The WCSS objective  $J = \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$  is:

1. **Assignment step** (fix  $\mu$ , optimize  $C$ ): Assigning each point to its nearest centroid minimizes  $J \Rightarrow J$  decreases or stays equal
2. **Update step** (fix  $C$ , optimize  $\mu$ ): The mean  $\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$  minimizes within-cluster sum of squares  $\Rightarrow J$  decreases or stays equal

Since  $J \geq 0$  and strictly decreasing at each step, and there are finitely many partitions of  $n$  points into  $K$  clusters: convergence is guaranteed.  $\square$

**NP-hardness:** Finding the globally optimal K-Means solution is NP-hard (Aloise et al., 2009). Lloyd's algorithm finds a local optimum; K-Means++ initialization provides an  $O(\log k)$ -competitive approximation guarantee.

---

This is a coordinate descent argument: alternating optimization of two blocks ( $C$  and  $\mu$ ) on a bounded objective.

## Connection to Expectation-Maximization

- K-Means = “Hard EM” for Gaussian Mixture Models
- Assumes: spherical Gaussians with equal variance
- E-step: assign points to nearest centroid (hard assignment)
- M-step: update centroids as cluster means


## Why This Matters

- Explains WHY convergence is guaranteed (EM always converges)
- Explains WHY K-Means assumes spherical clusters
- Opens door to soft clustering via GMM (next course)

---

Understanding EM connection deepens theoretical understanding

# K-Means: Visualization



Final state after convergence: points colored by cluster, X marks centroids

## Random Initialization

- Pick K random points as initial centroids
- Sensitive to choice, may get poor solution
- Run multiple times, keep best result

## K-Means++ (Default in scikit-learn)


- Smart initialization: spread out initial centroids
- First centroid: random
- Next centroids: probability proportional to squared distance

**Result:** Much better starting point, fewer iterations

---

K-Means++ achieves  $O(\log k)$ -competitive: expected cost  $\leq 8(\ln k + 2) \times$  optimal (Arthur & Vassilvitskii, 2007)

## Choosing K: Elbow Method



**Interpretation:** Look for the “elbow” where WCSS stops dropping sharply

## Choosing K: Silhouette Score

For each point  $i$

- $a(i)$ : average distance to points in same cluster
- $b(i)$ : average distance to points in nearest other cluster

Silhouette score


$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

- Range:  $[-1, 1]$
- $s \approx 1$ : point is well-matched to cluster
- $s \approx 0$ : point is on boundary
- $s < 0$ : point may be in wrong cluster
- Singleton cluster:  $s(i) = 0$  by convention

---

Average silhouette score summarizes overall clustering quality

## Silhouette Plot



All clusters should have similar width and scores above average line

# Should We Cluster? Hopkins Statistic

## Before Clustering: Test for Cluster Tendency

$$H = \frac{\sum u_i}{\sum u_i + \sum w_i}$$

- $u_i$  = distances from random points to nearest data point
- $w_i$  = distances from data points to nearest neighbor
- $H \approx 0.5$ : uniform distribution (no clusters)
- $H > 0.75$ : significant clustering tendency

**Use Case:** Before running K-Means, verify data HAS cluster structure

---

Don't cluster uniform data - Hopkins test catches this

## The Problem with Elbow Method

- Elbow is subjective - where exactly IS the elbow?
- No formal statistical test

## Gap Statistic (Tibshirani et al., 2001)

$$\text{Gap}_n(k) = E_n^*[\log W_k] - \log W_k$$


where  $W_k = \sum_{r=1}^k \frac{1}{2|C_r|} \sum_{i,j \in C_r} \|x_i - x_j\|^2$  and  $E_n^*$  is computed over reference distributions.

- Choose smallest K where:  $\text{Gap}(k) \geq \text{Gap}(k+1) - s_{k+1}$
- $s_{k+1}$  = standard error from bootstrap reference samples

---

Gap statistic provides statistical justification for K selection

# K-Means: Decision Regions



**K-Means creates Voronoi tessellation: regions where all points are closest to one centroid**

## What K-Means Assumes

- Clusters are spherical (isotropic)
- Clusters have similar sizes
- Clusters have similar densities

## When K-Means Fails

- Non-convex shapes (e.g., crescents)
- Very different cluster sizes
- Different densities
- Outliers (pull centroids away)

---

Consider DBSCAN or Gaussian Mixture Models for these cases

## Mini-Batch K-Means

- Uses random subsets for updates
- Much faster for large datasets
- Slightly worse results

## K-Medoids

- Centroids must be actual data points
- More robust to outliers
- Slower than K-Means

## K-Means for Mixed Data

- K-Modes: for categorical data
- K-Prototypes: mixed continuous and categorical

---

Mini-Batch K-Means: good for  $\geq 10k$  samples

## Basic Usage

- `from sklearn.cluster import KMeans`
- `kmeans = KMeans(n_clusters=3, random_state=42)`
- `labels = kmeans.fit_predict(X)`
- `centroids = kmeans.cluster_centers_`

## Key Parameters

- `n_clusters`: K (required)
- `init`: 'k-means++' (default) or 'random'
- `n_init`: number of runs (default 10)
- `max_iter`: max iterations per run

---

`inertia_` attribute gives WCSS after fitting

## KNN vs K-Means: Key Differences

Aspect	KNN	K-Means
Task	Classification/Regression	Clustering
Learning	Supervised (needs labels)	Unsupervised
K meaning	Number of neighbors	Number of clusters
Training	None (lazy)	Iterative optimization
Prediction	Compute distances	Assign to centroid
Output	Class label	Cluster ID

---

The “K” in KNN and K-Means mean completely different things!

## RFM Analysis (Industry Standard)

- Recency: Days since last transaction
- Frequency: Number of transactions in period
- Monetary: Total/average transaction value

## K-Means on RFM Features

- Standardize each RFM dimension (different scales!)
- Cluster into segments (typically K=4-6)
- Profile: “Champions” (high R,F,M) vs “At Risk” (low R)

## Business Value

- Customer Lifetime Value (CLV) prediction per segment
- Targeted retention campaigns by segment risk

---

RFM segmentation is foundational for CRM and marketing analytics

## CRITICAL: Class Imbalance Problem

- Fraud is typically <1% of transactions (100:1 ratio)
- Naive KNN majority vote: ALWAYS predicts non-fraud
- 99% accuracy but detects ZERO fraud!

## Solutions for Imbalanced Data

- **SMOTE**: Synthetic Minority Oversampling TEchnique
- **Weighted KNN**: Higher weight for minority class neighbors
- **Cost-sensitive**: False Negative costs  $>>$  False Positive
- **Anomaly score**: Use distance to K-th neighbor, not vote
- Use **Precision-Recall AUC**, NOT accuracy!

---

Always check class distribution before applying majority vote!

## Use KNN When

- You have labeled training data
- Local patterns matter (non-linear boundaries)
- Interpretability: “similar to these examples”
- Moderate dataset size


## Use K-Means When

- No labels available
- Looking for natural groupings
- Clusters are roughly spherical
- Need fast clustering of large data

---

K-Means often used as preprocessing before supervised learning

## KNN vs K-Means Decision Guide



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/l03\\_KNN\\_KMeans/07\\_decision\\_flowchart](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/l03_KNN_KMeans/07_decision_flowchart)

Start with KNN for classification, K-Means for clustering

## DBSCAN (Density-Based Spatial Clustering):

- **Core point:** has  $\geq \text{minPts}$  neighbors within radius  $\epsilon$
- **Border point:** within  $\epsilon$  of a core point but not itself core
- **Noise:** neither core nor border  $\Rightarrow$  automatically detected as outliers
- Advantage: discovers non-spherical clusters; does not require  $K$

## Hierarchical (Agglomerative) Clustering:

- Bottom-up: each point starts as its own cluster, iteratively merge closest pairs
- Linkage criteria: single (min), complete (max), Ward (minimize variance)
- Produces a **dendrogram**: cut at desired height to get  $K$  clusters

---

Use DBSCAN when clusters have irregular shapes. Use hierarchical when you want to explore multiple  $K$  values via the dendrogram.

## Instead of KNN

- Large data: use ball tree or approximate NN
- Need probability: logistic regression
- Many features: random forest

## Instead of K-Means

- Soft assignments: Gaussian Mixture Models
- Non-spherical: spectral clustering
- Streaming data: Mini-Batch K-Means

---

Choose algorithm based on data characteristics and problem requirements

## K-Nearest Neighbors

- Instance-based, lazy learner
- Scale features, choose K via cross-validation
- Works best with moderate features, moderate data size

## K-Means

- Iterative: assign points, update centroids
- K-Means++ for initialization, elbow/silhouette for K
- Assumes spherical, similar-size clusters

## Common Considerations

- Feature scaling is critical for both
- “K” means different things in each algorithm

---

Both are foundational algorithms: simple, interpretable, widely used

# References

## Textbooks

- James et al. (2021). *ISLR*, Chapters 2, 12
- Hastie et al. (2009). *ESL*, Chapters 13, 14

## Key Papers

- Arthur & Vassilvitskii (2007). K-Means++
- Cover & Hart (1967). Nearest Neighbor Pattern Classification

## Next Lecture

- L04: Random Forests
- Ensemble methods and feature importance