

## L04: Random Forests

Full Lecture: Ensemble Learning, Variance Reduction, and Fraud Detection


### Methods and Algorithms

MSc Data Science

# The Ensemble Approach

The cartoon captures the absurdity – and power – of combining models.

- One model is fragile; many models are robust
- Ensembles exploit the **wisdom of crowds**: independent errors cancel
- Random Forests, boosting, and stacking all follow this principle
- The key insight: **diversity** among models matters more than individual accuracy



XKCD #1885 by Randall Munroe (CC BY-NC 2.5)

By the end of this lecture, you will be able to:

1. **Derive** the variance reduction formula for ensemble averaging (Analyze)
2. **Evaluate** RF vs. boosting through the bias-variance lens (Evaluate)
3. **Analyze** feature importance: MDI, permutation, SHAP (Analyze)
4. **Apply** ensembles to fraud detection with class imbalance (Apply)
5. **Compare** RF, AdaBoost, gradient boosting, XGBoost (Evaluate)

## Bloom's Taxonomy Levels

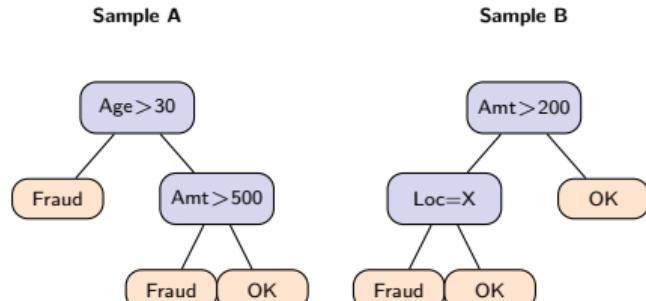
- **Analyze** – break down, derive, compare components
- **Evaluate** – judge tradeoffs, assess suitability
- **Apply** – use methods on real data problems

---

L04 covers ensemble methods from bagging to boosting with finance applications.

# Why Would a Single Tree Overfit Every Dataset It Touches?

- A fully grown tree has **low bias** (fits training data perfectly)
- But **high variance**: small data changes produce entirely different trees
- Bias-variance decomposition:  $MSE = \text{Bias}^2 + \text{Variance} + \text{Noise}$
- Trees are **unstable learners** – ideal building blocks for ensembles



Different training samples  $\Rightarrow$  completely different tree structures.

## Insight

High variance is a *feature*, not a bug – it means trees are sensitive enough to be improved by averaging.

---

Unstable learners benefit most from ensemble methods (Breiman, 1996).

# How Does a Tree Choose the Best Split? (Gini Impurity)

**Gini Impurity** measures how often a random sample would be misclassified:

$$G = 1 - \sum_{k=1}^K p_k^2$$

**Binary case:**  $G = 2p(1 - p)$

**Worked example (fraud data):**

- Node: 800 legit, 200 fraud  $\Rightarrow p_{\text{fraud}} = 0.2$
- $G = 1 - (0.8^2 + 0.2^2) = 1 - 0.68 = 0.32$
- Best split minimizes weighted Gini of children

**Split evaluation:**

- Left child: 50 fraud / 100 total  $\Rightarrow G_L = 0.50$
- Right child: 150 fraud / 900 total  $\Rightarrow G_R = 0.28$
- Weighted:  $\frac{100}{1000}(0.50) + \frac{900}{1000}(0.28) = 0.30$
- Gain:  $0.32 - 0.30 = 0.02$

**Properties:**

- $G = 0$ : pure node (one class only)
- $G = 0.5$ : maximum impurity (binary)
- Computationally cheaper than entropy

## Insight

Gini impurity is a *concave* function of class proportions – every split on a mixed node reduces impurity.

CART uses Gini by default; scikit-learn's `DecisionTreeClassifier(criterion='gini')`.

# What Does Entropy Add Beyond Gini?

## Shannon Entropy:

$$H = - \sum_{k=1}^K p_k \log_2 p_k$$

## Information Gain:

$$\text{IG} = H(\text{parent}) - \sum_j \frac{n_j}{n} H(\text{child}_j)$$

- Entropy is zero for pure nodes, maximal at uniform distribution
- IG selects the feature that reduces uncertainty the most
- C4.5 and ID3 algorithms use entropy; CART uses Gini
- **Regression trees:** use MSE reduction (variance reduction) as the splitting criterion instead of Gini/Entropy

## Gini vs. Entropy Comparison:

Property	Gini	Entropy
Range (binary)	[0, 0.5]	[0, 1]
Computation	Faster	log needed
Sensitivity	Majority class	Rare classes
Algorithm	CART	C4.5, ID3

In practice, Gini and entropy produce **nearly identical** trees.

## Insight

Entropy penalizes impure nodes slightly more than Gini – in fraud detection, this can surface rare-class splits earlier.

For binary classification,  $H \approx 2G$  near  $p = 0.5$ ; the choice rarely changes the final tree.

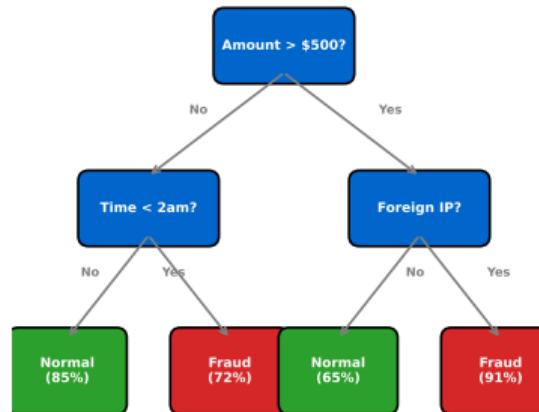
# What Does a Trained Decision Tree Actually Look Like?

**What you see:** A tree trained on financial features, splitting recursively on the most informative thresholds.

**Key pattern:** Each internal node tests one feature; each leaf assigns a class or value.  
Deeper trees fit noise.

**Takeaway:** A single tree is interpretable but brittle – small data shifts change the entire structure.

Decision Tree for Fraud Detection



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/04\\_Random\\_Forests/01\\_decision\\_tree](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/04_Random_Forests/01_decision_tree)

## Insight

Interpretability is the single tree's greatest strength – and its fragility is the motivation for ensembles.

**Pruning (max\_depth, min\_samples\_leaf) trades bias for variance reduction.**

# Why Does Averaging Multiple Models Reduce Variance?

**Key derivation.** For  $B$  models with variance  $\sigma^2$  and pairwise correlation  $\rho$ :

$$\text{Var}\left(\frac{1}{B} \sum_{b=1}^B X_b\right) = \rho\sigma^2 + \frac{1-\rho}{B} \sigma^2$$

**Two extreme cases:**

- $\rho = 1$  (identical models):  $\text{Var} = \sigma^2$  – no improvement
- $\rho = 0$  (independent models):  $\text{Var} = \sigma^2/B$  – perfect scaling

The **goal of bagging**: reduce  $\rho$  by training on different bootstrap samples while keeping  $\sigma^2$  large (deep trees).

**Derivation sketch:**

- $\text{Var}(\bar{X}) = \frac{1}{B^2} \sum_{i,j} \text{Cov}(X_i, X_j)$
- Diagonal terms:  $B \cdot \sigma^2$
- Off-diagonal:  $B(B-1)\rho\sigma^2$
- Combine:  $\frac{1}{B^2} [B\sigma^2 + B(B-1)\rho\sigma^2]$
- Simplify to the variance formula above

**Practical implication:**

With  $B = 500$  and  $\rho = 0.05$ :  
 $\text{Var} \approx 0.05\sigma^2 + \frac{0.95}{500}\sigma^2 \approx 0.052\sigma^2$

## Insight

The first term  $\rho\sigma^2$  cannot be reduced by adding more trees – only *decorrelation* can shrink it.

---

This is why Random Forests add feature randomization on top of bagging.


# How Does Bootstrap Sampling Create Diversity?

**What you see:** Bootstrap resampling draws  $n$  samples with replacement from  $n$  observations.

**Key pattern:** Each bootstrap sample contains  $\sim 63.2\%$  unique observations; the remaining  $\sim 36.8\%$  are left out (OOB).

**Takeaway:** Bootstrap creates model diversity through data perturbation – each tree sees a different “view” of the training set.

- $P(\text{not selected}) = (1 - 1/n)^n \rightarrow 1/e \approx 0.368$
- Duplicates force trees to emphasize different patterns
- OOB samples become free validation data



Each tree trained on  $\sim 63\%$  unique samples (with replacement)

[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/03\\_bootstrap](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/03_bootstrap)

## Insight

Bootstrap is a “poor man’s” way to approximate drawing from the true data-generating distribution.

Efron (1979) introduced bootstrap; Breiman (1996) applied it to tree ensembles.


# What Happens to the 36.8% of Samples Left Out?

**What you see:** OOB error converges as the number of trees grows, providing a built-in cross-validation estimate.

**Key pattern:** Each observation is OOB for  $\sim 36.8\%$  of trees. Aggregate their predictions to get an unbiased error estimate.

**Takeaway:** OOB error eliminates the need for a separate validation set – especially valuable when data is scarce.

- OOB error  $\approx$  leave-one-out CV error
- No data “wasted” on validation
- Monitored automatically in scikit-learn with `oob_score=True`



## Insight

OOB error is slightly pessimistic (each tree trained on  $\sim 63\%$  of data), making it a conservative estimate.

OOB error is one of Random Forests' most elegant properties – free, unbiased validation.


# What Is the One Trick That Makes Random Forests Better Than Bagging?

**Feature randomization.** At each split, consider only a random subset of features:

- Classification:  $m = \lfloor \sqrt{p} \rfloor$
- Regression:  $m = \lfloor p/3 \rfloor$
- This forces trees to use **different features**, reducing correlation  $\rho$

Why it works:

- Without randomization, every tree splits on the same dominant feature first
- Correlated trees  $\Rightarrow$  correlated errors  $\Rightarrow$  averaging helps less
- Feature subsampling **decorrelates** the trees



## Insight

The "random" in Random Forests refers to feature subsampling at each split – not random data (that is bagging).

Breiman (2001): "Random Forests." Machine Learning, 45(1), 5–32.

# How Does Feature Randomization Reduce Tree Correlation?

Recall the ensemble variance formula:

$$\text{Var}(\bar{f}) = \rho \sigma^2 + \frac{1 - \rho}{B} \sigma^2$$

Effect of feature randomization on  $\rho$ :

- Bagging alone:  $\rho \approx 0.5\text{--}0.9$  (trees still correlated)
- RF with  $m = \sqrt{p}$ :  $\rho \approx 0.05\text{--}0.2$
- Smaller  $m \Rightarrow$  lower  $\rho$ , but higher per-tree variance  $\sigma^2$

The sweet spot:

- As  $\rho \rightarrow 0$ :  $\text{Var} \rightarrow \sigma^2/B$  (ideal)
- Diminishing returns beyond  $B \approx 500$  trees
- The  $\rho\sigma^2$  floor is the irreducible ensemble variance

Method	$\rho$	$B$	Var
Single tree	1.0	1	$\sigma^2$
Bagging	0.6	500	$0.601\sigma^2$
RF ( $\sqrt{p}$ )	0.1	500	$0.102\sigma^2$
RF ( $\sqrt{p}$ )	0.1	1000	$0.101\sigma^2$

Note: going from 500 to 1000 trees barely helps once  $\rho$  dominates. The real gain comes from reducing  $\rho$ .

## Insight

Decorrelation (reducing  $\rho$ ) matters far more than adding trees (increasing  $B$ ) – this is RF's core innovation.

In practice, tune `max_features` before `n_estimators`.

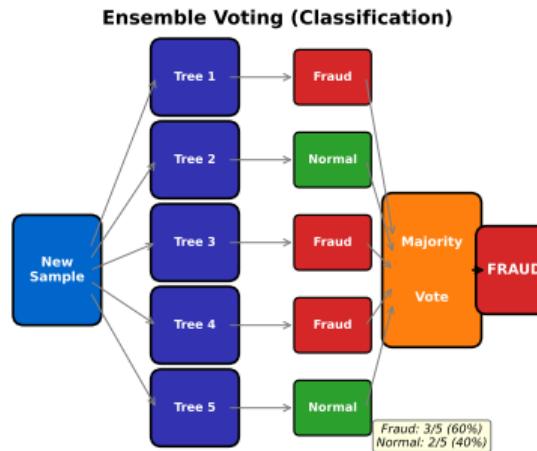
# How Do 500 Trees Vote on a Single Prediction?

**What you see:** Individual tree predictions and the ensemble's aggregated vote for a single observation.

**Key pattern:** Individual trees disagree, but the **majority vote** converges to the correct class as  $B$  grows.

**Takeaway:**

- Classification: majority vote  $\hat{y} = \text{mode}(\hat{y}_1, \dots, \hat{y}_B)$
- Regression: average  $\hat{y} = \frac{1}{B} \sum_{b=1}^B \hat{y}_b$
- Probability: fraction voting for each class
- Confidence grows with ensemble size and tree diversity



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/04\\_Random\\_Forests/05\\_ensemble\\_voting](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/04_Random_Forests/05_ensemble_voting)

## Insight

The Condorcet Jury Theorem guarantees: if each tree is better than random and errors are independent, the ensemble converges to perfect accuracy.

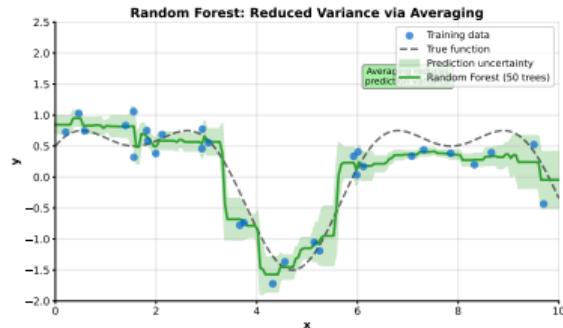

`predict_proba()` returns vote fractions – use these for ROC/AUC analysis.

# Can You See the Variance Drop from One Tree to a Forest?

**What you see:** A single tree's jagged boundary (top) vs. the forest's smooth boundary (bottom).

**Key pattern:** The single tree overfits local noise; the forest averages it out.

**Takeaway:** Variance reduction made visible – averaging decorrelated trees smooths the prediction surface.



## Insight

The forest does not learn a "better" tree – it averages *many different* trees to cancel out individual errors.

**Top: single tree (high variance). Bottom: 500-tree forest (low variance, similar bias).**

# What Is the Random Forest Algorithm in Pseudocode?

**Require:** Training data  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$ , number of trees  $B$ , feature subset size  $m$

```
1: for  $b = 1$  to  $B$  do
2:   Draw bootstrap sample  $\mathcal{D}_b$  of size  $n$ 
3:   Grow tree  $T_b$  on  $\mathcal{D}_b$ :
4:     At each node, select  $m$  features at random
5:     Find best split among these  $m$  features
6:     Split and recurse until stopping criterion
7:   Store  $T_b$ 
8: end for
```

**Ensure:** Predict new  $\mathbf{x}$ :

9: Classification:  $\hat{y} = \text{majority vote of } \{T_b(\mathbf{x})\}_{b=1}^B$

10: Regression:  $\hat{y} = \frac{1}{B} \sum_{b=1}^B T_b(\mathbf{x})$

## Complexity:

- Training:  $O(B \cdot n \log n \cdot m)$
- Prediction:  $O(B \cdot \text{depth})$
- Embarrassingly parallel – each tree independent

## Default hyperparameters:

- $B = 500$  (`n_estimators`)
- $m = \lfloor \sqrt{p} \rfloor$  (`max_features`)
- No `max_depth` (fully grown)
- `min_samples_leaf = 1`

## Insight

The algorithm is trivially parallelizable – training time scales linearly with cores, making RF fast even on large datasets.

**scikit-learn:** `RandomForestClassifier(n_jobs=-1)` uses all available CPU cores.


# Which Features Matter Most? (Mean Decrease in Impurity)

**What you see:** Feature importance ranked by Mean Decrease in Impurity (MDI) across all trees.

**Key pattern:** MDI sums the Gini reduction from each feature across all splits in all trees:

$$\text{MDI}_j = \frac{1}{B} \sum_{b=1}^B \sum_{t \in T_b} \Delta G_{t,j}$$

**Takeaway:** MDI is fast and built-in (`feature_importances_` in scikit-learn), but it has known biases.



## Insight

MDI favors features used in many splits – not necessarily the features most predictive of the outcome.

Always cross-check MDI with permutation importance or SHAP for reliable conclusions.

# Why Can MDI Be Misleading, and What Is Permutation Importance?

## MDI bias:

- Favors **high-cardinality** features (many unique values  $\Rightarrow$  more split opportunities)
- Inflated for correlated features (importance is "shared")
- Computed on training data – can reflect overfitting

## Permutation importance (model-agnostic):

1. Compute baseline score on held-out data
2. Shuffle feature  $j$ 's values randomly
3. Re-score; importance = drop in performance
4. Repeat for all features

	MDI	Permutation
Data used	Train	Test/OOB
Bias	High-card.	None
Cost	Free	$O(p \cdot B)$
Model	RF only	Any model
Correlations	Inflated	Shared

scikit-learn: `permutation_importance(model, X_test, y_test)`

## Insight

Permutation importance measures what the model *actually uses*; MDI measures what the model *could use*.

Strobl et al. (2007) first demonstrated MDI bias; always prefer permutation importance for final reporting.

# How Do SHAP Values Explain Individual Predictions?

Shapley value from cooperative game theory:

$$\phi_j = \sum_{S \subseteq N \setminus \{j\}} \frac{|S|! (p - |S| - 1)!}{p!} [f(S \cup \{j\}) - f(S)]$$

- $\phi_j$  = marginal contribution of feature  $j$ , averaged over all coalitions  $S$
- Satisfies efficiency, symmetry, linearity, and null-player axioms
- $\sum_j \phi_j = f(x) - \mathbb{E}[f(X)]$  (predictions decompose exactly)

## Insight

SHAP is the only feature attribution method with a *unique* solution satisfying all fairness axioms from game theory.

TreeSHAP (Lundberg et al., 2020):

- Exact Shapley values in  $O(TLD^2)$  – polynomial, not exponential
- Exploits tree structure to enumerate coalitions efficiently
- Local (per-prediction) and global (averaged) explanations

Example: Transaction flagged as fraud

$\phi_{amount} = +0.32$

$\phi_{location} = +0.18$

$\phi_{time} = -0.05$

⇒ Amount and location drove the flag.

Lundberg & Lee (2017): “A Unified Approach to Interpreting Model Predictions.” NeurIPS.


# How Does Boosting Learn from Its Mistakes?

**Core idea:** Train weak learners **sequentially**, each focusing on the errors of the previous ensemble.

- Bagging (RF): parallel, reduces variance, full trees
- **Boosting**: sequential, reduces bias, shallow trees (stumps)
- Each new learner “corrects” the residual errors
- Final prediction: weighted sum of all weak learners

Three major algorithms:

- AdaBoost – reweights misclassified samples
- Gradient Boosting – fits residuals of a loss function
- XGBoost – regularized gradient boosting



Sequential: each learner corrects the combined errors of all predecessors.

## Insight

Boosting reduces *bias* (makes weak learners strong); bagging reduces *variance* (stabilizes strong learners).

Schapire (1990) proved that weak learners can be “boosted” to arbitrary accuracy.

# What Is AdaBoost's Weight Update Rule?

**Learner weight** (how much to trust learner  $t$ ):

$$\alpha_t = \frac{1}{2} \ln \frac{1 - \epsilon_t}{\epsilon_t}$$

where  $\epsilon_t$  = weighted error rate of  $h_t$ .

**Sample weight update** (Freund & Schapire notation):

$$w_i \leftarrow w_i \cdot \exp(-\alpha_t y_i h_t(x_i))$$

where  $y_i \in \{-1, +1\}$  and  $h_t(x_i)$  is the weak learner prediction.

- When  $y_i \neq h_t(x_i)$ , the exponent is positive, increasing the weight
- Correctly classified: weight **decreased**; misclassified: weight **increased**
- Next learner focuses on hard examples
- Normalize weights to sum to 1

## Insight

AdaBoost's exponential loss makes it sensitive to outliers and label noise – a key limitation in noisy financial data.

**Properties of  $\alpha_t$ :**

- $\epsilon_t = 0$  (perfect):  $\alpha_t \rightarrow \infty$
- $\epsilon_t = 0.5$  (random):  $\alpha_t = 0$  (ignored)
- $\epsilon_t > 0.5$ :  $\alpha_t < 0$  (inverted)

**Final prediction:**

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

AdaBoost is equivalent to forward stagewise additive modeling with exponential loss.

**Freund & Schapire (1997): "A Decision-Theoretic Generalization of On-Line Learning."** JCSS.

# How Does Gradient Boosting Minimize an Arbitrary Loss?

## Additive update:

$$f_m(x) = f_{m-1}(x) + \eta \cdot h_m(x)$$

where  $h_m$  fits the **negative gradient** of the loss:

$$r_{im} = -\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \Big|_{f=f_{m-1}}$$

**Squared loss:**  $r_{im} = y_i - f_{m-1}(x_i)$  (residuals)

**Log loss:**  $r_{im} = y_i - \sigma(f_{m-1}(x_i))$

Learning rate  $\eta \in (0, 1]$  shrinks each tree's contribution for regularization.

## Algorithm:

1. Initialize  $f_0(x) = \arg \min_{\gamma} \sum L(y_i, \gamma)$
2. For  $m = 1, \dots, M$ :
  - Compute pseudo-residuals  $r_{im}$
  - Fit tree  $h_m$  to  $\{(x_i, r_{im})\}$
  - Update:  $f_m = f_{m-1} + \eta \cdot h_m$
3. Output  $f_M(x)$

## Key hyperparameters:

- $M$ : number of boosting rounds
- $\eta$ : learning rate (0.01–0.3)
- Tree depth: usually 3–8

## Insight

Gradient boosting performs *gradient descent in function space* – each tree is a step toward the loss minimum.

Friedman (2001): “Greedy Function Approximation: A Gradient Boosting Machine.” *Annals of Statistics*.

# What Makes XGBoost's Objective Function Special?

Regularized objective:

$$\mathcal{L} = \sum_{i=1}^n L(y_i, \hat{y}_i) + \sum_{k=1}^K \Omega(f_k)$$

where the regularization term is:

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$

- $T$  = number of leaves,  $w$  = leaf weights
- $\gamma$  penalizes tree complexity (pruning)
- $\lambda$  penalizes large leaf values (L2)

Second-order Taylor expansion:

$$\mathcal{L}^{(t)} \approx \sum_i [g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i)] + \Omega(f_t)$$

where  $g_i = \partial_{\hat{y}} L$ ,  $h_i = \partial_{\hat{y}}^2 L$ .

Optimal leaf weight:

$$w_j^* = -\frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda}$$

Using second-order info gives faster convergence than first-order gradient boosting.

## Insight

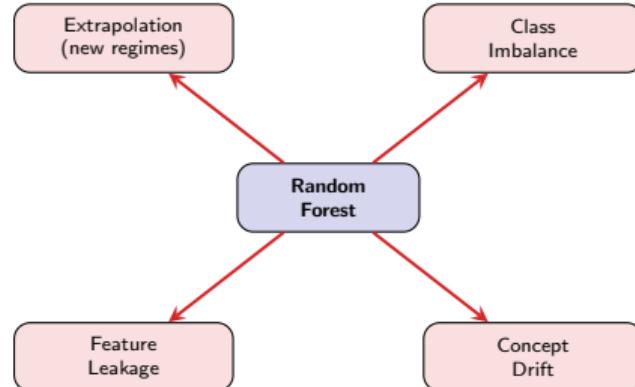
XGBoost's regularization prevents overfitting *structurally* (fewer leaves) and *numerically* (smaller weights) – crucial for noisy financial data.

Chen & Guestrin (2016): "XGBoost: A Scalable Tree Boosting System." KDD.

# When Does a Forest Fail Silently?

Four failure modes that produce **no error message** but wrong predictions:

- **Extrapolation:** RF cannot predict outside training range (piecewise constant). New market regimes break the model.
- **Class imbalance:** 99.9% legitimate  $\Rightarrow$  “always predict legit” achieves 99.9% accuracy but catches zero fraud.
- **Feature leakage:** Future information in training features (e.g., post-transaction flags) inflates apparent accuracy.
- **Concept drift:** Fraud patterns evolve; a model trained in 2023 degrades by 2024.



All four produce high accuracy on historical data but fail in production.

## Insight

Always evaluate on *temporally out-of-sample* data and monitor model performance continuously in production.

Silent failures are worse than crashes – they erode trust slowly and cause regulatory risk.

# How Do You Choose the Right Hyperparameters?

Param	Default	Guideline
n_estimators	100	500+; more is rarely worse
max_depth	None	Limit to 10–20 for speed
min_samples_split	2	Increase to reduce overfit
max_features	$\sqrt{p}$	Try $\log_2 p$ for large $p$
class_weight	None	'balanced' for imbalance

- **Tuning order:** max\_features → max\_depth → n\_estimators
- Use OOB score for quick evaluation
- Use cross-validation for final selection

## Overfitting signals:

- Training accuracy  $\gg$  test accuracy
- OOB error diverges from training error
- Feature importance dominated by noise features

## Underfitting signals:

- Both training and test accuracy low
- Increase max\_depth or n\_estimators
- Check if max\_features is too small

**Rule of thumb:** Start with defaults, increase n\_estimators to 500, then tune max\_features via grid search.

## Insight

Random Forests are remarkably robust to hyperparameters – defaults often perform within 1–2% of the optimum.

Probst et al. (2019): "Hyperparameters and Tuning Strategies for Random Forest." WIREs Data Mining.

# How Do Fraud Teams Use Feature Importance for Transaction Monitoring?

## Worked case study:

**Features:** transaction amount, merchant category, time of day, customer age, distance from home, transaction frequency (last 24h), card-present flag.

**Model:** RF with 500 trees, class\_weight='balanced', trained on 6 months of labeled data.

## Top-5 by permutation importance:

1. Distance from home ( $\Delta\text{AUC} = 0.08$ )
2. Transaction amount ( $\Delta\text{AUC} = 0.06$ )
3. Frequency last 24h ( $\Delta\text{AUC} = 0.04$ )
4. Time of day ( $\Delta\text{AUC} = 0.03$ )
5. Card-present flag ( $\Delta\text{AUC} = 0.02$ )

## Actionable insights:

- **Distance:** Transactions far from home are the strongest fraud signal
- **Amount:** Large transactions need extra scrutiny, but amount alone is insufficient
- **Frequency:** Rapid-fire transactions suggest card testing
- **Time:** Late-night transactions correlate with fraud
- **Card-present:** CNP (card-not-present) transactions are riskier

These findings are presented to fraud analysts as SHAP waterfall plots for each flagged transaction.

## Insight

Feature importance tells analysts *where to look*; SHAP tells them *why this specific transaction was flagged*.

---

Fraud teams combine model outputs with rule-based systems for final decisions.

# What Happens When 99.9% of Transactions Are Legitimate?

## The class imbalance problem:

- Fraud rate  $\approx 0.1\%$   $\Rightarrow$  1 fraud per 1000 transactions
- Classifier predicting “always legit” gets 99.9% accuracy
- **Accuracy is meaningless** for imbalanced data

## Solutions:

- `class_weight='balanced'`: upweight minority class inversely proportional to frequency
- **Stratified bootstrap**: ensure each tree's sample has proportional fraud cases
- **SMOTE**: synthetic minority oversampling (creates new fraud examples by interpolation)
- **Cost-sensitive learning**:  $C_{FN} \gg C_{FP}$  (missing fraud costs more than false alarms)

## Insight

In fraud detection, a false negative (missed fraud) costs 10–100x more than a false positive (unnecessary alert).

## Precision-recall tradeoff:

- **Precision**: Of flagged transactions, how many are actually fraud?
- **Recall**: Of all fraud, how many did we catch?
- Lowering threshold  $\Rightarrow$  higher recall, lower precision
- Banks target **high recall** (catch fraud) at acceptable precision

## Key metrics for imbalanced data:

- PR-AUC (precision-recall curve area)
- F1 or  $F_\beta$  score (weighted harmonic mean)
- Cost-weighted accuracy

Use PR-AUC, not ROC-AUC, when the positive class is rare (Davis & Goadrich, 2006).

# When Should You Choose RF Over Logistic Regression?


**What you see:** A decision flowchart guiding model selection based on data characteristics and requirements.

**Key pattern:**

- **Interpretability required?** ⇒ Logistic regression or single tree
- **Nonlinear relationships?** ⇒ RF or boosting
- **Many features, few samples?** ⇒ RF (implicit feature selection)
- **Maximum accuracy needed?** ⇒ XGBoost with tuning

**Takeaway:** No single model dominates; the right choice depends on the problem, data, and regulatory constraints.

## When to Use Random Forests



### Key:

Random Forest: Best for accuracy + feature importance  
Alternative: When interpretability is paramount

[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04\\_Random\\_Forests/07\\_decision\\_flowchart](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L04_Random_Forests/07_decision_flowchart)


## Insight

In regulated industries, the “best” model is often the one that is *explainable enough* to satisfy compliance – not the most accurate.

Many banks use RF internally but report logistic regression coefficients to regulators.

# Who Wins and Who Loses When Ensembles Replace Scorecards?

- **Fraud Team:** Wins – better detection rates, fewer missed cases, faster adaptation
- **Risk Manager:** Wins – lower losses, better capital allocation, richer risk models
- **Compliance:** Mixed – improved outcomes but harder to audit 500-tree models
- **Customer:** Mostly wins – fewer false declines, but less transparent decisions
- **Regulator:** Concerned – demands explainability, adverse action reasons, fairness audits



## Insight

The transition from scorecards to ensembles is a *governance challenge* as much as a technical one.

Successful adoption requires buy-in from all five stakeholders, not just data scientists.

# Can Regulators Trust a Model Made of 500 Trees?

## Regulatory landscape:

- **ECOA / Reg B (US):** Lenders must provide specific “adverse action reasons” for denial – hard with a forest
- **GDPR Article 22 (EU):** Right to “meaningful information about the logic involved” in automated decisions
- **EU AI Act:** High-risk AI systems (credit scoring) require transparency and human oversight

## Bridging the gap with SHAP:

- SHAP provides per-decision explanations
- Top- $k$  SHAP features  $\Rightarrow$  adverse action reasons
- Global SHAP summary  $\Rightarrow$  model documentation

## Insight

SHAP turns a “black box” into a “glass box” – the model is complex, but each decision is explainable.

## Practical compliance workflow:

1. Train RF for optimal performance
2. Compute SHAP values for every prediction
3. Map top SHAP features to human-readable reason codes
4. Document model in a “model risk management” framework
5. Monitor for fairness across protected classes

SR 11-7 (Fed) requires banks to validate and document all models used in decision-making.

The EU AI Act (2024) classifies credit scoring as “high-risk” – requiring explainability by law.

# Key Takeaways

## Mathematical Foundation

- Gini impurity and entropy guide tree splits; both yield similar results
- Variance of ensemble:  $\rho\sigma^2 + \frac{(1-\rho)}{B}\sigma^2$  – decorrelation is key
- OOB error provides free, unbiased validation

## Ensemble Methods Compared

- **Bagging/RF:** reduces variance via parallel, decorrelated trees
- **AdaBoost:** reduces bias via sequential reweighting
- **Gradient Boosting:** gradient descent in function space
- **XGBoost:** adds regularization and second-order optimization

## Practical Application

- Feature importance: MDI → permutation → SHAP (increasing reliability)
- Fraud detection demands high recall, cost-sensitive learning, and PR-AUC evaluation
- Class imbalance: class\_weight, SMOTE, stratified sampling
- SHAP bridges the gap between model complexity and regulatory explainability

---

Random Forests remain one of the strongest “out-of-the-box” classifiers two decades after Breiman (2001).


We opened with the absurdity of ensemble models (XKCD #1885).

Now you know **WHY** 500 trees beat one: **variance reduction through decorrelation**.

- One tree is interpretable but fragile
- Bagging stabilizes via bootstrap averaging
- Feature randomization decorrelates trees ( $\rho \downarrow$ )
- Boosting attacks bias where RF attacks variance
- SHAP makes the forest explainable for regulators

## Next: L05 – PCA & t-SNE

From ensemble predictions to **dimensionality reduction** – how to visualize and compress high-dimensional financial data.



*"The real power of ensemble models is not that each tree is good – it is that each tree is wrong in a different way."*

---

L05 preview: PCA finds directions of maximum variance; t-SNE preserves local structure for visualization.