

# L05: PCA & t-SNE

## Deep Dive: Theory, Derivations, and Financial Applications

### Methods and Algorithms

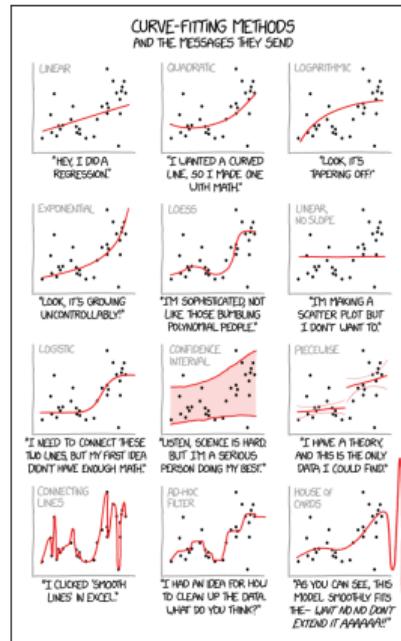
MSc Data Science

Spring 2026

# Outline

- 1 PCA Foundations
- 2 PCA in Finance
- 3 t-SNE Theory
- 4 PCA vs t-SNE Comparison
- 5 Finance Applications
- 6 Implementation
- 7 Practice
- 8 Summary

# Finding Signal in the Noise



XKCD #2048 “Curve Fitting” by Randall Munroe (CC BY-NC 2.5)

**By the end of this deep dive, you will be able to:**

1. Derive PCA from variance maximization via Lagrangian optimization and prove the SVD–PCA equivalence theorem
2. Evaluate dimensionality reduction methods using quantitative criteria (variance explained, reconstruction error, trustworthiness)
3. Analyze the t-SNE gradient, crowding problem, and perplexity–entropy relationship
4. Critique PCA limitations for nonlinear data and apply yield curve decomposition in fixed-income risk management

**Finance Applications:**

- Yield curve PCA (level/slope/curvature)
- Portfolio risk factor decomposition
- Market regime detection via t-SNE

---

Bloom's Levels 4–5: Analyze, Evaluate, Create

# PCA: Variance Maximization Objective

**Goal:** Find orthogonal directions of maximum variance in the data.

**Mean-Centering (Required First Step):**

$$X_c = X - \bar{X} \quad \text{where } \bar{X}_j = \frac{1}{n} \sum_{i=1}^n X_{ij}$$

**Key Properties of PCA:**

- **Linear:** Components are linear combinations of original features
- **Uncorrelated:** Principal components have zero cross-correlation
- **Partially reversible:** Lossy reconstruction when  $k < p$  (discards low-variance directions)

**PCA does NOT require Gaussian data** — it is optimal for Gaussian distributions, but valid for any distribution with finite second moments.

---

Centering is not optional: PCA on uncentered data maximizes distance from origin, not variance.

**Covariance Matrix** (from centered data):

$$\Sigma = \frac{1}{n-1} X_c^T X_c \quad \text{where } X_c \in \mathbb{R}^{n \times p}$$

**Eigendecomposition:**

$$\Sigma \mathbf{v}_k = \lambda_k \mathbf{v}_k, \quad k = 1, \dots, p$$

- $\mathbf{v}_k$  = eigenvector (principal direction)
- $\lambda_k$  = eigenvalue (variance along that direction),  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_p \geq 0$

**Projection to  $k$  dimensions:**

$$Z = X_c W_k \quad \text{where } W_k = [\mathbf{v}_1, \dots, \mathbf{v}_k] \in \mathbb{R}^{p \times k}$$

---

$\Sigma$  is symmetric positive semi-definite, guaranteeing real non-negative eigenvalues and orthogonal eigenvectors.

# Why Eigenvectors? The Optimality Proof

Find direction  $w$  maximizing projected variance:

$$\max_w w^T \Sigma w \quad \text{subject to } \|w\| = 1$$

Lagrangian:

$$\mathcal{L} = w^T \Sigma w - \lambda(w^T w - 1)$$

First-order condition:

$$\nabla_w \mathcal{L} = 2\Sigma w - 2\lambda w = 0 \implies \Sigma w = \lambda w$$

**Result:** The optimal  $w$  is an eigenvector of  $\Sigma$ .

Substituting back:  $w^T \Sigma w = w^T \lambda w = \lambda$ . Maximum variance is achieved at  $\lambda_1$  (largest eigenvalue), so PC1 = eigenvector for  $\lambda_1$ .

---

Constrained optimization proves eigenvectors are the unique optimal solution (up to sign).

**Theorem:** The principal components of  $X_c$  are the right singular vectors of  $X_c$ .

**Proof sketch:**

1. SVD of centered data:  $X_c = USV^T$  where  $S = \text{diag}(s_1, \dots, s_r)$
2. Form  $X_c^T X_c = VS^T U^T USV^T = VS^2 V^T$
3. Covariance:  $C = \frac{1}{n-1} X_c^T X_c = \frac{1}{n-1} VS^2 V^T$
4. Therefore  $CV = V \cdot \frac{S^2}{n-1}$  — columns of  $V$  are eigenvectors of  $C$

**Eigenvalues:**  $\lambda_k = s_k^2 / (n - 1)$  where  $s_k$  are singular values.

**Why SVD is preferred:** Avoids forming  $X_c^T X_c$ , which squares the condition number. SVD is numerically more stable for ill-conditioned data.

---

All practical implementations (scikit-learn, R, MATLAB) use SVD internally, not eigendecomposition.

## Explained Variance Ratio (EVR):

$$\text{EVR}_k = \frac{\lambda_k}{\sum_{j=1}^p \lambda_j}, \quad \text{Cumulative: } \text{CVR}_k = \sum_{i=1}^k \text{EVR}_i$$

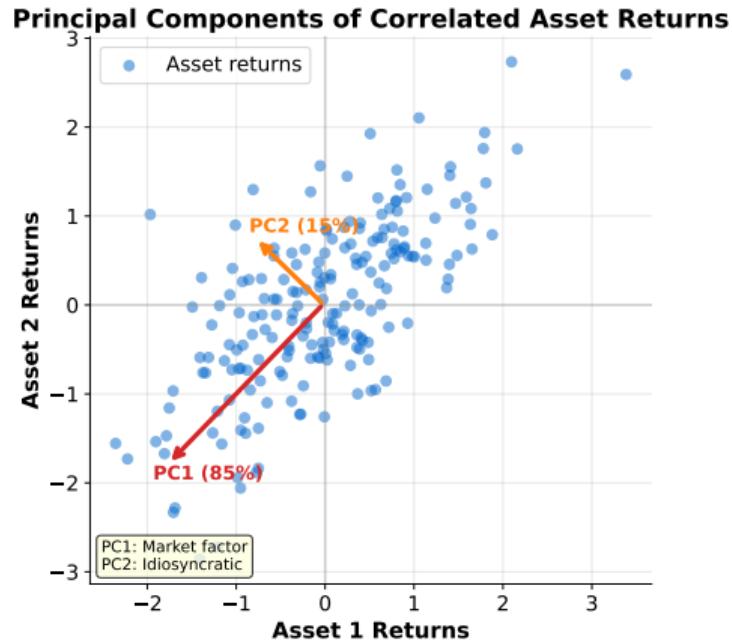
Rules for choosing  $k$ :

- **Variance threshold:** Keep 80–95% of total variance (domain-dependent)
- **Scree plot elbow:** Visual inspection for diminishing returns
- **Kaiser criterion:** Retain components with  $\lambda > 1$  — **valid only for correlation matrix** (standardized data)
- **Parallel analysis:** Compare eigenvalues against random data permutations (most rigorous)

---

In finance, 3 components explain 98%+ of yield curve variance; for equities, 5–10 may be needed.

# Principal Components Visualization



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L05\\_PCA\\_TSNE/02\\_principal\\_components](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L05_PCA_TSNE/02_principal_components)

**PC1 captures the dominant variance direction; PC2 captures orthogonal residual variation.**

# Reconstruction from Principal Components

From  $k$  components back to original space:

$$\hat{X} = ZW_k^T + \bar{X} \quad (\text{add the mean back for original scale})$$

Reconstruction Error (Frobenius norm):

$$\|X - \hat{X}\|_F^2 = \sum_{i=k+1}^p \lambda_i$$

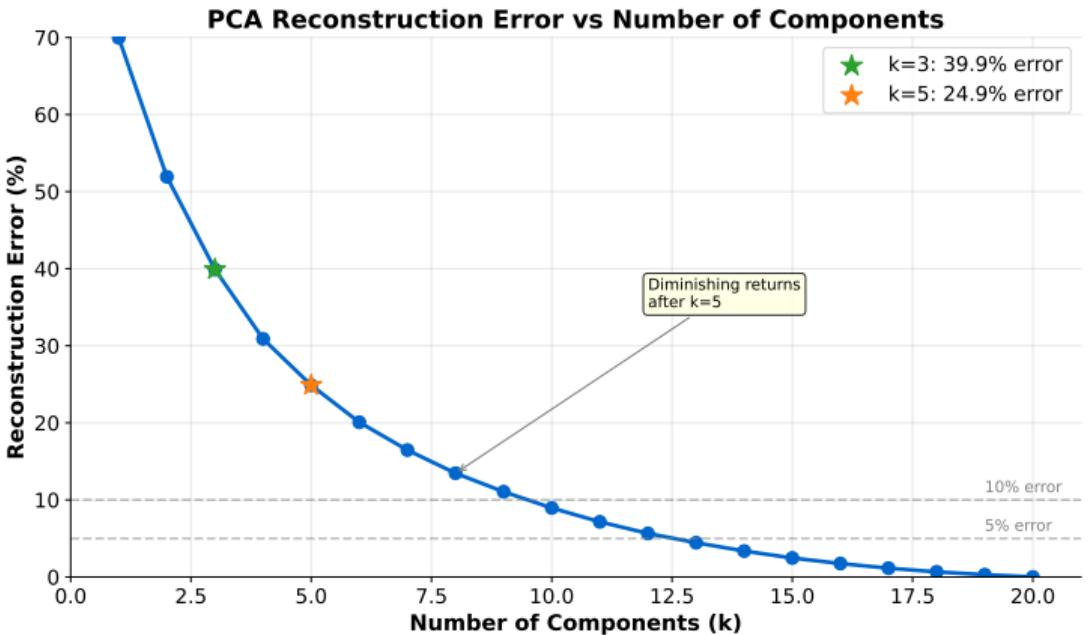
Key insights:

- Error equals the sum of *discarded* eigenvalues
- PCA gives the *optimal* rank- $k$  approximation (Eckart–Young theorem)
- More components = less error, but diminishing returns

---

The Eckart–Young theorem guarantees no other linear method can do better for a given  $k$ .

# Reconstruction Error vs. Number of Components



Adding more components always reduces error, but with diminishing marginal improvement.

# Yield Curve PCA: The Canonical Example

**Yield curves decompose into ~3 principal components:**

- **PC1 = Level** (parallel shift): explains ~85% of variance
- **PC2 = Slope** (steepening/flattening): explains ~10%
- **PC3 = Curvature** (butterfly): explains ~3%

**Together: 98%+ of all yield curve movements in 3 numbers.**

**Industry Applications:**

- Used daily in bank risk management systems worldwide
- Foundation for duration-neutral hedging strategies
- Basis for curve trades and butterfly spreads

---

**Litterman & Scheinkman (1991): the foundational paper establishing level/slope/curvature decomposition.**

# Yield Curve PCA: Worked Example

## Typical PC Loadings by Maturity:

Maturity	PC1 (Level)	PC2 (Slope)	PC3 (Curvature)
1Y	0.42	-0.58	0.50
2Y	0.44	-0.37	-0.20
5Y	0.46	0.06	-0.62
10Y	0.45	0.41	-0.10
30Y	0.43	0.60	0.56

### Interpretation:

- **PC1:** Near-uniform loadings — all rates move together (parallel shift)
- **PC2:** Negative short-end, positive long-end — curve steepens or flattens
- **PC3:** Positive at extremes, negative in middle — butterfly movement

---

Loadings from US Treasury yields; signs are conventional (eigenvectors are unique up to sign).

## Applying PCA to equity return covariance:

- **PC1 = Market factor:** Uniform loadings, explains ~60% of variance
- **PC2 = Sector rotation:** Growth vs. value or sector tilts
- **PC3+:** Increasingly idiosyncratic risk factors

## Example: 10-Stock Portfolio

- PC1 explains 62% of portfolio variance (broad market exposure)
- First 3 PCs explain 81% (market + two sector factors)
- Remaining 7 PCs: stock-specific noise

## Use Cases:

- Risk factor modeling and attribution
- Dimensionality reduction for trading signal extraction

---

PCA-based risk factors are model-free alternatives to explicit factor models (Fama–French).

## Bootstrap Confidence Intervals:

- Resample data with replacement, recompute PCA each time
- Track loading stability and eigenvalue confidence bands

## Parallel Analysis (most rigorous for $k$ ):

- Compare observed eigenvalues to eigenvalues from random data
- Keep components where  $\lambda_{\text{data}} > \lambda_{\text{random}}$

## Cross-Validation:

- Split data, fit PCA on training set, evaluate reconstruction on holdout

For t-SNE: Run multiple times with different random seeds; clusters that are unstable across runs are likely artifacts.

---

MSc-level: always quantify uncertainty in dimensionality reduction choices.

# PCA Limitations

## When PCA Falls Short:

- **Non-linear relationships:** Curved manifolds projected incorrectly
- **Cluster misalignment:** Maximum variance may not separate clusters
- **Outlier sensitivity:** A single outlier can rotate principal components

## Solutions and Alternatives:

- **Kernel PCA:** Implicit non-linear mapping via kernel trick
- **Robust PCA:** Decomposes data into low-rank + sparse (outlier-resistant)
- **t-SNE / UMAP:** Non-linear methods optimized for visualization

**PCA assumes linear structure — Gaussian NOT required**, but PCA is only provably optimal (in the maximum-variance sense) for Gaussian data.

---

Non-linear data demands non-linear methods — motivating the t-SNE section that follows.

## t-Distributed Stochastic Neighbor Embedding (van der Maaten & Hinton, 2008)

### Intuition:

- Convert pairwise distances to probabilities (“who is my neighbor?”)
- **High-D:** Use Gaussian kernel to define neighbor probabilities
- **Low-D:** Use Student's t-distribution (heavy tails)
- Minimize KL divergence between the two probability distributions

### Critical distinction:

- t-SNE is a **visualization method**, NOT a preprocessing step
- Output coordinates have no interpretable axes or distances
- Cannot embed new points without re-running on entire dataset

---

t-SNE reveals IF clusters exist; it does not tell you HOW clusters relate to each other.

**High-dimensional conditional similarity:**

$$p_{j|i} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\sigma_i^2)}$$

**Symmetrization** (joint probability):

$$p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$$

**Low-dimensional similarity** (Student's t-distribution, 1 d.f.):

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$

**Objective:** Minimize KL divergence  $\text{KL}(P\|Q) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$

---

Symmetrization ensures each point contributes equally regardless of local density.

## Why Student's t instead of Gaussian in low-D?

### The problem:

- In high-D, moderate distances are *common* (volume grows exponentially)
- Mapping to 2D with Gaussian: moderate distances collapse to small distances
- Result: all points crowd into center, clusters become indistinguishable

### The solution — heavy-tailed t-distribution:

- Heavy tails allow dissimilar points to be placed *far apart* in 2D
- Nearby points stay nearby (Gaussian and t agree at short distances)
- Far-away points get pushed apart (t-distribution decays slower)

This is the key innovation of t-SNE over the original SNE (Hinton & Roweis, 2002).

---

The crowding problem is fundamental to all high-to-low dimensional embeddings, not just t-SNE.

**KL Divergence** (asymmetric):

$$\text{KL}(P\|Q) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

- Penalizes large  $p_{ij}$  with small  $q_{ij}$  heavily (nearby  $\rightarrow$  must stay nearby)
- Tolerates large  $q_{ij}$  with small  $p_{ij}$  (distant  $\rightarrow$  can be anywhere)
- This asymmetry is why t-SNE preserves *local* structure

**Gradient** (the force on each embedding point):

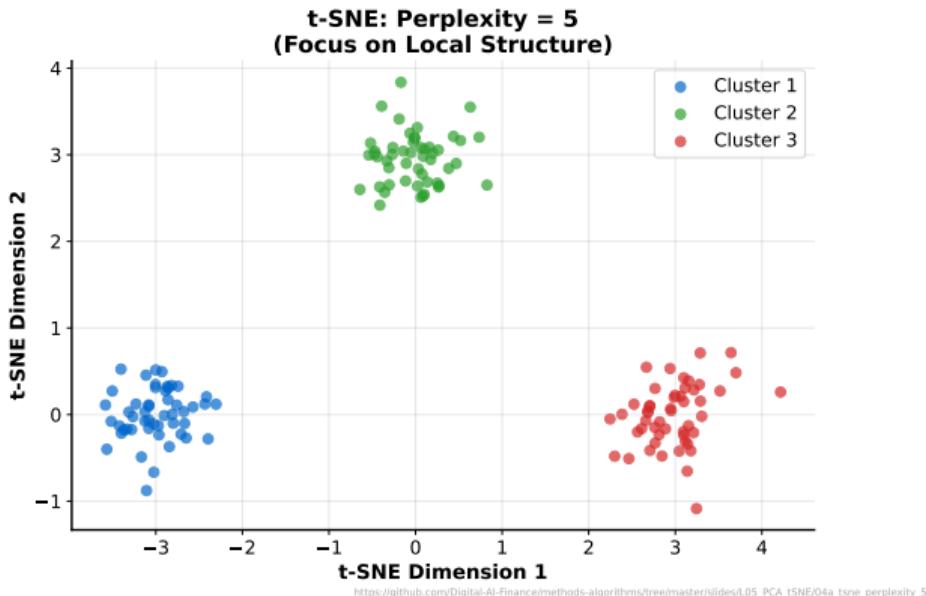
$$\frac{\partial \text{KL}}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)(1 + \|y_i - y_j\|^2)^{-1}$$

**Interpretation:** Attractive forces ( $p_{ij} > q_{ij}$ ) pull neighbors closer; repulsive forces ( $p_{ij} < q_{ij}$ ) push non-neighbors apart.

---

Optimized via gradient descent with momentum; early exaggeration amplifies attractive forces in initial iterations.

## Perplexity: Low (Local Focus)



**Perplexity = 5**

$$\text{Perp} = 2^{H(\sigma_i)}$$

Effectively considers  $\sim 5$  nearest neighbors per point.

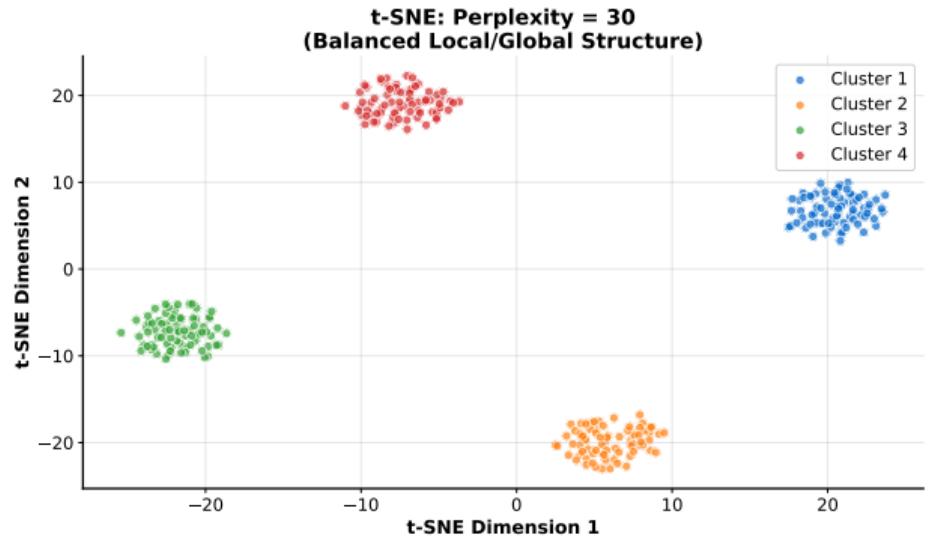
**Effect:**

- Very tight, fragmented clusters
- May split true clusters
- Captures fine local detail

---

**Low perplexity: good for detecting sub-structure within clusters, but risks over-fragmentation.**

## Perplexity: Default (Balanced)



**Perplexity = 30**

The standard default value for most implementations.

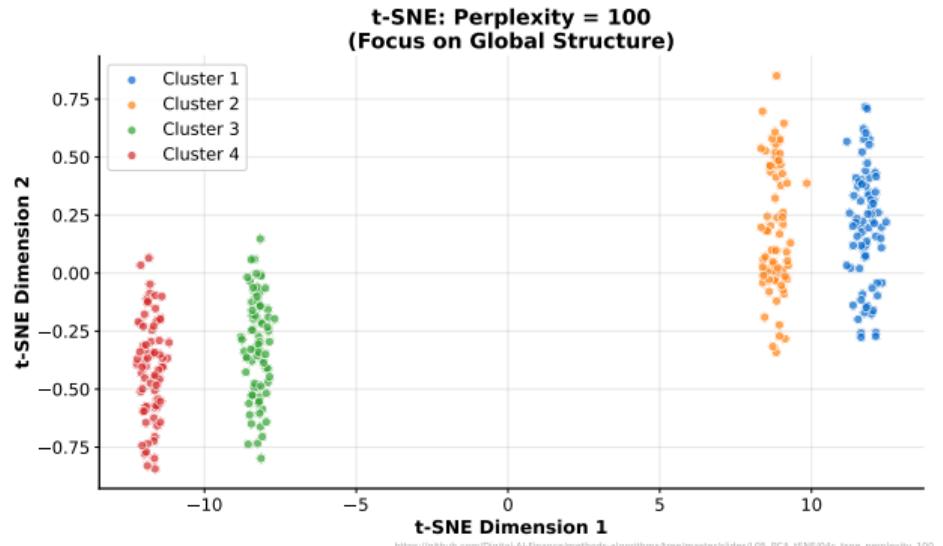
**Effect:**

- Balanced local and global structure
- Clusters are well-separated
- Robust starting point

---

Perplexity 30 is the recommended starting point; always compare with other values.

## Perplexity: High (Global Focus)



**Perplexity = 100**

Considers ~100 neighbors — broader context.

**Effect:**

- More spread-out embedding
- Clusters may merge
- Better global relationships

---

**High perplexity: approaches a more global view but may obscure fine cluster boundaries.**

# Perplexity Guidelines

**Perplexity** =  $2^H$  where  $H = -\sum_j p_{j|i} \log_2 p_{j|i}$  is entropy.

Intuitively: the effective number of neighbors each point “considers.”

**Practical ranges:**

- **Low (5–10):** Fine-grained local detail; risk of fragmentation
- **Medium (30–50):** Balanced; recommended default range
- **High (100+):** Global structure emphasis; clusters may blend

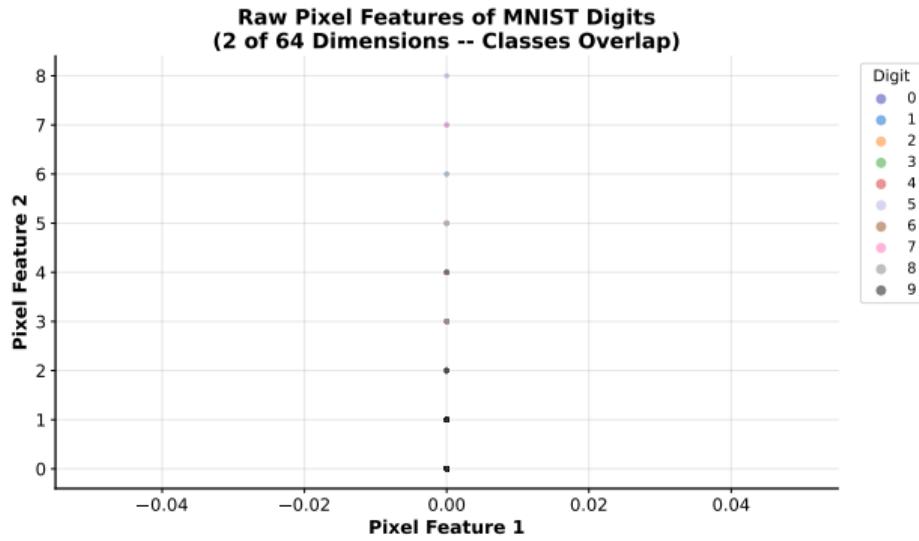
**Hard constraints and guidelines:**

- Must be  $< n$  (number of data points)
- Larger datasets tolerate higher perplexity
- **Always run multiple perplexities** to validate findings

---

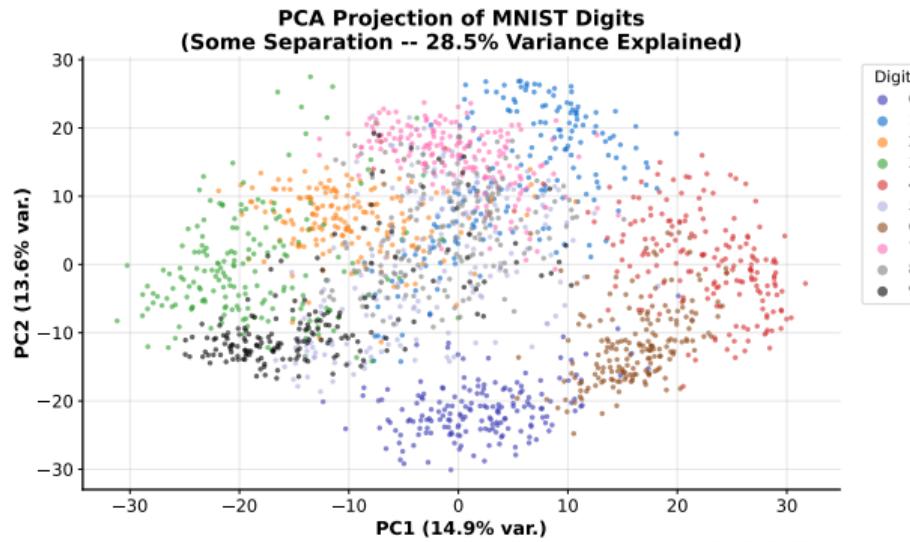
If results change dramatically with perplexity, the clusters may not be robust.

## Cluster Preservation: Original Data



MNIST digits (64 dimensions): classes overlap severely when viewed in raw pixel space.

## Cluster Preservation: PCA Projection



PCA finds max-variance directions — provides partial digit separation but clusters overlap.

## Important Limitations:

- **Non-deterministic:** Different random seeds produce different layouts
- **Complexity:**  $O(n^2)$  naive; Barnes-Hut approximation gives  $O(n \log n)$
- **No out-of-sample:** Cannot embed new points without full re-computation
- **Hyperparameter sensitive:** Perplexity, learning rate, iterations all matter

## Best Practices:

- **PCA first:** Reduce to 30–50 dimensions before t-SNE (faster, denoises)
- **Multiple seeds:** Run 3–5 times, keep structures that persist
- **Don't over-interpret:** Cluster sizes and inter-cluster distances are meaningless

---

t-SNE shows IF clusters exist; cluster sizes, gaps, and densities are artifacts of the algorithm.

## Uniform Manifold Approximation and Projection (McInnes et al., 2018)

### Advantages over t-SNE:

- **Faster:**  $O(n)$  after nearest-neighbor graph construction
- **Better global structure:** Preserves more large-scale relationships
- **Out-of-sample embedding:** Can map new points via learned transform

### Key Hyperparameters:

- `n_neighbors`: Similar role to perplexity (local vs. global balance)
- `min_dist`: Controls how tightly points cluster (0.0 = tight, 1.0 = spread)

Based on Riemannian geometry and algebraic topology (fuzzy simplicial sets).

---

UMAP is increasingly the default choice for visualization in production ML pipelines.

# When to Use Which Method

## Use PCA When:

- Preprocessing features for downstream ML models
- Linear relationships are expected or sufficient
- Reversibility needed (reconstruction, denoising)
- Speed matters (real-time, large datasets)

## Use t-SNE / UMAP When:

- Visualizing high-dimensional data in 2D/3D
- Exploring cluster structure and local neighborhoods
- Non-linear manifolds expected in the data
- Exploratory data analysis (not final inference)

**Common Pipeline:** Standardize → PCA (30–50 dims) → t-SNE/UMAP (2D)

---

PCA and t-SNE are complementary, not competitors — use them together in a pipeline.

## Pipeline:

1. Compute rolling features: volatility, correlations, returns (window = 20–60 days)
2. Standardize features (zero mean, unit variance)
3. PCA to 15–20 dimensions (remove noise)
4. t-SNE to 2D for visualization

## Typical regimes discovered:

- **Calm:** Low volatility, moderate correlations
- **Volatile:** Elevated volatility, sector dispersion
- **Crisis:** High volatility, high correlation (“all correlations go to 1”)

**Validation required:** Cluster labels should align with known market events (2008, 2020).

---

Regime detection is exploratory — clusters must be validated against economic fundamentals.

# PCA Preprocessing for ML Pipelines

## Standard Pipeline:

StandardScaler → PCA(`n_components=0.95`) → Classifier

## Benefits:

- Reduces multicollinearity (orthogonal features)
- Removes noise in low-variance components
- Speeds up training for high-dimensional data

## Cautions:

- PCA is unsupervised — may discard features that are low-variance but highly predictive
- Always compare ML performance with and without PCA preprocessing
- Consider supervised alternatives: LDA, feature selection

---

PCA preprocessing is a bias–variance tradeoff: reduces overfitting but may lose signal.

# Dimensionality Reduction: Method Comparison

Method	Type	Speed	New Points?	Best For
PCA	Linear	Fast	Yes	Preprocessing, denoising
Kernel PCA	Non-linear	Medium	Approximate	Non-linear structure
t-SNE	Non-linear	Slow	No	2D visualization
UMAP	Non-linear	Fast	Yes	Visualization + embedding
Autoencoder	Non-linear	Slow	Yes	Complex non-linear features

## Decision heuristic:

- Need reversibility or preprocessing? → PCA
- Need visualization? → UMAP (or t-SNE)
- Need non-linear feature extraction? → Autoencoder

---

No single method dominates — choice depends on goal (preprocessing vs. visualization vs. feature learning).

## Core API:

- `PCA(n_components=k)` — keep  $k$  components (integer)
- `PCA(n_components=0.95)` — keep 95% variance (float)
- `pca.explained_variance_ratio_` — variance per component
- `pca.inverse_transform(Z)` — reconstruct original space

## Implementation details:

- Uses randomized truncated SVD internally (not eigendecomposition)
- Automatically centers the data (subtracts mean)
- For sparse data: use TruncatedSVD instead (no centering)

**Key practice:** Always `StandardScaler().fit_transform(X)` before PCA when features have different units or scales.

---

PCA on covariance matrix (raw data) vs. correlation matrix (standardized) gives different results.

## t-SNE (`sklearn.manifold.TSNE`):

- `TSNE(n_components=2, perplexity=30, random_state=42)`
- Key params: `perplexity`, `learning_rate`, `n_iter`
- Set `random_state` for reproducibility; run multiple seeds to validate

## UMAP (`umap-learn` package):

- `UMAP(n_components=2, n_neighbors=15, min_dist=0.1)`
- Key params: `n_neighbors` ( $\approx$  perplexity), `min_dist`
- Supports `.transform()` for new data points

**Best practice pipeline:** PCA to 30–50 dims first, then t-SNE/UMAP to 2D. This improves speed and removes noise.

---

Install UMAP via: `pip install umap-learn` (not `umap`, which is a different package).

## Exercise 1: Yield Curve Decomposition

- Apply PCA to historical yield curve data (5 maturities)
- Interpret PC1/PC2/PC3 loadings as level/slope/curvature
- Compute cumulative variance explained

## Exercise 2: PCA vs. t-SNE on Digits

- Compare PCA (2D) and t-SNE (2D) projections on MNIST digits
- Evaluate cluster separation qualitatively and via silhouette score

## Exercise 3: Perplexity Sensitivity

- Run t-SNE with perplexity  $\in \{5, 15, 30, 50, 100\}$
- Identify which clusters are robust across perplexities

---

See course notebook: `L05_pca_tsne.ipynb` for starter code and datasets.

## Key Takeaways

### PCA:

- Maximizes variance via eigendecomposition / SVD
- Optimal linear dimensionality reduction (Eckart–Young)
- Finance: yield curve = level + slope + curvature (3 PCs, 98%+ variance)

### t-SNE:

- Neighbor embedding with KL divergence objective
- Crowding problem solved by t-distribution heavy tails
- Visualization only — do not use for preprocessing or inference

**Pipeline:** Standardize → PCA (preprocessing) → t-SNE/UMAP (visualization)

**Finance:** Yield curve PCA, portfolio risk factors, market regime detection

---

Dimensionality reduction is both a standalone tool and an essential preprocessing step.

*“After reducing 200 dimensions to 3,  
the risk manager asked:  
‘Which 197 did we lose?’*

*Answer: ‘The ones that were just noise...  
statistically speaking.’”*

— Inspired by XKCD #2400 “Statistics” by Randall Munroe

---

XKCD #2400 by Randall Munroe (CC BY-NC 2.5). The real question: how do you know it was noise?

## References: Textbooks and Papers

### Textbooks:

- Jolliffe, I.T. (2002). *Principal Component Analysis*, 2nd ed. Springer.
- James et al. (2021). *ISLR*, Chapter 12: Unsupervised Learning.
- Hastie et al. (2009). *ESL*, Chapter 14: Unsupervised Learning.

### Foundational Papers:

- Pearson, K. (1901). On Lines and Planes of Closest Fit to Systems of Points in Space.
- van der Maaten, L. & Hinton, G. (2008). Visualizing Data using t-SNE. *JMLR*, 9, 2579–2605.
- McInnes, L. et al. (2018). UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction.

---

The van der Maaten & Hinton (2008) paper has 40,000+ citations — one of the most cited ML papers.

# References: Documentation and Finance

## Software Documentation:

- scikit-learn: `sklearn.decomposition.PCA`
- scikit-learn: `sklearn.manifold.TSNE`
- UMAP documentation: <https://umap-learn.readthedocs.io/>

## Finance Applications:

- Litterman, R. & Scheinkman, J. (1991). Common Factors Affecting Bond Returns. *Journal of Fixed Income*, 1(1), 54–61.

## Interactive Resources:

- Wattenberg, M. et al. (2016). How to Use t-SNE Effectively. *Distill*. <https://distill.pub/2016/misread-tsne/>

---

The Distill article is essential reading for understanding t-SNE pitfalls and best practices.

# Appendix

## Advanced Topics and Proofs

Supplementary material for self-study and reference

---

**Appendix slides are not covered in lecture — provided for advanced students and exam preparation.**

## Full Eigenvalue Derivation

**PC1:** Solve  $\max_{\mathbf{w}_1} \mathbf{w}_1^T \Sigma \mathbf{w}_1$  s.t.  $\|\mathbf{w}_1\| = 1$ .

Lagrangian gives  $\Sigma \mathbf{w}_1 = \lambda_1 \mathbf{w}_1$ . Solution:  $\mathbf{w}_1$  = eigenvector for  $\lambda_1$  (largest).

**PC2:** Solve  $\max_{\mathbf{w}_2} \mathbf{w}_2^T \Sigma \mathbf{w}_2$  s.t.  $\|\mathbf{w}_2\| = 1$  and  $\mathbf{w}_2^T \mathbf{w}_1 = 0$ .

Lagrangian:  $\mathcal{L} = \mathbf{w}_2^T \Sigma \mathbf{w}_2 - \lambda(\mathbf{w}_2^T \mathbf{w}_2 - 1) - \mu(\mathbf{w}_2^T \mathbf{w}_1)$

First-order:  $2\Sigma \mathbf{w}_2 - 2\lambda \mathbf{w}_2 - \mu \mathbf{w}_1 = 0$

Multiply by  $\mathbf{w}_1^T$ :  $\mu = 2\mathbf{w}_1^T \Sigma \mathbf{w}_2 = 0$  (since  $\Sigma$  symmetric and  $\mathbf{w}_1 \perp \mathbf{w}_2$ ).

Thus  $\Sigma \mathbf{w}_2 = \lambda_2 \mathbf{w}_2$  with  $\lambda_2$  = second largest eigenvalue.

**Induction:** PC $k$  is the eigenvector for  $\lambda_k$  with orthogonality to all previous PCs.

---

The orthogonality constraint propagates cleanly because  $\Sigma$  is symmetric.

## SVD–PCA Equivalence: Full Proof

**Step 1 — SVD:**  $X_c = USV^T$  where  $U \in \mathbb{R}^{n \times n}$ ,  $S \in \mathbb{R}^{n \times p}$ ,  $V \in \mathbb{R}^{p \times p}$ .

**Step 2 — Gram matrix:**  $X_c^T X_c = (USV^T)^T (USV^T) = VS^T U^T USV^T = VS^2 V^T$   
(since  $U^T U = I$ , and  $S^T S = S^2$  collects  $s_k^2$  on diagonal)

**Step 3 — Covariance:**  $C = \frac{1}{n-1} X_c^T X_c = V \cdot \frac{S^2}{n-1} \cdot V^T$

This is the eigendecomposition of  $C$ : columns of  $V$  are eigenvectors,  $\frac{s_k^2}{n-1}$  are eigenvalues.

**Step 4 — Eigenvalues:**  $\lambda_k = s_k^2 / (n - 1)$ .

**Numerical note:**  $\text{cond}(X_c^T X_c) = \text{cond}(X_c)^2$ . SVD avoids this squaring, reducing round-off error for ill-conditioned data.

---

This equivalence is why `numpy.linalg.svd` is preferred over `numpy.linalg.eig` for PCA.

## t-SNE Gradient Derivation

**Starting from the KL objective:**  $\text{KL}(P\|Q) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}} = \sum_{i \neq j} p_{ij} \log p_{ij} - \sum_{i \neq j} p_{ij} \log q_{ij}$

**Substitute**  $q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{Z}$  where  $Z = \sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}$ :

$$-\sum_{i \neq j} p_{ij} \log q_{ij} = \sum_{i \neq j} p_{ij} \log(1 + \|y_i - y_j\|^2) + \log Z$$

**Differentiate w.r.t.**  $y_i$  and simplify:

$$\frac{\partial \text{KL}}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)(1 + \|y_i - y_j\|^2)^{-1}$$

**Barnes-Hut approximation:** Reduces  $O(n^2)$  repulsive term to  $O(n \log n)$  via space-partitioning trees (quadtree in 2D, octree in 3D).

---

The gradient has a physical interpretation: spring forces (attractive) vs. Coulomb repulsion.

**Perplexity** is defined via entropy of the conditional distribution:

$$\text{Perp}(P_i) = 2^{H(P_i)} \quad \text{where } H(P_i) = - \sum_j p_{j|i} \log_2 p_{j|i}$$

**Finding  $\sigma_i$ :** For each point  $i$ , find  $\sigma_i$  such that  $\text{Perp}(P_i)$  equals the user-specified perplexity. Solved via **binary search** on  $\sigma_i$ .

**Adaptive bandwidth:**

- Dense regions  $\rightarrow$  small  $\sigma_i$  (tight Gaussian)
- Sparse regions  $\rightarrow$  large  $\sigma_i$  (wide Gaussian)
- Each point adapts to its local density

**Why this matters:** Without adaptive  $\sigma_i$ , sparse-region points would have no meaningful neighbors, and dense-region points would consider too many.

---

Binary search for  $\sigma_i$  typically converges in 20–50 iterations per point.

# Kernel PCA

**Idea:** Apply PCA in a feature space  $\phi(x)$  without explicitly computing  $\phi$ .

**Kernel matrix:**  $K_{ij} = \kappa(x_i, x_j) = \langle \phi(x_i), \phi(x_j) \rangle$

**Eigendecomposition:** Solve  $K\alpha = n\lambda\alpha$  (dual formulation).

**Common kernels:**

- **Polynomial:**  $\kappa(x, y) = (x^T y + c)^d$
- **RBF (Gaussian):**  $\kappa(x, y) = \exp(-\gamma \|x - y\|^2)$
- **Sigmoid:**  $\kappa(x, y) = \tanh(\alpha x^T y + c)$

**Complexity:**  $O(n^2)$  for kernel matrix,  $O(n^3)$  for eigendecomposition — prohibitive for large  $n$ . Approximations: Nystrom method, random Fourier features.

---

Kernel PCA bridges PCA and t-SNE: non-linear but with a well-defined mathematical framework.

**Trustworthiness**  $T(k)$  — are embedded neighbors true neighbors?

$$T(k) = 1 - \frac{2}{nk(2n - 3k - 1)} \sum_{i=1}^n \sum_{j \in \mathcal{U}_k(i)} (r(i, j) - k)$$

where  $\mathcal{U}_k(i)$  = points in  $k$ -NN of  $i$  in low-D but *not* in high-D, and  $r(i, j)$  = rank of  $j$  w.r.t.  $i$  in high-D.

**Continuity**  $C(k)$  — are true neighbors still embedded neighbors?

$$C(k) = 1 - \frac{2}{nk(2n - 3k - 1)} \sum_{i=1}^n \sum_{j \in \mathcal{V}_k(i)} (\hat{r}(i, j) - k)$$

where  $\mathcal{V}_k(i)$  = points in  $k$ -NN of  $i$  in high-D but *not* in low-D.

**Interpretation:** Both range  $[0, 1]$ ; higher is better. Use  $k = 10\text{--}50$ . Compare across methods, perplexities, and random seeds.

---

Trustworthiness penalizes “false neighbors”; continuity penalizes “missing neighbors” in the embedding.

# References and Further Reading

## Advanced Theory:

- Jolliffe, I.T. & Cadima, J. (2016). Principal Component Analysis: A Review and Recent Developments. *Phil. Trans. R. Soc. A*, 374.
- Hinton, G. & Roweis, S. (2002). Stochastic Neighbor Embedding. *NeurIPS*.
- Kobak, D. & Berens, P. (2019). The Art of Using t-SNE for Single-Cell Transcriptomics. *Nature Communications*.

## Finance Applications:

- Litterman, R. & Scheinkman, J. (1991). Common Factors Affecting Bond Returns.
- Alexander, C. (2008). *Market Risk Analysis, Vol. I: Quantitative Methods in Finance*. Wiley.
- Lopez de Prado, M. (2018). *Advances in Financial Machine Learning*. Wiley. Ch. 8: Feature Importance.

---

Lopez de Prado (2018) discusses PCA denoising for covariance matrices in portfolio optimization.