

## L05: PCA & t-SNE

Deep Dive: Theory, Implementation, and Applications

Methods and Algorithms

Spring 2026

# Outline

1 Problem

2 Method

3 Solution

4 Practice

5 Decision Framework

6 Implementation

7 Summary

### Principal Component Analysis (PCA)

- Find orthogonal directions of maximum variance
- Project data onto these directions
- Reduce dimensions while preserving information

### Key Properties:

- Linear transformation
- Components are uncorrelated
- Partially reversible (lossy reconstruction when  $k < p$ )

---

PCA: one of the most fundamental tools in data science

## Covariance Matrix:

$$\Sigma = \frac{1}{n-1} X_c^T X_c \quad \text{where } X_c = X - \bar{X} \text{ (mean-centered)}$$

## Eigendecomposition:

$$\Sigma v = \lambda v$$

where  $v$  = eigenvector (principal direction),  $\lambda$  = eigenvalue (variance)

## Projection:

$$Z = X_c W_k \quad \text{where } X_c \text{ is centered, } W_k = [v_1, \dots, v_k]$$

---

Eigenvalues tell us how much variance each component captures

# Why Eigenvectors? Optimality Proof

**Find direction  $w$  maximizing variance:**

$$\max_w w^T \Sigma w \quad \text{subject to } \|w\| = 1$$

**Lagrangian:**

$$L = w^T \Sigma w - \lambda(w^T w - 1)$$

**First-order condition:**

$$\nabla_w L = 2\Sigma w - 2\lambda w = 0$$

**Result:**

$$\Sigma w = \lambda w \quad (\text{eigenvalue equation})$$

**Maximum variance = largest eigenvalue  $\lambda_1$**

---

The constrained optimization proof shows eigenvectors are optimal

## SVD–PCA Equivalence: The Computational Connection

**Theorem:** The principal components of  $X_c$  are the right singular vectors of  $X_c$ .

**Proof sketch:**

1. The covariance matrix:  $C = \frac{1}{n-1} X_c^\top X_c$
2. SVD of  $X_c$ :  $X_c = U\Sigma V^\top$
3. Then:  $X_c^\top X_c = V\Sigma^\top U^\top U\Sigma V^\top = V\Sigma^2 V^\top$
4. Therefore:  $C = \frac{1}{n-1} V\Sigma^2 V^\top$
5. Since  $CV = V \cdot \frac{\Sigma^2}{n-1}$ , columns of  $V$  are eigenvectors of  $C$

**Eigenvalues:**  $\lambda_k = \sigma_k^2 / (n - 1)$  where  $\sigma_k$  are singular values

**Computational advantage:** SVD is numerically more stable than eigendecomposition of  $X_c^\top X_c$  (avoids squaring condition number)

---

Practical implementations (scikit-learn, R) use SVD internally, not eigendecomposition.

## Proportion of Variance:

$$\text{Explained Variance Ratio}_i = \frac{\lambda_i}{\sum_{j=1}^p \lambda_j}$$

## Cumulative Variance:

$$\text{Cumulative}_k = \sum_{i=1}^k \frac{\lambda_i}{\sum_{j=1}^p \lambda_j}$$

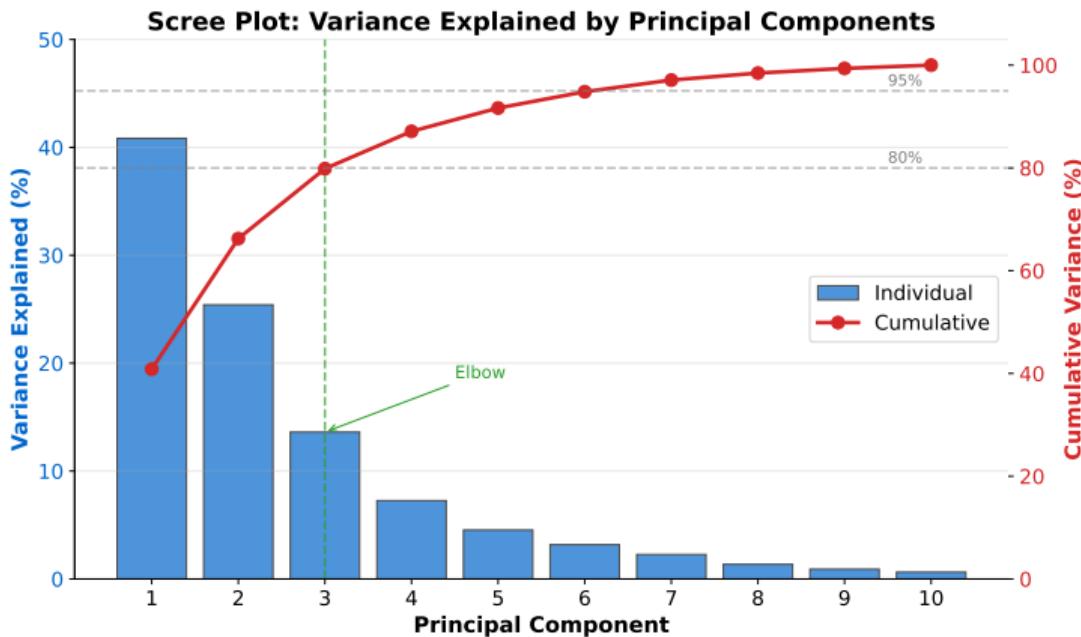
## Rules of thumb for choosing k:

- Keep 80-95% of total variance
- Use scree plot “elbow” method
- Kaiser criterion:  $\lambda > 1$  (valid for correlation matrix / standardized data only)

---

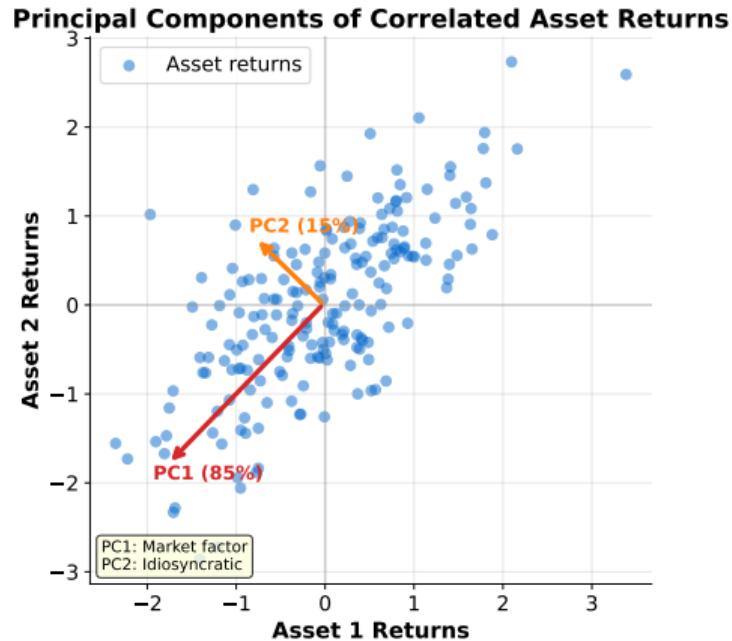
## Balance dimensionality reduction with information preservation

# Scree Plot



Look for the “elbow” where variance explained drops off

# Principal Components Visualization



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L05\\_PCA\\_TSNE/02\\_principal\\_components](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L05_PCA_TSNE/02_principal_components)

PC1 captures the dominant trend, PC2 the residual variation

From k components back to original space:

$$\hat{X} = ZW_k^T + \bar{X} \quad (\text{add mean back for original scale})$$

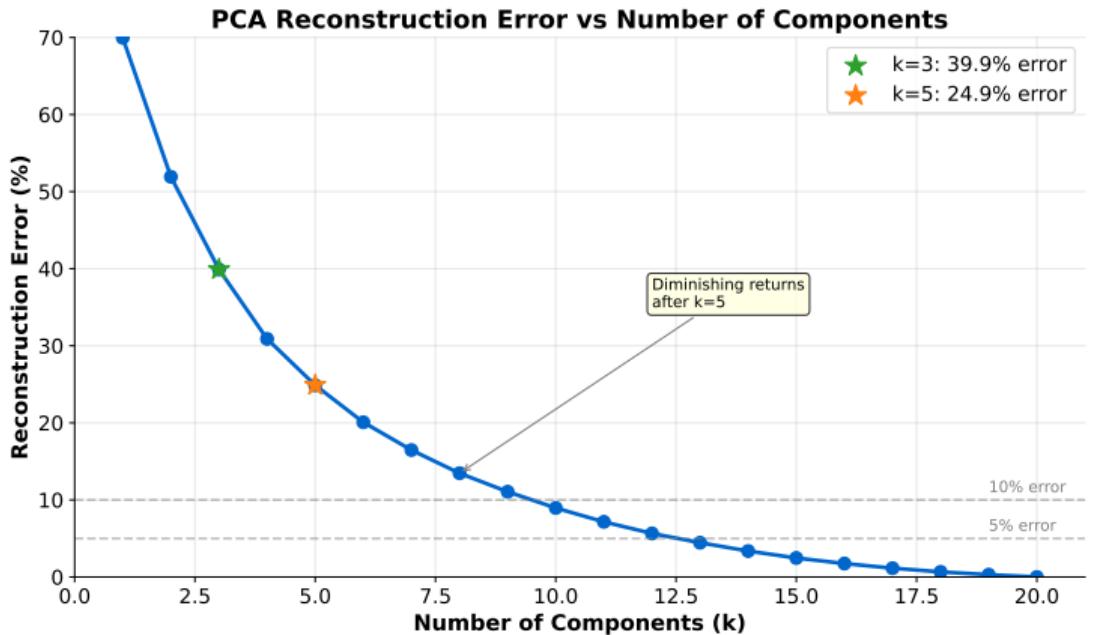
Reconstruction Error:

$$\text{Error} = \|X - \hat{X}\|_F^2 = \sum_{i=k+1}^p \lambda_i$$

---

Reconstruction error = sum of discarded eigenvalues

# Reconstruction Error vs Components



Adding more components always reduces error (but diminishing returns)

## Bootstrap Confidence Intervals:

- Resample data, recompute PCA, track loading stability

## Parallel Analysis:

- Compare eigenvalues to random data
- Keep components where  $\lambda_{data} > \lambda_{random}$

## Cross-Validation for k:

- Split data, train PCA, test reconstruction on holdout

## For t-SNE:

- Run multiple times with different seeds
- Unstable clusters may be artifacts

---

MSc-level: always quantify uncertainty in dimensionality reduction

### Portfolio Risk Decomposition:

- PC1 often represents “market factor”
- PC2-3 may capture sector/size factors
- Higher PCs: idiosyncratic risk

### Example: 10-Stock Portfolio

- PC1 loadings  $\approx [0.31, 0.32, 0.30, \dots]$  — uniform (market factor)
- PC2: tech stocks positive, banks negative (sector rotation)
- First 3 PCs explain  $\sim 75\%$  of portfolio variance

### Applications:

- Risk factor modeling
- Dimensionality reduction for trading signals
- Noise reduction in time series
- Feature extraction for ML models

---

PCA reveals latent structure in financial data

## The Canonical Finance Example

- Yield curves decompose into ~3 principal components
- PC1 = **Level** (parallel shift): explains ~85% variance
- PC2 = **Slope** (steepening/flattening): explains ~10% variance
- PC3 = **Curvature** (butterfly): explains ~3% variance

**Together explain 98%+ of yield curve movements**

- Used daily in every bank's risk management system
- Foundation for interest rate hedging strategies

---

**Level/slope/curvature: industry-standard yield curve decomposition**

## Typical Loadings Pattern:

- **PC1 (Level):** Flat loadings across maturities—all rates move together
- **PC2 (Slope):** Negative short end, positive long end—curve steepens/flattens
- **PC3 (Curvature):** Negative at ends, positive in middle—butterfly trades

## Trading Applications:

- Duration-neutral hedging (hedge PC1 exposure)
- Curve trades (exploit PC2 movements)
- Butterfly spreads (exploit PC3 movements)

---

PCA on interest rates is foundational fixed-income knowledge

## When PCA Falls Short:

- Non-linear relationships (curved manifolds)
- Cluster structure not aligned with variance
- Discrete or categorical data
- Outliers heavily influence results

## Solutions:

- Kernel PCA (implicit non-linear mapping via kernel trick)
- Robust PCA (outlier-resistant)
- t-SNE/UMAP (for visualization)

---

PCA assumes linear structure; Gaussian NOT required but PCA is optimal for Gaussian data

### t-Distributed Stochastic Neighbor Embedding

- Non-linear dimensionality reduction
- Optimized for visualization (2D/3D)
- Preserves local neighborhood structure

#### Key Idea:

- Convert distances to probabilities
- In high-D: Gaussian similarities
- In low-D: t-distribution similarities
- Minimize KL divergence between distributions

---

t-SNE: visualization method, NOT for preprocessing

**High-dimensional similarity:**

$$p_{j|i} = \frac{\exp(-\|x_i - x_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|x_i - x_k\|^2 / 2\sigma_i^2)}$$

**Symmetrize:**  $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$  (joint probability)

**Low-dimensional similarity (t-distribution):**

$$q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq l} (1 + \|y_k - y_l\|^2)^{-1}}$$

**Objective: Minimize KL divergence**

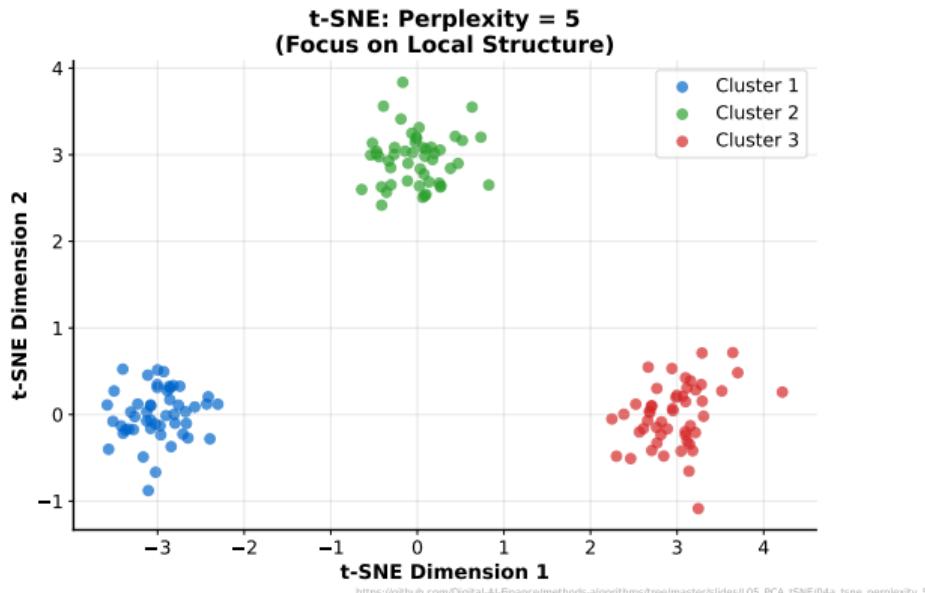
$$KL(P||Q) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

Note: KL is asymmetric—penalizes large  $p_{ij}$  with small  $q_{ij}$  more than reverse. This preserves local structure (nearby points stay nearby).

---

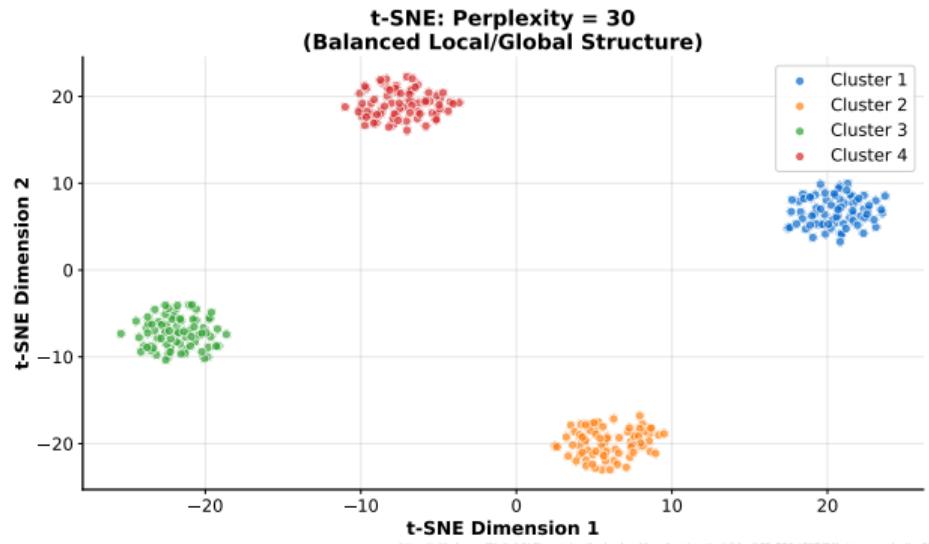
t-distribution solves the crowding problem: in high-D, moderate distances become small in low-D, causing points to collapse. Heavy tails allow dissimilar points to spread apart.

## Perplexity: Low (Local Focus)



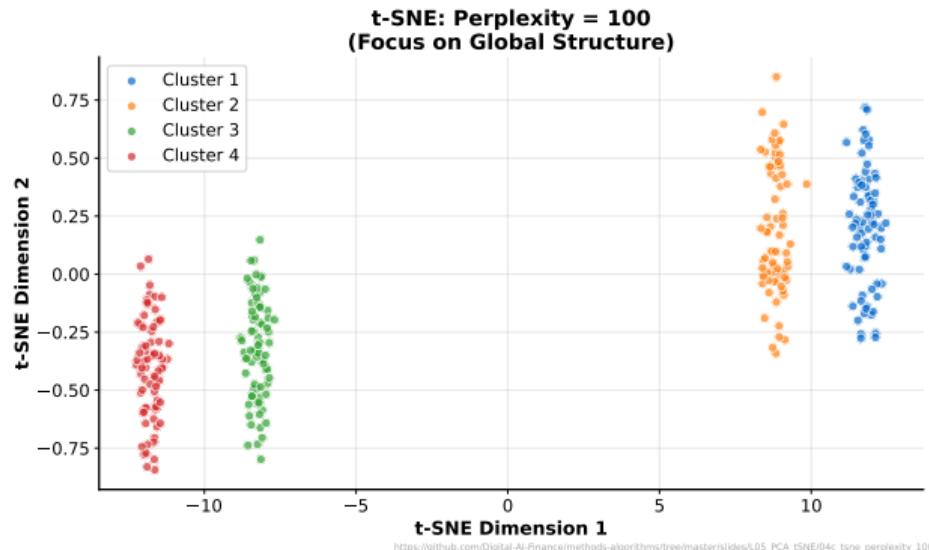
Low perplexity (5): tight clusters, focus on nearest neighbors only

## Perplexity: Default (Balanced)



Default perplexity (30): balanced local and global structure

## Perplexity: High (Global Focus)



High perplexity (100): more spread, clusters may merge

# Perplexity Guidelines

**Perplexity** controls the balance between local and global structure:

Perplexity =  $2^H$  where  $H$  is entropy; roughly the effective number of neighbors considered.

- Low perplexity (5-10): Focus on very local structure
- Medium perplexity (30-50): Balanced (default)
- High perplexity (100+): More global structure

## Guidelines:

- Should be smaller than number of points
- Larger datasets can use higher perplexity
- Run multiple perplexities to validate findings

---

Results can vary significantly with perplexity choice

## Important Limitations:

- Non-deterministic (run multiple times)
- Cluster sizes are not meaningful
- Distances between clusters are not meaningful
- Slow for large datasets ( $O(n^2)$ ); Barnes-Hut approximation gives  $O(n \log n)$ )

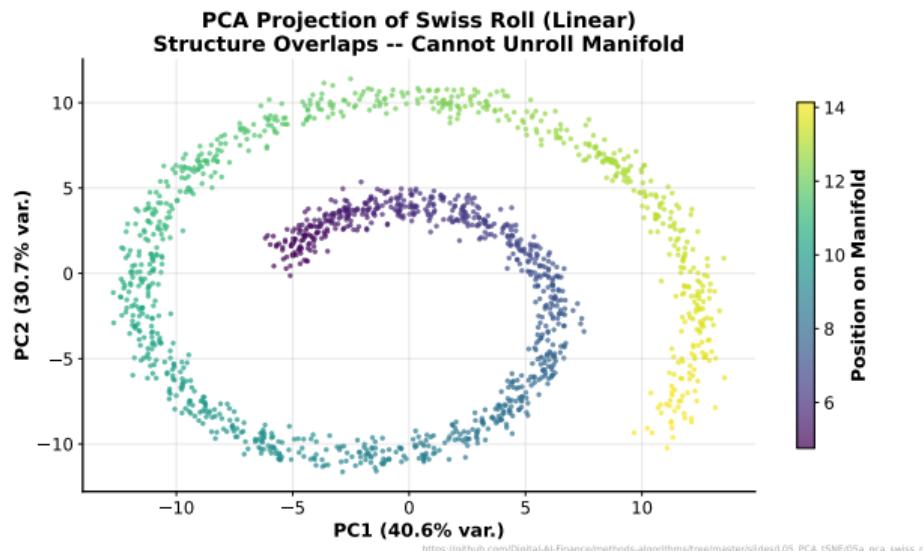
## Best Practices:

- Use PCA first to reduce to 30-50 dims
- Run multiple times with different seeds
- Don't over-interpret cluster sizes/distances
- Use for exploration, not final conclusions

---

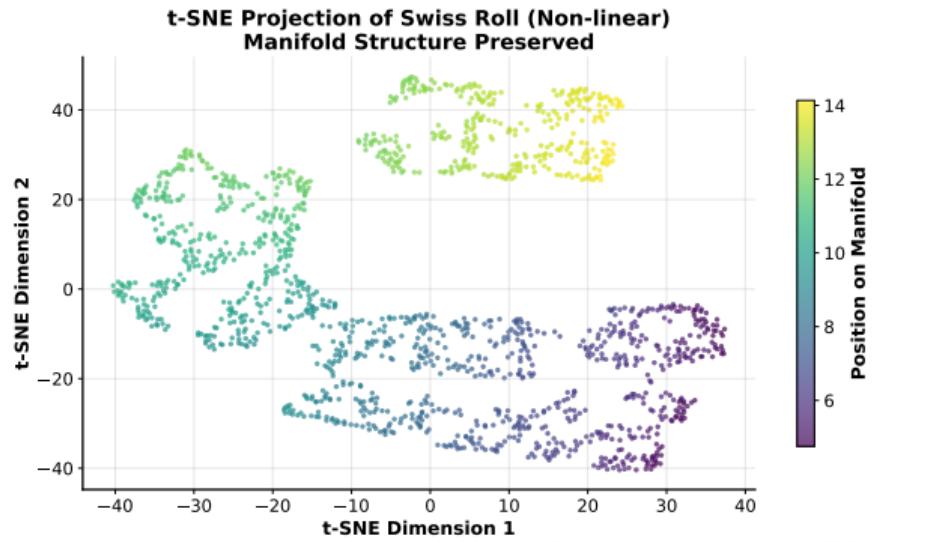
t-SNE shows IF clusters exist, not HOW they relate

## Part 4: PCA on Non-linear Manifolds



**PCA (linear) cannot unroll the Swiss roll - structure overlaps**

# t-SNE on Non-linear Manifolds



**t-SNE (non-linear)** successfully unrolls the manifold structure

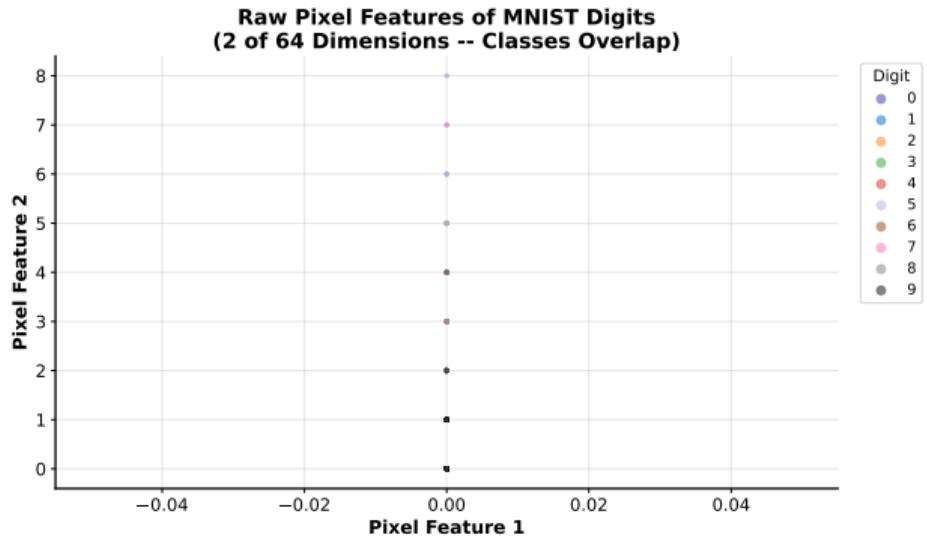
## Comparison Table

Aspect	PCA	t-SNE
Type	Linear	Non-linear
Speed	Fast $O(np^2)$	Slow $O(n^2)$
Deterministic	Yes	No
Preserves	Global variance	Local neighbors
Reversible	Lossy if $k < p$	No
Use for ML	Yes (preprocessing)	No
Visualization	Okay	Excellent

---

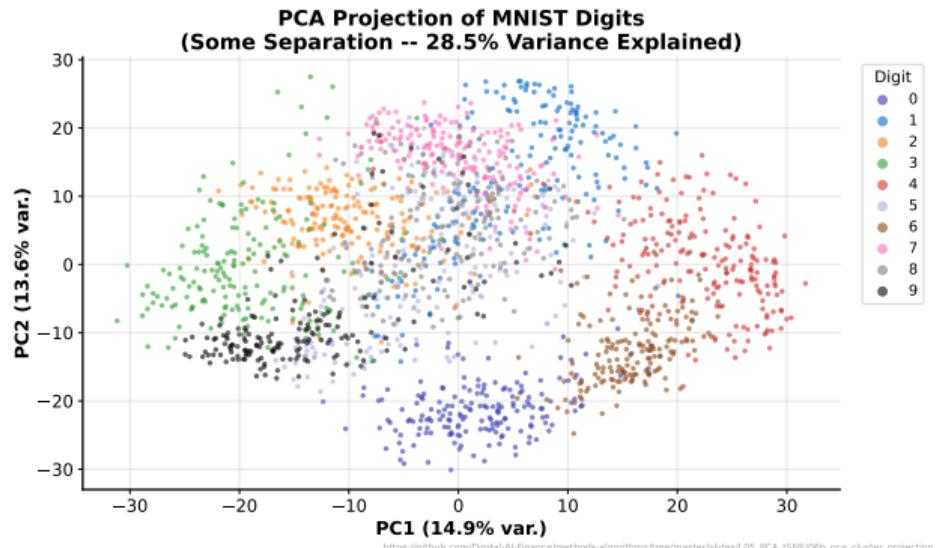
Use PCA for preprocessing, t-SNE for visualization only

## Cluster Preservation: Original Data



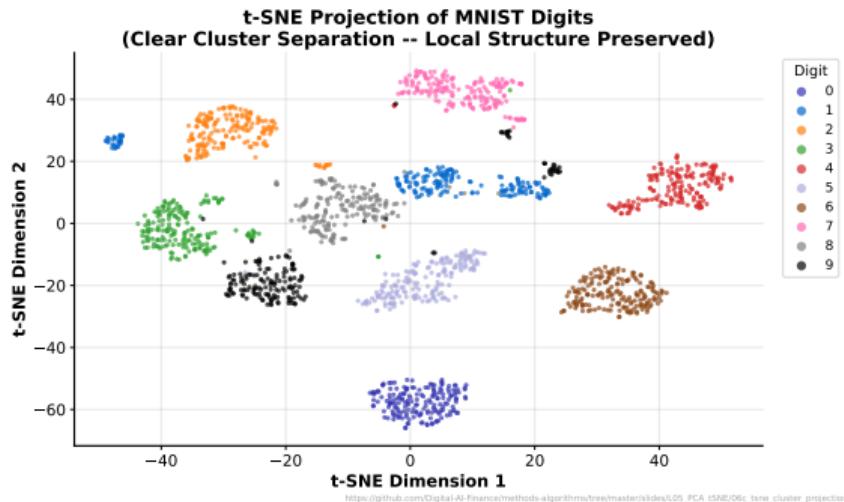
MNIST digits (64 dims): classes overlap when viewed in raw pixel space

## Cluster Preservation: PCA Projection



PCA finds max-variance directions – partial digit separation

## Cluster Preservation: t-SNE Projection



**t-SNE preserves local neighborhoods – clear digit cluster separation**

# When to Use Which

## Use PCA When:

- Preprocessing for ML (reduce features)
- Linear relationships expected
- Need reversibility (reconstruction)
- Speed matters

## Use t-SNE When:

- Visualizing high-dimensional data
- Looking for cluster structure
- Non-linear manifolds expected
- Exploratory analysis

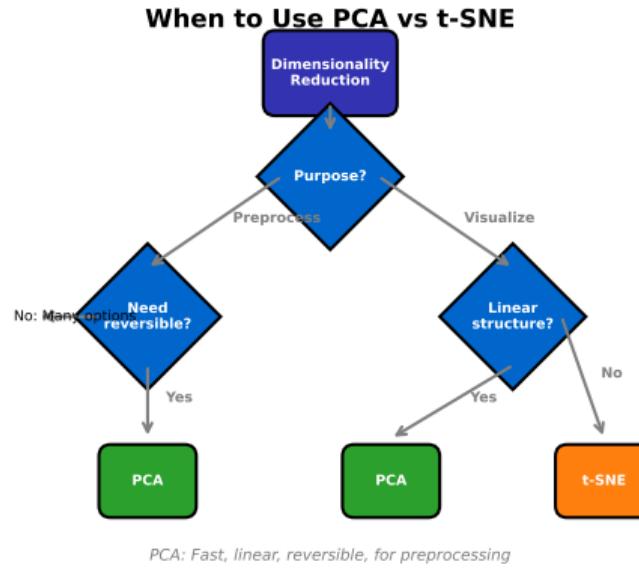
---

Often use both: PCA first to 30-50 dims, then t-SNE for visualization

### Open the Colab Notebook

- Exercise 1: Apply PCA to high-dimensional finance data
- Exercise 2: Visualize clusters with t-SNE
- Exercise 3: Compare PCA vs t-SNE for different datasets

Link: <https://colab.research.google.com/> See course materials



[https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L05\\_PCA\\_tSNE/07\\_decision\\_flowchart](https://github.com/Digital-AI-Finance/methods-algorithms/tree/master/slides/L05_PCA_tSNE/07_decision_flowchart)

---

Consider purpose: preprocessing (PCA) vs visualization (t-SNE)

## Part 5: Implementation

### PCA in scikit-learn:

- `PCA(n_components=k)`: Keep k components
- `PCA(n_components=0.95)`: Keep 95% variance
- `pca.explained_variance_ratio_`: Variance per component
- `pca.inverse_transform()`: Reconstruct original
- Note: sklearn uses truncated SVD internally (more stable than eigendecomposition)

### t-SNE in scikit-learn:

- `TSNE(n_components=2, perplexity=30)`
- Always normalize data first
- Consider PCA preprocessing for speed

---

Standardize data before PCA; normalize before t-SNE

## Uniform Manifold Approximation and Projection

- Faster than t-SNE
- Better preserves global structure
- Can embed new points (unlike t-SNE)
- Hyperparameters: n\_neighbors, min\_dist

### When to use UMAP:

- Large datasets (faster than t-SNE)
- Need to embed new data points
- Want more preserved global structure

---

UMAP increasingly popular; see McInnes et al. (2018) for comparison

## PCA:

- Linear, fast, reversible
- Use for preprocessing and feature extraction
- Choose k by variance explained or elbow

## t-SNE:

- Non-linear, slow, visualization-only
- Excellent for exploring cluster structure
- Don't interpret distances or sizes literally

**Common Pipeline:** Standardize → PCA (30-50) → t-SNE (2D)

---

**Next:** Embeddings and Reinforcement Learning

# References

## Textbooks:

- James et al. (2021). *ISLR*, Chapter 12: Unsupervised Learning
- Hastie et al. (2009). *ESL*, Chapter 14: Unsupervised Learning

## Original Papers:

- Pearson (1901). On Lines and Planes of Closest Fit
- van der Maaten & Hinton (2008). Visualizing Data using t-SNE
- McInnes et al. (2018). UMAP

## Documentation:

- scikit-learn: `sklearn.decomposition.PCA`
- scikit-learn: `sklearn.manifold.TSNE`

---

t-SNE paper: one of the most influential visualization papers